

Autonomous Vehicles
Assignment I: Use of ROS bags

Tommaso Bocchietti 10740309

A.Y. 2024/25



POLITECNICO
MILANO 1863

Contents

| | | |
|----------|---|----------|
| 1 | Introduction | 3 |
| 2 | Assignment Part A | 3 |
| 2.1 | Request | 3 |
| 2.2 | Analysis | 3 |
| 2.2.1 | Minimum distance from obstacles | 4 |
| 2.2.2 | Estimate of /cmd_vel command | 5 |
| 3 | Assignment Part B | 5 |
| 3.1 | Request | 6 |
| 3.2 | Analysis | 6 |
| 4 | Conclusions | 8 |

List of Figures

| | | |
|---|---|---|
| 1 | Gazebo world used for the simulation. Credit to https://emanual.robotis.com/ | 3 |
| 2 | Minimum distance from obstacles over time. | 4 |
| 3 | Velocities of the robot over time, purely based on the odometry data. | 5 |
| 4 | Comparison of the trajectories performed by the original simulation and the one performed by the replayed commands. | 7 |
| 5 | Comparison of the velocities performed by the original simulation and the one performed by the replayed commands. | 8 |

1 Introduction

The aim of this work is to gain insight into the use of `rosbags` for the analysis of data collected by autonomous vehicles.

The following sections provide a detailed description of the requests associated with this assignment, the approach taken to fulfill them, and the discussion of the results obtained.

Tools As for the tools used, `ROS1` (Robot Operating System) is employed as the main framework for data collection, while their analysis is performed using `MATLAB`. Notice that with the current setup used by the author, `MATLAB 2024a` is running in Windows 10, while `ROS1` is running in the `WSL2` (Windows Subsystem for Linux) environment, specifically with the `Ubuntu 20.04` distribution.

2 Assignment Part A

In this section, we provide a brief overview of the requests associated with the first part of the assignment, along with a description of the approach taken to fulfill them. Discussion of the results obtained is also included.

2.1 Request

Starting from the data collected in the provided `rosbag` file, the goal of this first part of assignment is to evaluate at each time instant, the following quantities:

- **Minimum distance** of the vehicle from the obstacles in the environment;
- **Estimate of /cmd.vel command** sent to the vehicle during the simulation.

We also recall that the `rosbag` file was obtained by running a `Turtlebot3` robot of the model `burger` in a simulated environment, specifically the `turtlebot3_world` provided by the `turtlebot3_gazebo` package.

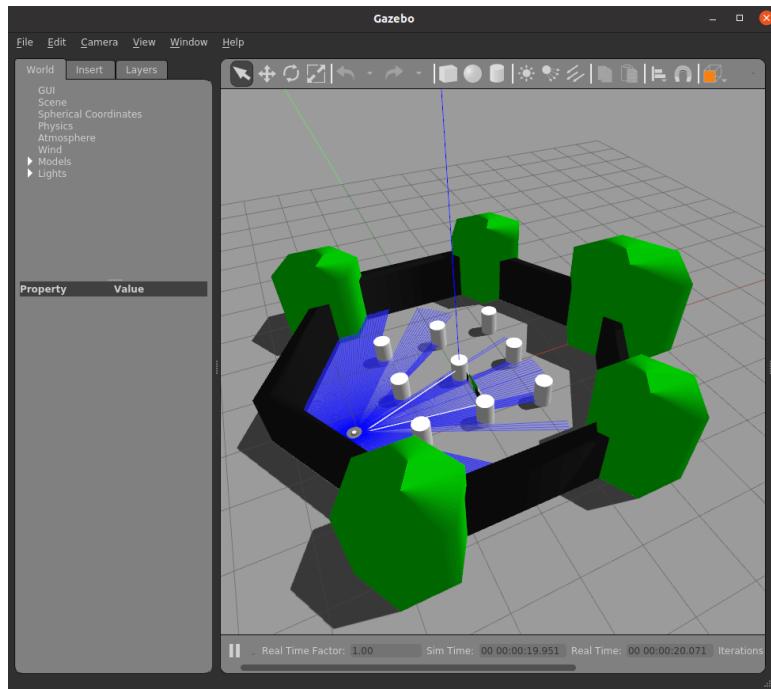


Figure 1: Gazebo world used for the simulation. Credit to <https://emanual.robotis.com/>

2.2 Analysis

At first, the `rosbag` file is loaded into `MATLAB` using the `rosbag` function, which allows to read and manipulate `rosbag` files in a convenient way. A preliminary analysis of the data contained in the `rosbag` file is performed showing the following topics:

| | NumMessages | MessageType |
|--------|--------------------|-----------------------|
| /clock | 110,710 | rosgraph_msgs/Clock |
| /imu | 110,240 | sensor_msgs/Imu |
| /odom | 3,316 | nav_msgs/Odometry |
| /scan | 552 | sensor_msgs/LaserScan |
| /tf | 3,316 | tf2_msgs/TFMessage |

Table 1: Topics contained in the `rosbag` file.

2.2.1 Minimum distance from obstacles

Given the presence of the `/scan` topic, which contains the laser scan data, we can use it to compute the minimum distance from obstacles in the environment. The `scan` message contains a field called `ranges`, which is an array of distances measured by the laser scanner at different angles. The minimum distance can be computed by taking the minimum value of this array, which represents the closest obstacle detected by the laser scanner at that time instant. The following code snippet shows how to extract the `ranges` field from the `/scan` topic and compute the minimum distance at each time instant:

```

1 scan = select(bag, 'Topic', '/scan');
2 scan_msgs = readMessages(scan, 'DataFormat', 'struct');
3 scan_time = scan.MessageList.Time - scan.StartTime;
4 scan_ranges = cell2mat(cellfun(@(msg) msg.Ranges(:)', scan_msgs, 'UniformOutput', false)
    );
5 min(scan_ranges, [], 2)
```

Listing 1: Extracting the `ranges` field from the `/scan` topic and computing the minimum distance at each time step.

One can also plot the minimum distance over time, as shown in Figure 2.

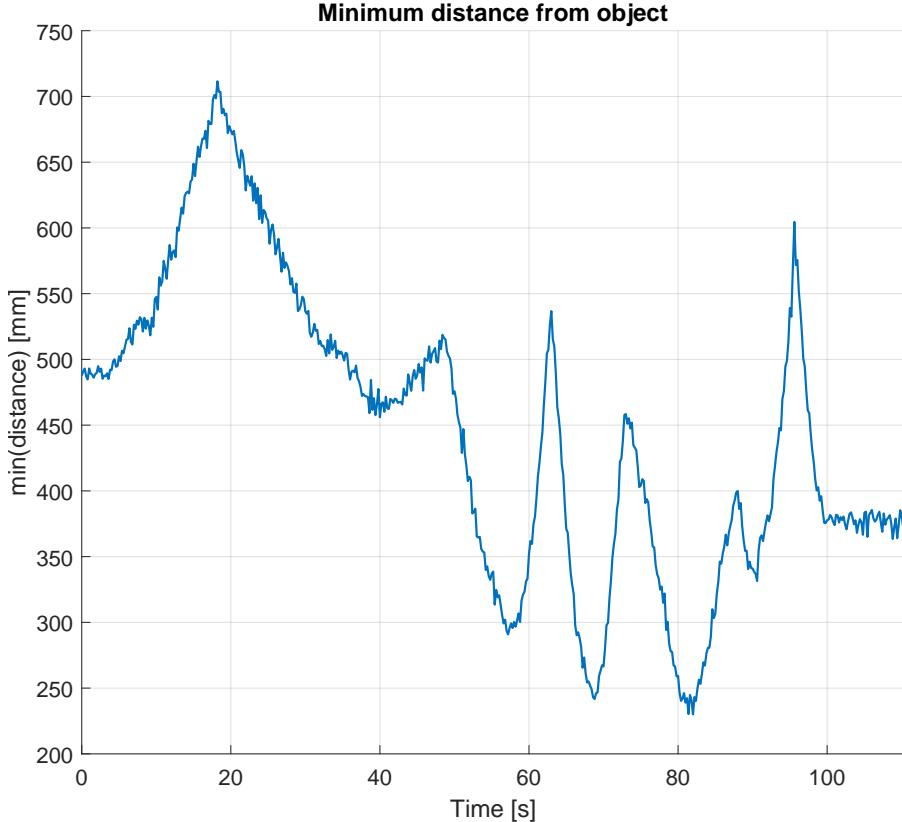


Figure 2: Minimum distance from obstacles over time.

It's important to notice that, being the data coming from a simulated environment, we can expect the accuracy of the laser scanner to be almost perfect (unless the simulation environment add noise on purpose), while on a real vehicle we could face issues like sensor noise or hit of minimum and/or maximum range of the laser scanner.

This could lead to a less accurate estimation of the minimum distance from obstacles, which could be a problem for the robot navigation and obstacle avoidance.

2.2.2 Estimate of /cmd_vel command

The /cmd_vel topic usually contains the velocity commands sent to the robot, which are typically represented as linear and angular velocities.

As an exercise, given that in the provided rosbag file the /cmd_vel topic has been removed, we can estimate the linear and angular velocities of the robot using the /odom topic, which contains the odometry data. The odom message contains the position and orientation of the robot in the world frame, as measured by the odometry system.

Again, with a similar fashion as done for the /scan topic, we can extract the information from the /odom topic, save them in a matrix form and plot them.

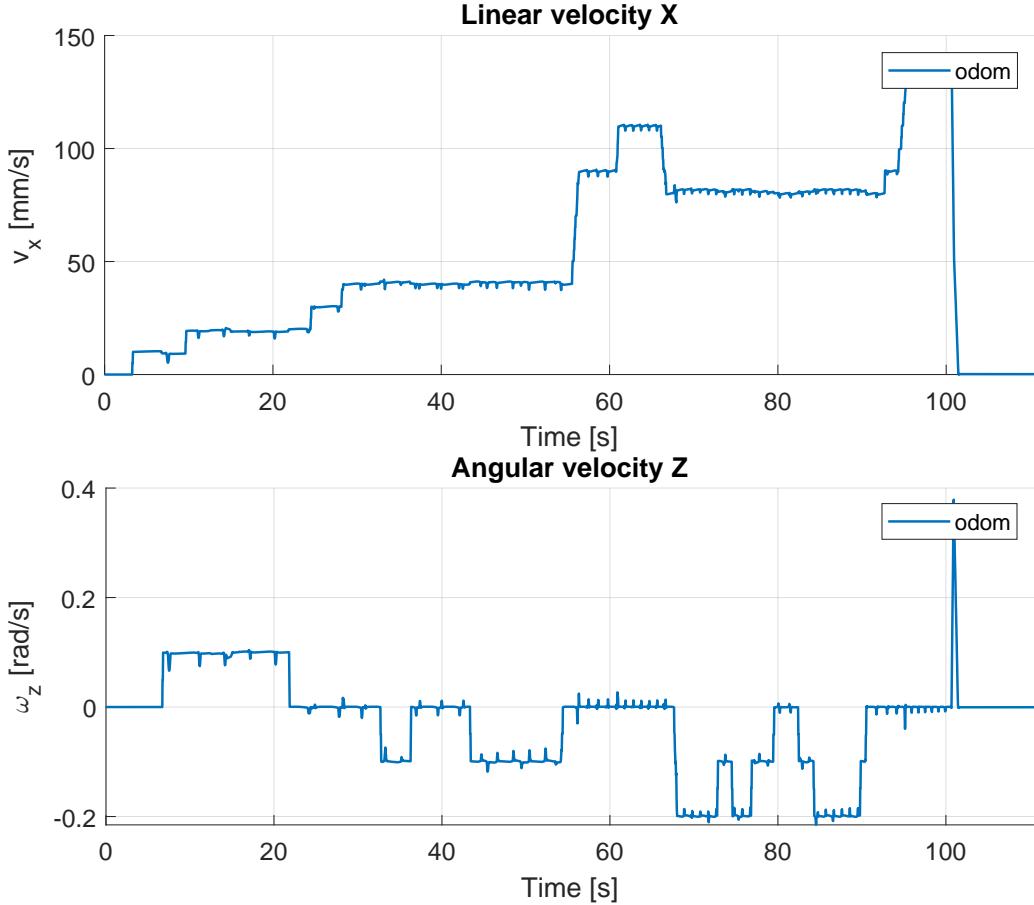


Figure 3: Velocities of the robot over time, purely based on the odometry data.

Notice that one could have also possibly used the /imu topic to estimate both position and velocities of the robot (via single and double integration of the acceleration data). However, due to the presence of noise in the IMU data and cumulative errors due to integration, this approach is definitely not recommended as it easily results in drifted output estimation.

On the other hand, a sensor fusion approach based for example on the Kalman filter could have been used to combine the two sources of data (IMU and odometry) ultimately increasing the overall accuracy.

3 Assignment Part B

In this section, we provide a brief overview of the requests associated with the second part of the assignment, along with a description of the approach taken to fulfill them. Discussion of the results obtained is also included in the following sections.

3.1 Request

Given that in the previous section we have estimated the control commands sent to the robot during the original simulation, in this second part we are asked to evaluate the accuracy of such commands.

3.2 Analysis

In order to evaluate the accuracy of the control commands sent to the robot during the simulation, we need to rerun a simulation by our own, using the same world and the same robot model, and then compare both the trajectory and the velocities with the ones saved in the original `rosbag` file.

We recall that our working environment is setup so that `ROS1` is running in the `WSL2` (Windows Subsystem for Linux) environment, specifically in `Ubuntu 20.04`, while `MATLAB 2024a` is running in Windows 10. Based on these considerations, we can run the simulation in at least three different ways:

- Create a publisher script in `MATLAB` that sends the commands to `ROS1` running in `WSL2`;
- Create a publisher in `Simulink` that sends the commands to `ROS1` running in `WSL2`;
- Perform the analysis directly in `WSL2` using the native commands of `rosbag` to automatically replay the estimated commands.

At first, we tried to create a publisher script in `MATLAB` that was sending the commands over the bridged network, pacing the sender based on the original timing of the `/odom` topic. However, we encountered some major issues with the timing of the commands, which were sent at almost unpredictable intervals, leading to a completely different behavior of the robot in the simulation.

In order to avoid or at least reduce the issues related to the timing of the commands, we decided to directly move to the third option, which is to perform the analysis directly in `WSL2` using the native commands of `rosbag` to automatically replay the estimated commands. At first, the `/odom` topic has been ported to a new `rosbag` file using a simple `python` script, which is able to read the original `rosbag` file and write the `/odom` topic to a new `rosbag` file while also reshaping the data to match the type of the `/cmd_vel` topic. In particular, the following relationship table has been implemented:

| <code>/odom</code> | <code>/cmd_vel</code> |
|--|----------------------------|
| <code>msg.Twist.Twist.Linear.X</code> | <code>msg.Linear.X</code> |
| <code>msg.Twist.Twist.Angular.Z</code> | <code>msg.Angular.Z</code> |

Table 2: Relationship between the `/odom` and `/cmd_vel` topics.

Once the content of the `/odom` topic has been saved as `/cmd_vel` topic in a new `rosbag` file, we can use the `rosbag play` command to replay the commands in the simulation.

In Figure 4 we can see the comparison of the trajectories performed by the original simulation and the one performed by the replayed commands.

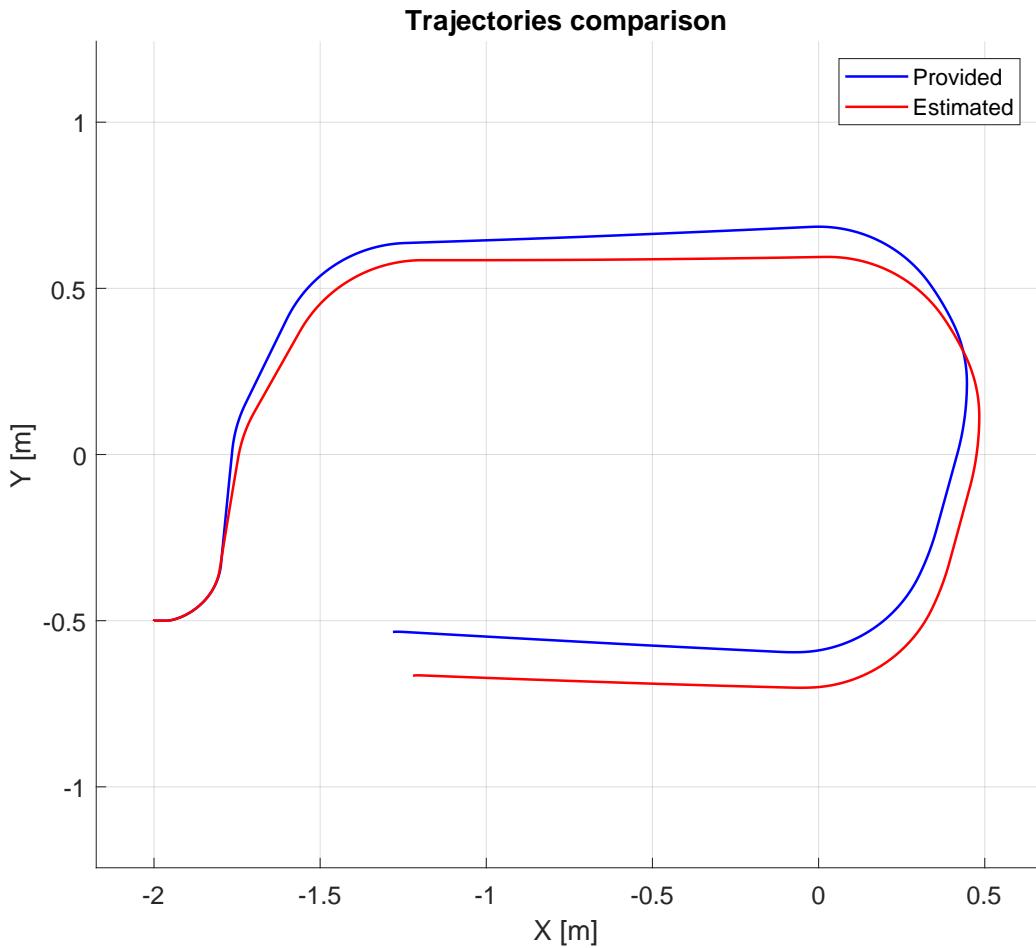


Figure 4: Comparison of the trajectories performed by the original simulation and the one performed by the replayed commands.

One can clearly see that the two trajectories are not completely overlapping.

Looking at the values of the actual velocities of the robot during the simulation, we can see a minor discrepancy between the original simulation and the one performed by the replayed commands, similarly to the behaviour observer for the trajectories. Figure 5 shows the comparison of the velocities performed by the original simulation and the one performed by the replayed commands.

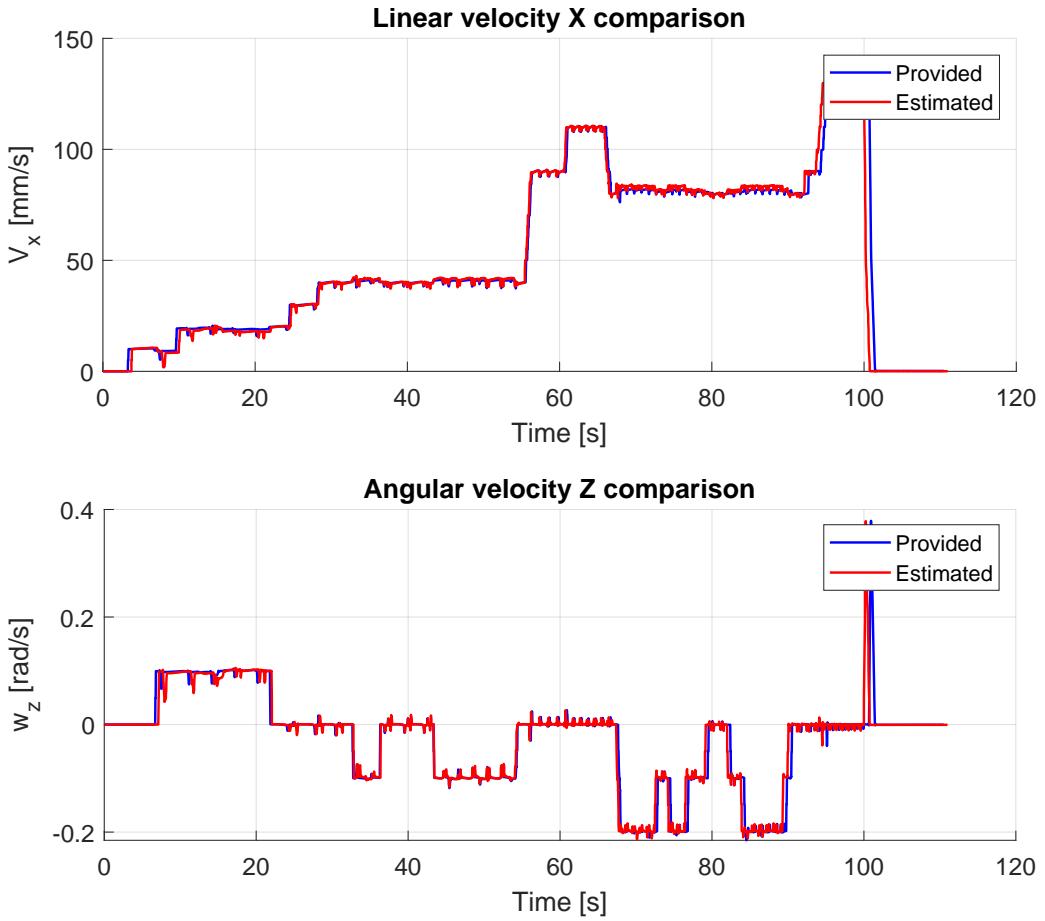


Figure 5: Comparison of the velocities performed by the original simulation and the one performed by the replayed commands.

Again, some minor issues related to timing can be observed. For some reason that the author is not able to determine, the `rosbag play` command is not able to replay the commands at the exact same pace as the original simulation, leading to a slight shift in the velocities.

What is also interesting to notice, are the juggling of the velocities along the whole simulation. These are, in fact, not desired behaviors and most probably caused by the simulated environment itself.

4 Conclusions

In this short report, we have presented the results of the first assignment of the course on Autonomous Vehicles. We have shown how to analyze data collected by a simulated robot in a Gazebo environment, using the `rosbag` command to store the data. We have also demonstrated how to estimate the control commands sent to the robot during the simulation, and how to evaluate their accuracy by comparing trajectories and odometer readings. Results obtained show that the estimated commands are representative of the original ones, but not completely overlapping due to poor timing performance caused (probably) by the `WSL2` environment that can't fully afford the real-time performance required by the simulation.

Further investigations and testing could be done by switching to a native `Linux` environment, which could provide better performance and more accurate results.

References

- [1] Stanford Artificial Intelligence Laboratory et al. Robotic operating system.

Autonomous Vehicles

Assignment II: Feedback control of a differential drive robot

Tommaso Bocchietti 10740309

A.Y. 2024/25



POLITECNICO
MILANO 1863

Contents

| | | |
|----------|--------------------------------|-----------|
| 1 | Introduction | 3 |
| 2 | System overview | 3 |
| 2.1 | Request | 3 |
| 2.2 | Waypoints | 3 |
| 2.3 | Simulink model | 4 |
| 3 | Proportional controller | 5 |
| 3.1 | Implementation | 6 |
| 3.2 | Results | 6 |
| 4 | Lyapunov controller | 8 |
| 4.1 | Lyapunov control law | 8 |
| 4.2 | Implementation | 9 |
| 4.3 | Results | 9 |
| 5 | Conclusions | 11 |

List of Figures

| | | |
|---|--|----|
| 1 | Waypoints used for the simulation. | 4 |
| 2 | Simulink model used for the implementation of the feedback control loop. | 4 |
| 3 | Inside view of the controller block. | 5 |
| 4 | Trajectory of the robot during the simulation with the proportional controller. | 7 |
| 5 | Velocity profiles of the robot during the simulation with the proportional controller. | 7 |
| 6 | Schematic representation of the unicycle-like system and the reference frame. | 8 |
| 7 | Trajectory of the robot during the simulation with the Lyapunov controller. | 10 |
| 8 | Velocity profiles of the robot during the simulation with the Lyapunov controller. | 11 |

1 Introduction

The aim of this work is to gain insight into the development of a feedback control loop for an autonomous vehicle, specifically a Turtlebot3 robot of the model **burger**. Two different control strategies have been implemented, namely a simple proportional controller and a non-linear controller based on Lyapunov stability theory. As for the simulation environment, the Turtlebot3 robot was simulated in a Gazebo world, specifically the `turtlebot3_world` provided by the `turtlebot3_gazebo` package.

The following sections provide a detailed description of the requests associated with this assignment, the approach taken to fulfill them, and the discussion of the results obtained.

Tools As for the tools used, **ROS1** (Robot Operating System) is employed as the main framework for communication between the different components of the system. **Simulink** is used as the main agent in the loop, sending control commands to the vehicles based on the telemetry data received from the vehicle itself. **MATLAB** instead is used to perform the analysis of the data collected during the simulation, and to visualize the results. Notice that with the current setup used by the author, **MATLAB 2024a** and **Simulink** are running in Windows 10, while **ROS1** is running in the **WSL2** (Windows Subsystem for Linux) environment, specifically in the **Ubuntu 22.04** distribution.

2 System overview

In this section, we provide a brief overview of the requests associated with this assignment, along with a description of the approach taken to fulfill them. In the successive two sections instead, we describe in detail the control strategies implemented, and the analysis of the results obtained.

2.1 Request

For this assignment, we are asked to implement a feedback control strategy that is able to drive a Turtlebot3 robot of the model **burger** through some predefined waypoints in the environment.

In particular, the waypoints are specified as a list of 3D coordinates in the world frame, where the first two coordinates represent the position of the vehicle in the plane, and the third coordinate represents the orientation of the vehicle in the world frame.

Also, along with this core request, we are asked to implement two control strategy such as:

- A simple proportional controller, subjected to a set of constraint;
- A more advanced controller (no particular constraints on its choice).

2.2 Waypoints

Before proceeding with the actual implementation of the control system, we show in Figure 1 the waypoints used for this assignment. These waypoints are provided along with the text of the assignment.

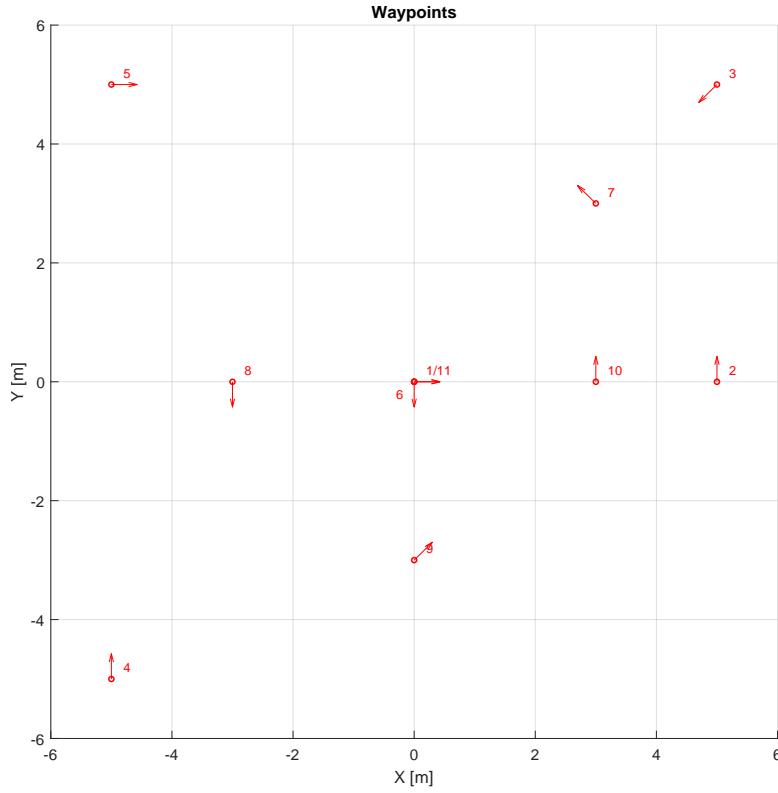


Figure 1: Waypoints used for the simulation.

In Figure 1, we can see the waypoints used for the simulation, which are represented as red dots in the world frame. The arrow associated to each waypoint represents the orientation of the vehicle at that waypoint, which is given by the third coordinate of the waypoint list.

2.3 Simulink model

As already mentioned in the previous section, the logic of the feedback control loop is implemented in **Simulink**, which is a graphical programming environment for modeling, simulating and analyzing dynamic systems. In Figure 2 we can see the Simulink model used for this assignment.

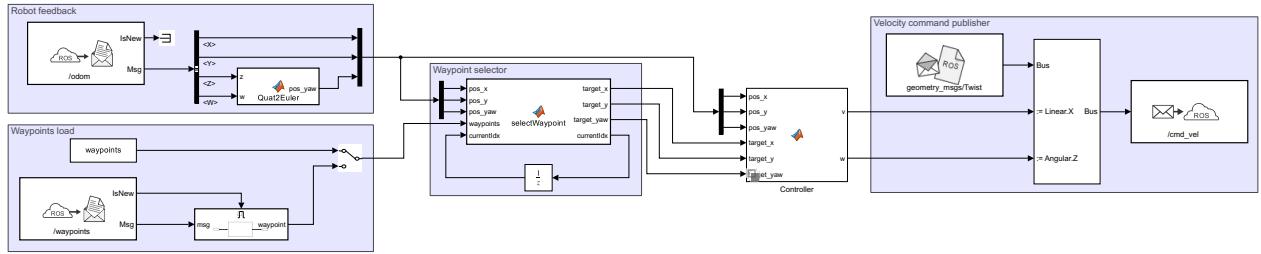


Figure 2: Simulink model used for the implementation of the feedback control loop.

One can easily recognize the different components of the system thanks to the colored areas used to group the different blocks. The main components of the system are:

- Sensors readings: these blocks are used to subscribe to the `/odom` topic and to extract the data needed for the control system;
- Waypoint loader: this block is used to load the waypoints from either a workspace variable or a custom topic `/waypoints` on which the list of waypoints is published;
- Waypoint selector: this block is used to select the current waypoint from the list of waypoints, based on the current position of the vehicle and its history;

- Controller: this block is responsible for the logic of the feedback control loop. It computes the control commands to be sent to the vehicle based on the current position and orientation of the vehicle, and the current waypoint to be reached. The controller is implemented as a MATLAB function block, which allows for a more flexible implementation of the control logic. More details on the implementation of the controller are provided in the next sections;
- Command publisher: this block is used to publish the control commands to the `/cmd_vel` topic, which is used by the vehicle to receive the control commands.

Notice that the commands sent to the vehicle are comprehensive of both linear and angular velocities. In fact, a differential drive model is used internally by the robot to compute the wheel velocities based on the linear and angular velocities provided by the control system.

Inside the controller block, we can find the two different control strategies implemented for this assignment, namely a simple proportional controller and a more advanced controller based on Lyapunov analysis. See Figure 3 for an inside view of the controller block.

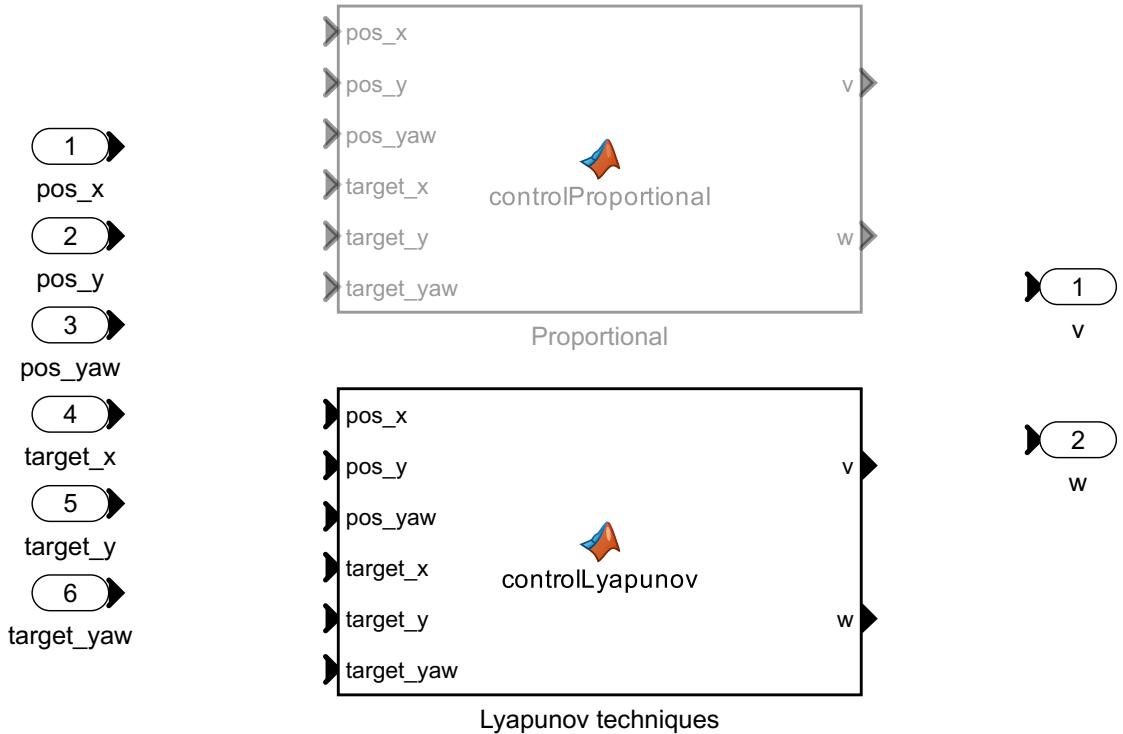


Figure 3: Inside view of the controller block.

3 Proportional controller

As explained in the previous section, one of the requests associated with this assignment is to implement a simple proportional controller for the `Turtlebot3` robot.

For this task, we have opted for a simple two-stage proportional controller. In particular, based on the distance of the robot from the current target waypoint, we apply a control input so that:

- If the robot is far from the target waypoint, we force its direction to be aligned with the target position, and we apply a linear velocity proportional to the distance from the target;
- Instead, when the robot is close to the target waypoint, we apply a control input that forces the robot to rotate in place, so that it can align its direction with the direction required by the target waypoint.

Some control constraints are also applied. In particular, we limit the maximum linear velocity to 0.2 [m/s] and the maximum angular velocity to 0.4 [rad/s]. Moreover, we consider the robot to be close to the target waypoint when the Euclidean distance from the target waypoint is less than 0.05 [m].

3.1 Implementation

The code for the implementation of the controller is provided in the Listing 1.

```
1 function [v, w] = controlProportional(pos_x, pos_y, pos_yaw, target_x, target_y,
2     target_yaw)
3 % This controller align the robot with the target waypoint and guide it to
4 % reach it. Once within the threshold zone, correct the yaw of the robot.
5 % It uses simple Proportional controller for both actions.
6
7 bound = @(value, limit) max(min(value, limit), -limit);
8
9 % Gains
10 K_v = 1.2;
11 K_w = 1.5;
12 max_linear_speed = 0.2;
13 max_angular_speed = 0.4;
14
15 dx = target_x - pos_x;
16 dy = target_y - pos_y;
17 dist_to_target = hypot(dx, dy);
18 angle_to_target = atan2(dy, dx);
19
20 % Control logic
21 if dist_to_target > 0.05
22     v = K_v * dist_to_target;
23     w = K_w * wrapToPi(angle_to_target - pos_yaw);
24 else
25     v = 0;
26     w = K_w * wrapToPi(target_yaw - pos_yaw);
27 end
28
29 v = bound(v, max_linear_speed);
30 w = bound(w, max_angular_speed);
31
```

Listing 1: Code for the implementation of the proportional controller.

3.2 Results

Given the above implementation of the proportional controller, we can now analyze the results obtained during the simulation.

The results of the obtained trajectory are shown in Figure 4, where we report the waypoints highlighted in red and the trajectory of the robot drawn in black.

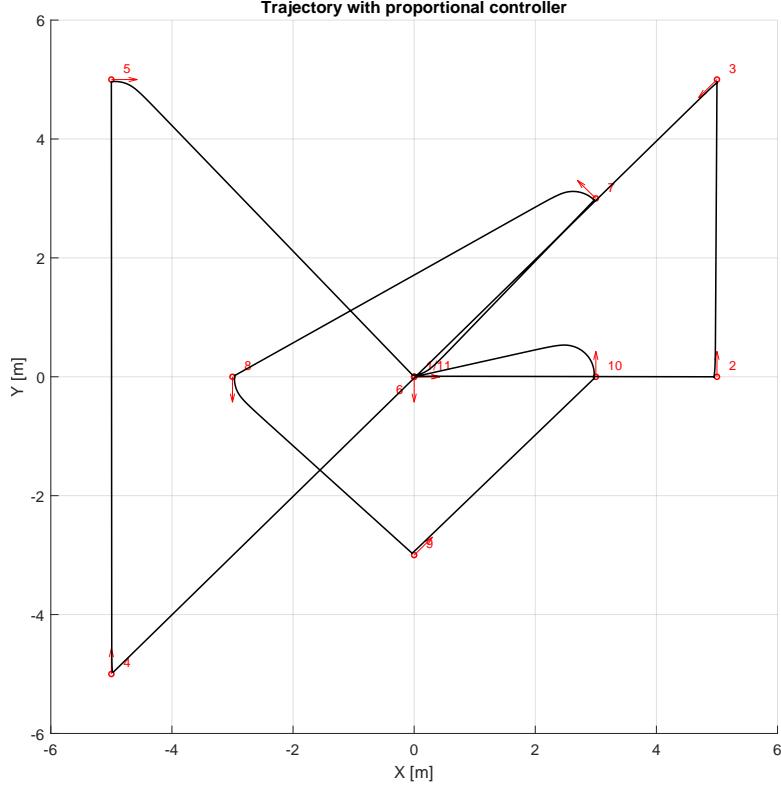


Figure 4: Trajectory of the robot during the simulation with the proportional controller.

One can easily observe that all the waypoints have been reached successfully.

When it comes to the velocity profiles, both in linear and angular dimension, it's easy to visualize the constraints applied to the controller. The linear velocity is limited to 0.2 [m/s], while the angular velocity is limited to 0.4 [rad/s]. This results in a trapezoidal profile for the velocities, as shown in Figure 5.

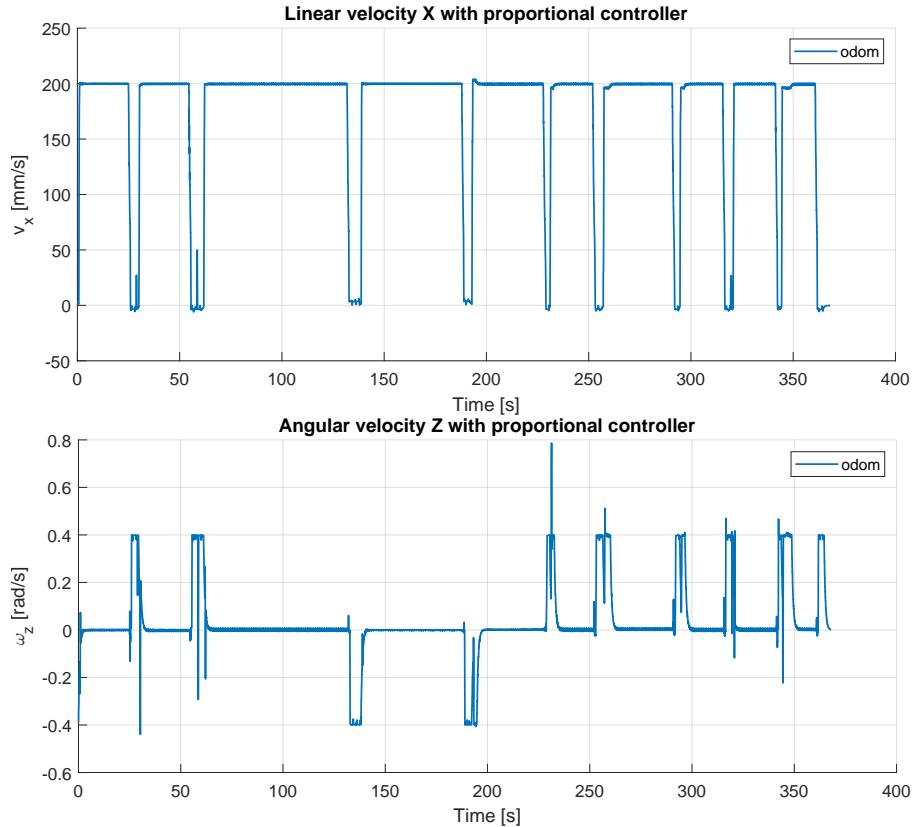


Figure 5: Velocity profiles of the robot during the simulation with the proportional controller.

One can also observe some severe (but luckily not too frequent) peaks in the velocities profiles that violate the constraints imposed. Similarly to the behavior observer in previous assignment, this fact find its explanation in poor performances of the machine on which the simulation is run. The simulation is run on a virtual machine with limited computational power, which results in some momentary oscillations in the robot's dynamics which are well captured by its telemetries. We think that a switch to a native environment of Ubuntu 20.04 on a Linux kernel machine would solve this issue.

4 Lyapunov controller

In this section, we proceed with the implementation of a Lyapunov-based controller for the already defined system.

Lyapunov control theory is a powerful tool for designing controllers that proofs the stability of non-linear systems. The main idea is to find a Lyapunov function associated with the system and prove its convergence to the desired equilibrium point. Notice that global asymptotic stability is guaranteed only if the Lyapunov function is positive definite and radially unbounded.

Once the Lyapunov function is defined and its properties are verified, we can derive a control law that ensures the system's stability. Again, many methods from non-linear control theory exists such as the backstepping method, the feedback linearization method, or again passivation-based methods.

A final important remark here is that the control law derived from Lyapunov theory is not unique and generally speaking is also much more complex than the one derived from traditional linear control theory. Even if stability is guaranteed, the performance of the system may not be optimal and a proper tuning of the controller is not always straightforward.

4.1 Lyapunov control law

Based on the work of [1], we can derive a Lyapunov control law for a generic unicycle-like system. Given that the low level firmware of the Turtlebot3 robot allows us to send control inputs in the form of linear and angular velocities, we can consider the turtlebot as a unicycle-like system and apply the Lyapunov control law to it.

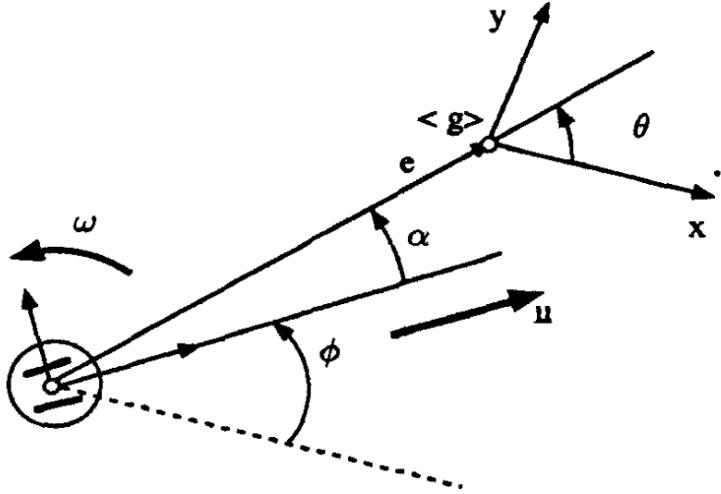


Figure 6: Schematic representation of the unicycle-like system and the reference frame.

Based on the notation used in Figure 6, we can define the following Lyapunov function:

$$v = k_1 \cos(\alpha)e \quad (1)$$

$$w = k_2 \alpha + k_1 \frac{\sin(\alpha) \cos(\alpha)}{\alpha} (\alpha + k_3 \theta) \quad (2)$$

Where k_1 , k_2 , and k_3 are positive constants, e is the Euclidean distance from the target waypoint, α is the angle between the robot heading and the target waypoint, and θ is the angle between the vector connecting the robot to the target waypoint and the waypoint direction (see Figure 6).

It turns out that with a proper remapping of the control inputs, we can obtain a Lyapunov function that is positive definite and radially unbounded, hence ensuring the global asymptotic stability of the system.

4.2 Implementation

The code for the implementation of the Lyapunov controller is provided in the Listing 2.

```

1 function [v, w] = controlLyapunov(pos_x, pos_y, pos_yaw, target_x, target_y, target_yaw)
2 % This controller implements a Lyapunov-based control law for the unicycle-like
3 % system. It uses a Lyapunov function to derive the control inputs for the robot.
4 % The controller is designed to ensure global asymptotic stability of the system.
5 % Inspired by: Aicardi et al. (1995)
6
7 bound = @(value, limit) max(min(value, limit), -limit);
8
9 % Limits
10 max_linear_speed = 0.8;
11 max_angular_speed = 1.6;
12
13 % Gains
14 K_rho = 1.0;
15 K_alpha = 1.5;
16 K_delta = 1.0;
17
18 dx = target_x - pos_x;
19 dy = target_y - pos_y;
20 rho = hypot(dx, dy);
21 alpha = wrapToPi(atan2(dy, dx) - pos_yaw);
22 theta = wrapToPi(atan2(dy, dx) - target_yaw);
23
24 % Control law (from Aicardi paper)
25 v = K_rho * rho * cos(alpha);
26 w = K_alpha * alpha + K_rho * (sin(alpha) * cos(alpha)) + K_rho * (sin(alpha) * cos(
    alpha)) * K_delta * theta / max(abs(alpha), 1e-3) * sign(alpha);
27
28 v = bound(v, max_linear_speed);
29 w = bound(w, max_angular_speed);
30
31 end
```

Listing 2: Code for the implementation of the Lyapunov controller.

4.3 Results

The results of the obtained trajectory are shown in Figure 7, where we can see the waypoints highlighted in red and the trajectory of the robot drawn in black.

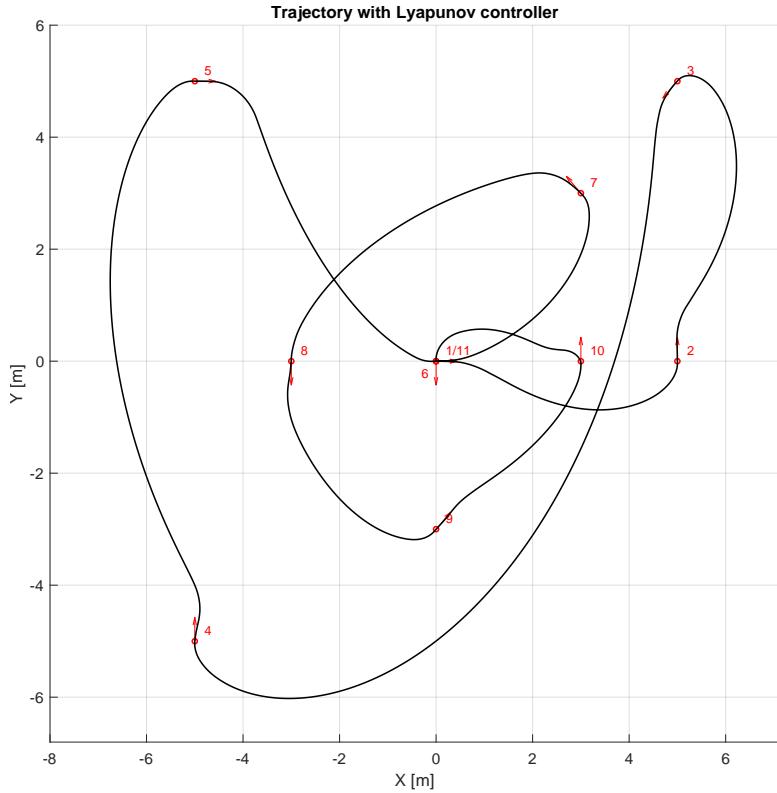


Figure 7: Trajectory of the robot during the simulation with the Lyapunov controller.

With respect to the proportional controller, we can see that the Lyapunov controller is able to smoothly go through the waypoints and to reach them avoiding stop-and-go behavior. The if-else logic is not needed any more given that the control law is able to smoothly connect the waypoints weighting the timing and amplitude for steering and velocity commands based on the 3 parameters k_1 , k_2 , and k_3 . One could also modify the parameters to obtain a more aggressive or conservative behavior of the robot, with consequent changes in the trajectory.

When it comes to the velocities profile, we observe that linearity is lost and non-linear behavior is present.

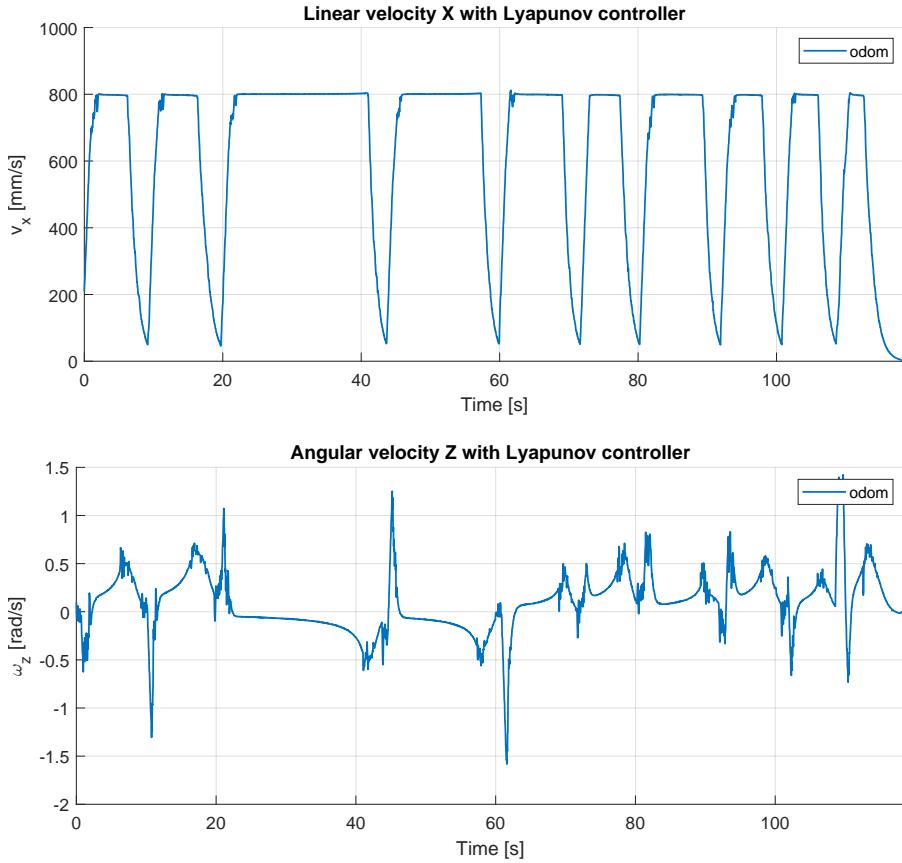


Figure 8: Velocity profiles of the robot during the simulation with the Lyapunov controller.

Notice that with respect to the simulation run with the proportional controller, the linear velocity is now not limited to 0.2[m/s], but it can reach up to 0.8[m/s], and similarly the angular velocity can reach up to 1.6[rad/s]. The author has decided to increase the limits of the velocities in order to allow the robot to reach the waypoints in a shorter time. Nevertheless, the no-slip condition generally required for the unicycle-like system has been checked and verified even with the presence of stronger inertial forces due to the higher velocities.

5 Conclusions

In this short report, we have presented the results of the second assignment of the course on Autonomous Vehicles. We have shown how to implement a feedback control strategy for a mobile robot, specifically the **Turtlebot3** robot, in a simulated environment using **ROS1** and **Gazebo**.

Even if not extensively discussed in this report, the **Simulink** system has allowed to easily load waypoints from different sources (i.e. workspace or ROS topic) and no differences in the results have been observed.

Most importantly, we have shown how to implement a simple proportional controller and a Lyapunov-based controller for the robot. While both have shown to be effective in guiding the robot to the target waypoints, the Lyapunov controller has demonstrated a native smoothness in the trajectory, avoiding stop-and-go behavior. Some unwanted spikes in the telemetry of the robot have been observed during the simulation, which could be due to the fact that the simulation was run on a **Windows** environment connected to a **WSL** (Windows Subsystem for Linux). Further investigations and testing could be done by switching to a native **Linux** environment, which could provide better performance and more accurate results.

References

- [1] M. Aicardi, G. Casalino, A. Bicchi, and A. Balestrino. Closed loop steering of unicycle like vehicles via lyapunov techniques. *IEEE Robotics & Automation Magazine*, 2(1):27–35, 1995.
- [2] Stanford Artificial Intelligence Laboratory et al. Robotic operating system.

Autonomous Vehicles
Assignment IV: Algorithms for grid-based motion planning

Tommaso Bocchietti 10740309

A.Y. 2024/25



POLITECNICO
MILANO 1863

Contents

| | | |
|----------|------------------------------------|-----------|
| 1 | Introduction | 3 |
| 2 | Grid-based planning | 3 |
| 2.1 | Requests | 3 |
| 2.2 | Graph search algorithms | 3 |
| 2.2.1 | Dijkstra's algorithm | 4 |
| 2.2.2 | A* algorithm | 6 |
| 3 | Algorithm Testing | 7 |
| 3.1 | Map 1 | 8 |
| 3.2 | Map 2 | 10 |
| 3.3 | Map 3 | 11 |
| 3.4 | Map 4 | 12 |
| 4 | Conclusions | 13 |
| A | MATLAB Graph implementation | 14 |

List of Figures

| | | |
|---|--|----|
| 1 | Generic graph search algorithm flowchart | 4 |
| 2 | Test maps used for the algorithm testing. | 8 |
| 3 | Map 1: Graph and path found by the algorithms. | 9 |
| 4 | Map 2: Graph and path found by the algorithms. | 10 |
| 5 | Map 3: Graph and path found by the algorithms. | 11 |
| 6 | Map 4: Graph and path found by the algorithms. | 12 |
| 7 | Profiler output for the graph implementation. | 16 |

1 Introduction

The aim of this work is to gain insight into the implementation of some of the most basic grid-based planning algorithms, specifically the **Dijkstra** and **A*** algorithms.

The following sections will provide a detailed description of the requests associated with this assignment, the approach taken to fulfill them, and the discussion of the results obtained.

Tools As for the tools used, **MATLAB** is employed as the main platform for implementing the algorithms and performing the analysis of the data collected during the simulation.

2 Grid-based planning

In the field of robotics, a variety of algorithms are available for path planning, each with its own strengths and weaknesses. One of the most ancient and well known class of algorithms is the so called *Grid-based*, which is based on the discretization of the environment into a grid of cells. Once the environment is represented as a grid, the problem of path planning is reduced to a graph search problem, where the cells of the grid are the nodes of the graph and the edges are the connections between adjacent cells.

2.1 Requests

Among the graph search algorithms, two of the most widely known are the **Dijkstra** and its extension **A*** algorithm. The goal of this assignment can be summarized as follows:

- Implement both the **Dijkstra** and **A*** algorithms;
- Allow the possibility of orthogonal and diagonal movements;
- Test the algorithms on a set of predefined maps and compare the performances.

2.2 Graph search algorithms

In general, given a graph $G = (V, E)$, where V is the set of vertices and E is the set of edges, a graph search algorithm is used to find a path from a starting vertex $s \in V$ to a goal vertex $g \in V$. The main idea is that the algorithm explores the graph by expanding nodes and evaluating their neighbors until it finds the goal node or exhausts all possibilities.

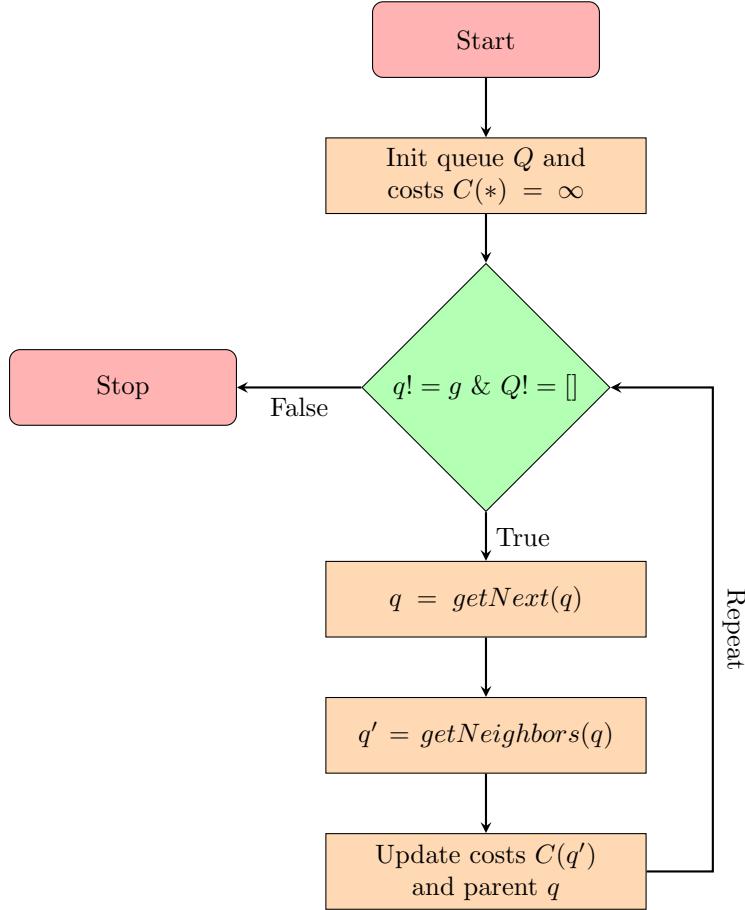


Figure 1: Generic graph search algorithm flowchart

The flowchart in Figure 1 summarizes the main steps of a generic graph search algorithm. The algorithm starts by initializing the open list of nodes and the costs of all nodes to infinity. Then, it enters a loop where based on the previous node visited, its neighbors are evaluated and added to the open list. For each neighbor, the algorithm updates the cost if lower than what already stored and in case updates also the parent node. The loop continues until the open list is empty, or the goal node is reached.

Grid-based algorithms have been proven to be complete and optimal, meaning that they will always find the shortest path if one exists. However, they can be computationally expensive, especially in large environments. There exist many variations built on top of the core algorithm presented in Figure 1. But, what really differentiates them, is the criteria used to select the next node to analyze from the open list. Some of the most common approaches are:

- **Depth-first search (DFS):** explores as far as possible along each branch before backtracking. It uses a stack to keep track of the nodes to be explored. This approach can easily get stuck in loops.
- **Breadth-first search (BFS):** explores all the neighbors of a node before moving on to the next level. It uses a queue to keep track of the nodes to be explored.
- **Best-first (a.k.a. Dijkstra):** selects the node with the lowest cost from the start. This approach is optimal and complete, but can be slow in large environments.
- **A^{*}:** an extension of Dijkstra's algorithm that uses a heuristic to estimate the cost from the current node to the goal node. This approach is also optimal and complete, and is often faster than Dijkstra's algorithm in practice, requiring a lower number of nodes to be explored.

In the following sections, we will focus on the Dijkstra and A^{*} algorithms, explaining their mechanisms in choosing the next node to be explored and how they can be implemented to solve the path planning problem in a grid-based environment.

2.2.1 Dijkstra's algorithm

As already mentioned, Dijkstra's algorithm leverages the core structure of the graph search algorithm presented in Figure 1, but with a specific strategy for selecting the next node to explore.

In particular, at each iteration, the algorithm selects the node among the nodes in the open list that has the lowest cost from the start node. The idea is so to select $q' \in Q$ such that:

$$q' = \min_{q \in Q} C(q) \quad (1)$$

Where $C(q)$ is the cost of reaching node q from the start node s . The cost is calculated as the sum of the cost of the previous node q and the cost of moving from q to q' .

From an implementation point of view, the algorithm can be easily implemented using a priority queue to store the nodes in the open list. The priority queue allows for efficient retrieval of the node with the lowest cost, which is crucial for the performance of the algorithm. Listing 1 shows a possible implementation of the Dijkstra's algorithm. The algorithm takes as input the graph G , the start node s , and the goal node g .

```

1 function [exist, distances, parents] = dijkstra(G, node_initial, node_final)
2 % DIJKSTRA Finds the shortest path between two nodes in a graph
3 % using Dijkstra's algorithm.
4
5 N = length(G.nodes);
6 distances = inf(N, 1);
7 parents = NaN(N, 1);
8 visited = false(N, 1);
9
10 % Initialization
11 ID_current = node_initial.ID;
12 distances(ID_current) = 0;
13
14 while (ID_current ~= node_final.ID && ~all(visited))
15
16     unvisited = find(~visited);
17     [~, ID_current] = min(distances(unvisited));
18     ID_current = unvisited(ID_current);
19
20     visited(ID_current) = true;
21
22     % For each neighbor not yet visited
23     IDs_neighbors = G.adjacents{ID_current};
24     IDs_neighbors = IDs_neighbors(~visited(IDs_neighbors));
25     for ID_neighbor = IDs_neighbors'
26
27         distance = ...
28             distances(ID_current) + ...
29             G.getEdgeByConnection(ID_current, ID_neighbor).weight;
30
31         if distance < distances(ID_neighbor)
32             distances(ID_neighbor) = distance;
33             parents(ID_neighbor) = ID_current;
34         end
35
36     end
37
38 end
39
40 exist = ~isinf(distances(node_final.ID));
41
42 end

```

Listing 1: Code for the implementation of Dijkstra's algorithm.

The algorithm starts by initializing the distances of all nodes to infinity, except for the start node, which is set to zero. Then, it enters a loop where it selects the node with the lowest cost from the open list and explores its neighbors. For each neighbor, it updates the cost if the new cost is lower than the previously stored cost and updates the parent node accordingly. The loop continues until the goal node is reached, or all nodes have been visited. Finally, the algorithm returns the distances and parents of each node, which can be used to reconstruct the shortest path from the start node to the goal node.

2.2.2 A* algorithm

A* algorithm is an extension of Dijkstra's algorithm that uses a heuristic to estimate the cost from the current node to the goal node. The heuristic is a function that provides an estimate of the cost to reach the goal from a given node, and it is used to guide the search towards the goal more efficiently.

In particular, the A* algorithm selects the next node to explore based on the sum of the cost from the start node to the next node and the estimated cost from the next node to the goal node. The selection is done as follows:

$$q' = \min \arg_{q \in Q} (C(q) + h(q)) \quad (2)$$

Where $h(q)$ is the heuristic function that estimates the cost from node q to the goal node g . The heuristic function can be any function that provides an estimate of the cost, but it is important that it is admissible, meaning that it never overestimates the true cost to reach the goal. For the sake of simplicity, and in particular leveraging the meaningful geometrical interpretation of the grid-based environment, we will use the Euclidean distance as the heuristic function. The A* algorithm can be implemented similarly to Dijkstra's algorithm, with the addition of the heuristic function. Listing 2 shows a possible implementation of the A* algorithm. The algorithm takes in the same input as the Dijkstra's one.

```

1 function [exist, distances, parents] = astar(G, node_initial, node_final)
2 % ASTAR Finds the shortest path between two nodes in a graph
3 % using the A* algorithm.
4
5 N = length(G.nodes);
6 distances = inf(N, 1);
7 costs      = inf(N, 1);
8 parents    = NaN(N, 1);
9 visited    = false(N, 1);
10
11 % Initialization
12 ID_current = node_initial.ID;
13 distances(ID_current) = 0;
14 costs(ID_current) = Node.euclideanDistance(node_initial.state, node_final.state);
15
16 while (ID_current ~= node_final.ID && ~all(visited))
17
18     unvisited = find(~visited);
19     [~, ID_current] = min(costs(unvisited));
20     ID_current = unvisited(ID_current);
21
22     visited(ID_current) = true;
23
24     % For each neighbor not yet visited
25     IDs_neighbors = G.adjacents{ID_current};
26     IDs_neighbors = IDs_neighbors(~visited(IDs_neighbors));
27     for ID_neighbor = IDs_neighbors'
28
29         distance = ...
30             distances(ID_current) + ...
31             G.getEdgeByConnection(ID_current, ID_neighbor).weight;
32
33         if distance < distances(ID_neighbor)
34
35             distances(ID_neighbor) = distance;
36             parents(ID_neighbor) = ID_current;
37
38             node_neighbor = G.getNodeByID(ID_neighbor);
39             costs(ID_neighbor) = distance + Node.euclideanDistance(node_neighbor.state,
node_final.state);
40
41     end
42

```

```

43     end
44
45 end
46
47 exist = ~isinf(distances(node_final.ID));
48
49 end

```

Listing 2: Code for the implementation of A* algorithm.

One can clearly notice that the structure of the A* algorithm is basically the same as the Dijkstra's algorithm, except for the addition of the heuristic function in the cost calculation. Notice also that the heuristic function is only used to update the cost of the nodes in the open list, and it is not used to update the distances of the nodes from the start node.

Given that now the algorithms is somehow guided, one can expect that the A* algorithm will be faster than the Dijkstra one, as fewer nodes will be explored.

3 Algorithm Testing

Now that the algorithm has been implemented, we can proceed presenting at first the test environment that have been used to test them and then the results of the tests.

Figures 2 show the four image maps used to test the algorithm. Even if the images here below are of size 300x300, we preferred to scale down to 30x30 given that each pixel is then converted into a node of the graph and so to avoid huge graphs and long computation times. Notice also that while the white areas are free space, the one marked in black or gray are to be considered as obstacles. Finally, the green dot is the starting point of the vehicle, while the red one is the goal.

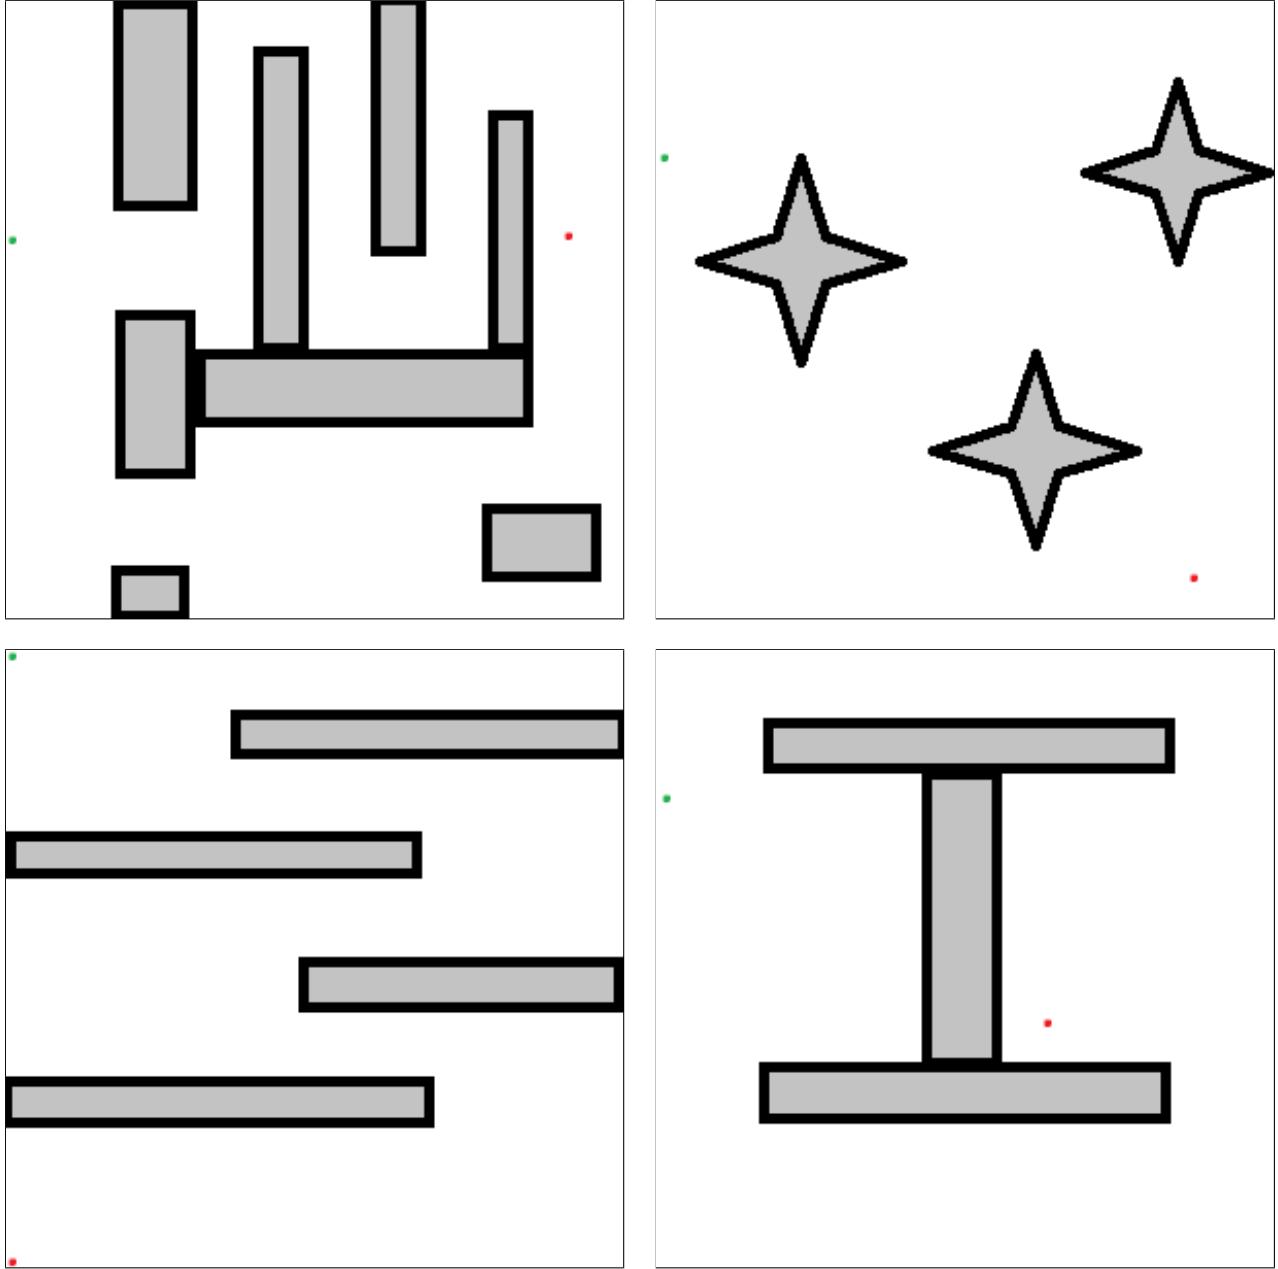


Figure 2: Test maps used for the algorithm testing.

In the following, results obtained running both the A* and the Dijkstra algorithms on each of the four maps are presented.

The plotting of both the (geometrical) connected graph and the best path found is also shown. Notice that each test case has been run allowing or disabling the diagonal movement of the vehicle, which can be achieved during the graph generation phase modifying the connection of the nodes.

Tables presenting the main metrics of the tests are also shown. The metrics are the following:

- **Elapsed time:** here only the time needed to run the searching algorithm is considered (i.e. we are excluding both the generation of the graph and the plotting of the results);
- **Nodes analyzed:** number of nodes analyzed (visited) before reaching the goal;
- **Path length:** length of the path found by the algorithm. This also correspond to the path cost, given that the cost of each edge is equal to its geometrical length.

3.1 Map 1

Results obtained running the A* and Dijkstra algorithms on Map 1 are shown in Figures 3 and Table 1. About the figures, the upper row shows the graph generated running the Dijkstra algorithm while the lower one shows

the graph generated running the A* algorithm. The left column shows the graph generated allowing only orthogonal movements, while the right one shows the graph generated allowing also diagonal movements. The path found by the algorithm is shown in red, while the graph is shown in black.

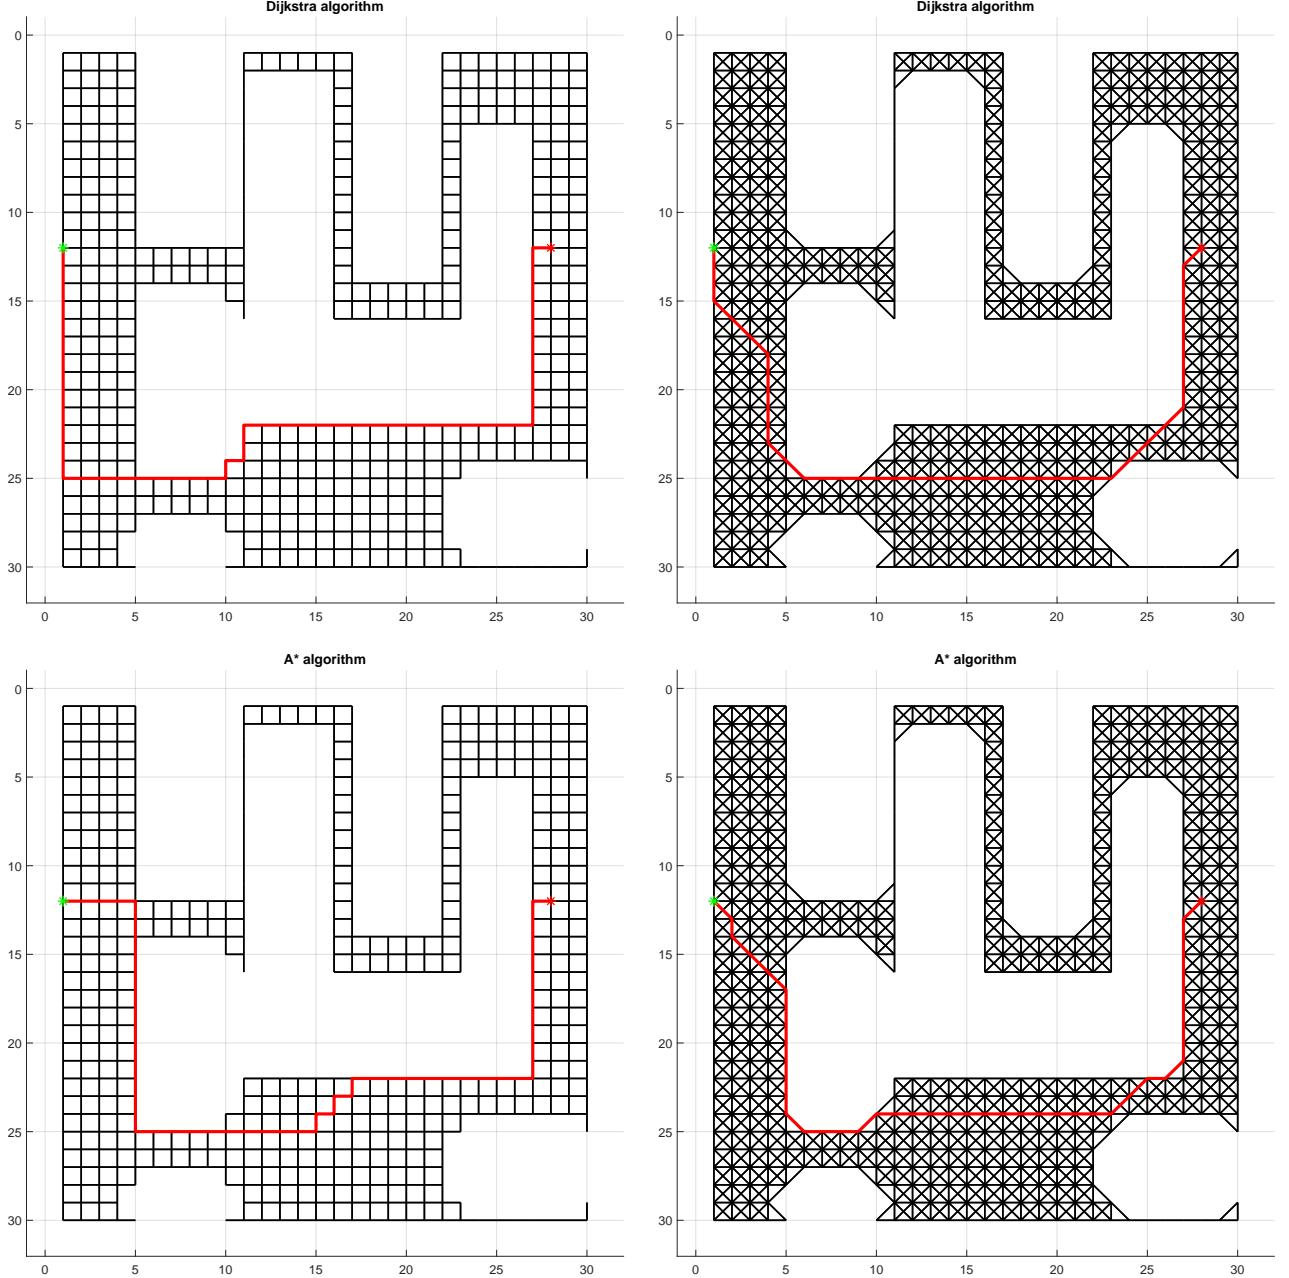


Figure 3: Map 1: Graph and path found by the algorithms.

| Algorithm | Diagonal | Elapsed time (ms) | Nodes analyzed | Path length (m) |
|-----------|----------|-------------------|----------------|-----------------|
| Dijkstra | No | 624 | 456 | 53 |
| | Yes | 1592 | 453 | 47 |
| A* | No | 339 | 378 | 53 |
| | Yes | 1227 | 339 | 47 |

Table 1: Map 1: Results of the tests.

From the results presented in Table 1 and in Figure 3, we can already notice that while the cost of the path found by the two algorithms is equivalent, the number of nodes considered by the A* algorithm is significantly lower than the one considered by the Dijkstra algorithm. This behavior is expected due to the use of the heuristic function to guide the search of the A* algorithm. The advantage in terms of computational time is also evident, with the A* algorithm outperforming the Dijkstra algorithm.

3.2 Map 2

Results obtained running the A* and Dijkstra algorithms on Map 2 are shown in Figures 4 and Table 2. The same layout of the previous map is used for the figures and the same metrics are used for the table.

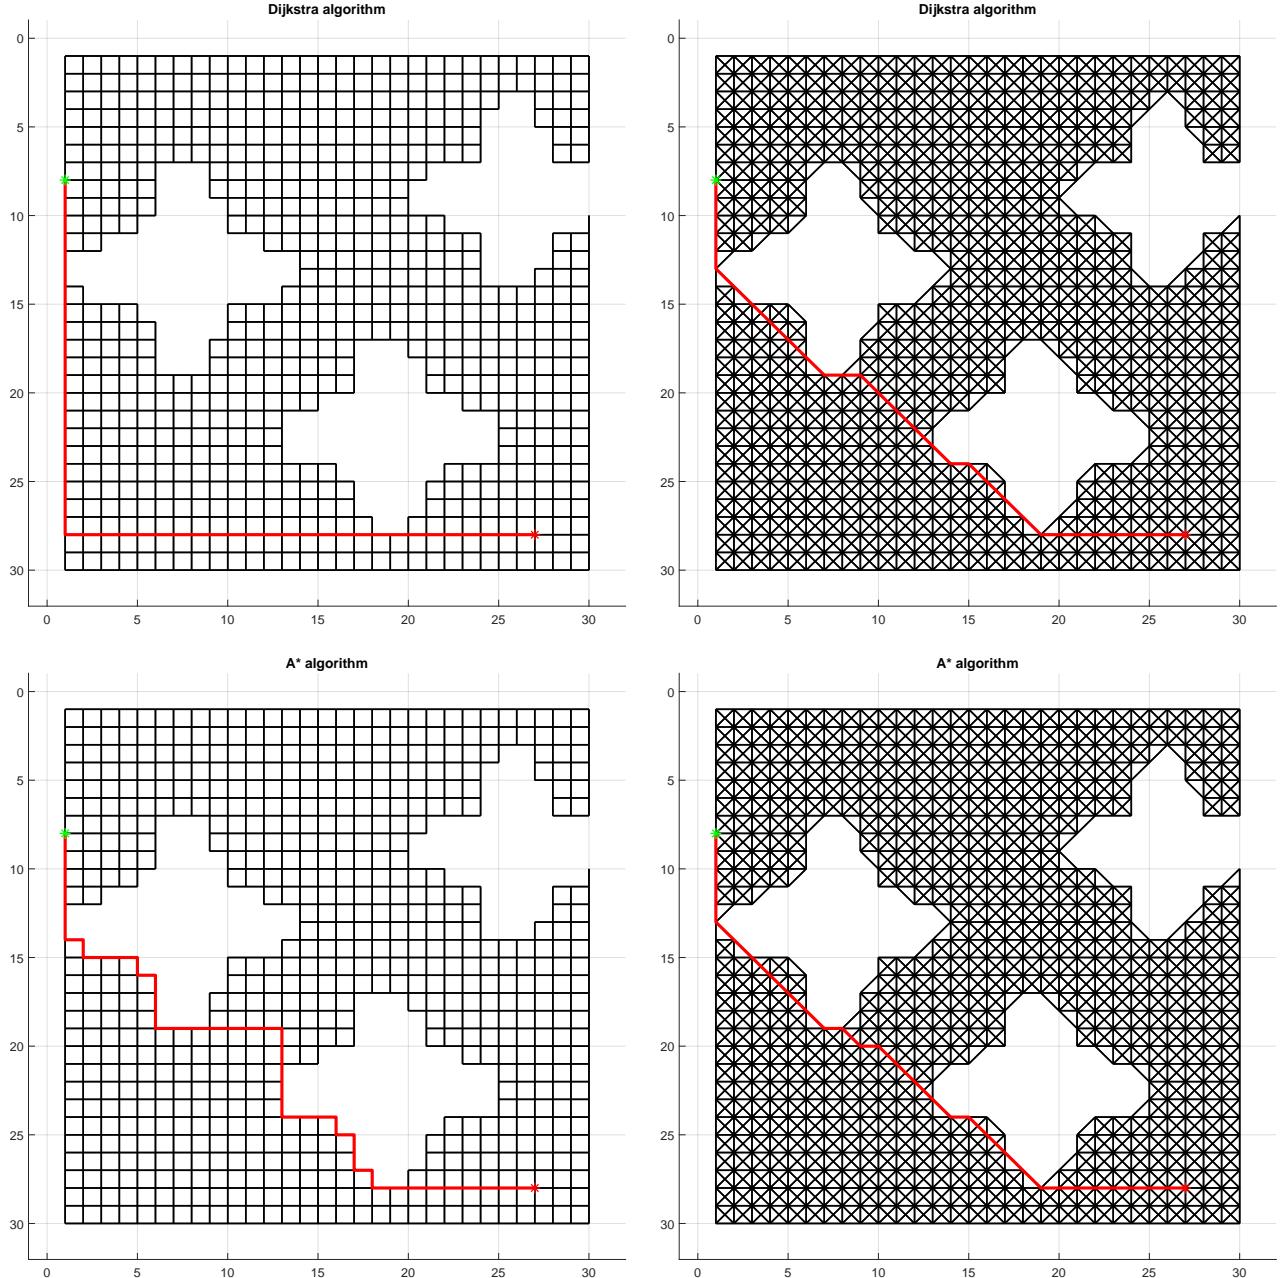


Figure 4: Map 2: Graph and path found by the algorithms.

| Algorithm | Diagonal | Elapsed time (ms) | Nodes analyzed | Path length (m) |
|-----------|----------|-------------------|----------------|-----------------|
| Dijkstra | No | 1009 | 731 | 46 |
| | Yes | 3963 | 742 | 37 |
| A* | No | 732 | 490 | 46 |
| | Yes | 1434 | 233 | 37 |

Table 2: Map 2: Results of the tests.

Similarly to the previous map, the results presented in Table 2 and in Figure 4 show that while the cost of the path found by the two algorithms is equivalent, the number of nodes considered by the A* algorithm is significantly lower than the one considered by the Dijkstra algorithm. Similar conclusions can be drawn in terms of computational time and performance of the two algorithms.

As an additional note, in this case we can also appreciate that the number of analyzed nodes by A* is significantly lower than the one analyzed by Dijkstra, particularly in case of diagonal movements. This is, however, expected. Looking at the final path, we observe that it is almost equivalent to the diagonal connecting ideally the start and goal points. This means that the heuristic function used by A* is even more effective in this case, given that very few obstacles are present along the optimal path designed by the heuristic function itself and so fewer dead ends are encountered. In other words, if the optimal path turns out to be a straight line connecting the start and goal points, the A* algorithm would analyze a number of nodes equal to the number of nodes along the straight line, while the Dijkstra algorithm would still analyze a much larger number of nodes given that no information is used to guide the search.

3.3 Map 3

Results obtained running the A* and Dijkstra algorithms on Map 3 are shown in Figures 5 and table. The same layout of the previous map is used for the figures and the same metrics are used for the table.

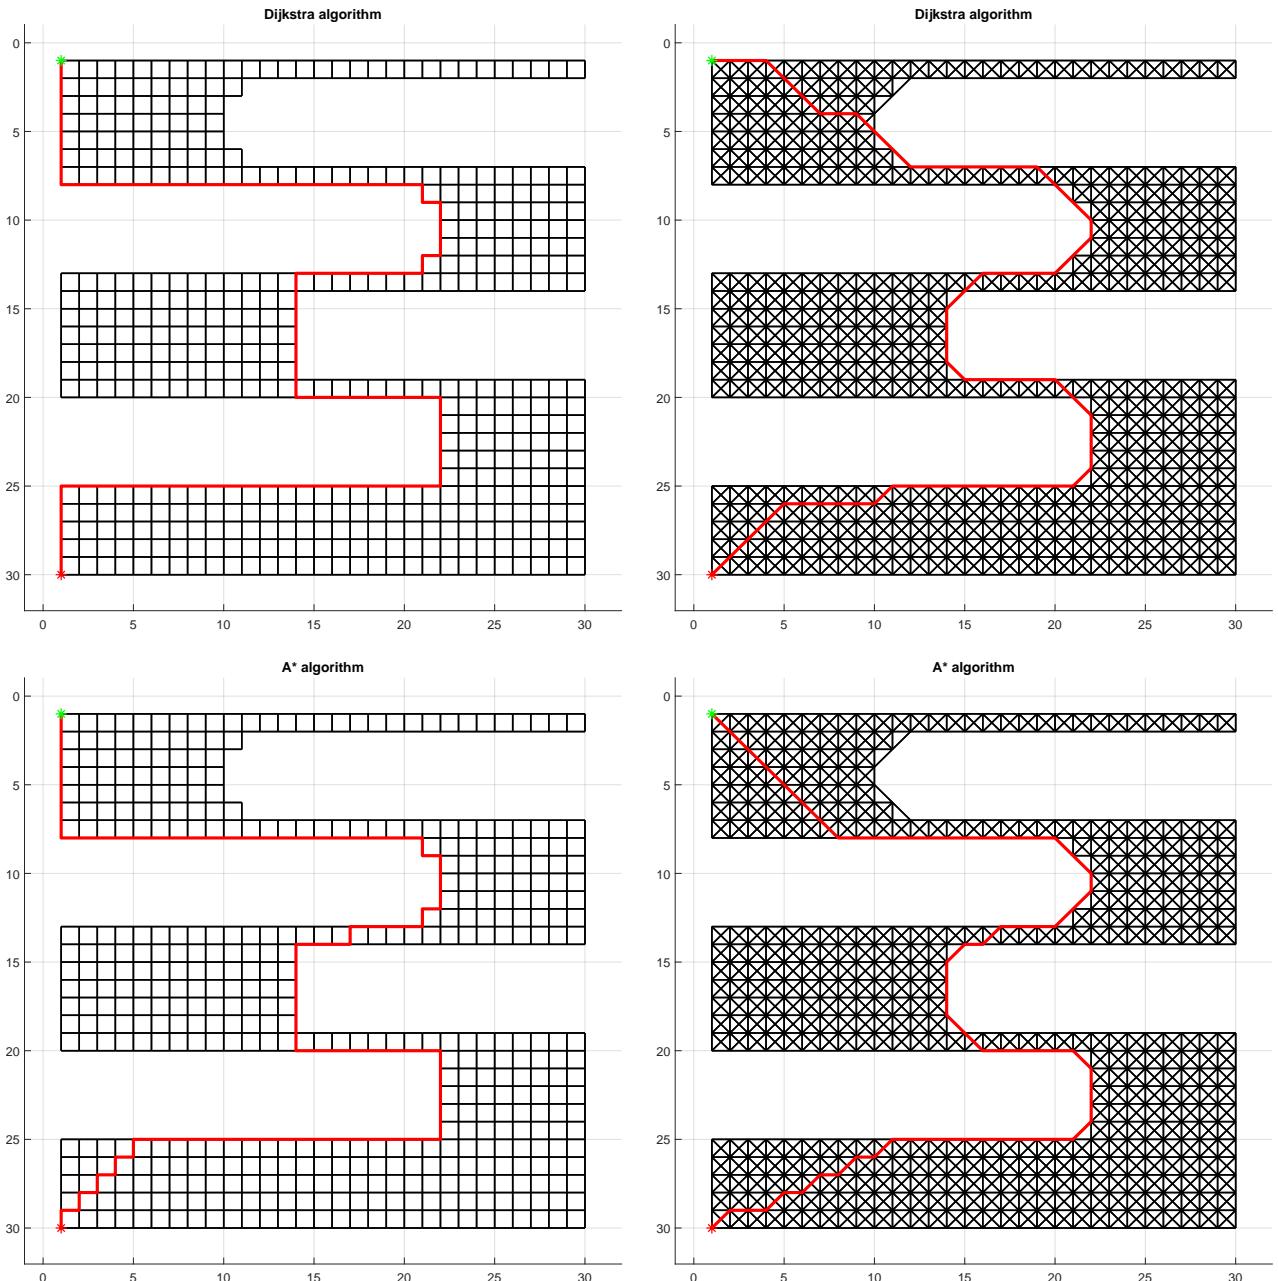


Figure 5: Map 3: Graph and path found by the algorithms.

| Algorithm | Diagonal | Elapsed time (ms) | Nodes analyzed | Path length (m) |
|-----------|----------|-------------------|----------------|-----------------|
| Dijkstra | No | 567 | 592 | 87 |
| | Yes | 2461 | 592 | 74 |
| A* | No | 496 | 497 | 87 |
| | Yes | 1780 | 452 | 74 |

Table 3: Map 3: Results of the tests.

Again, similar considerations as before can be drawn from the results presented in Table 3 and in Figure 5.

3.4 Map 4

Results obtained running the A* and Dijkstra algorithms on Map 4 are shown in Figures 6 and table. The same layout of the previous map is used for the figures and the same metrics are used for the table.

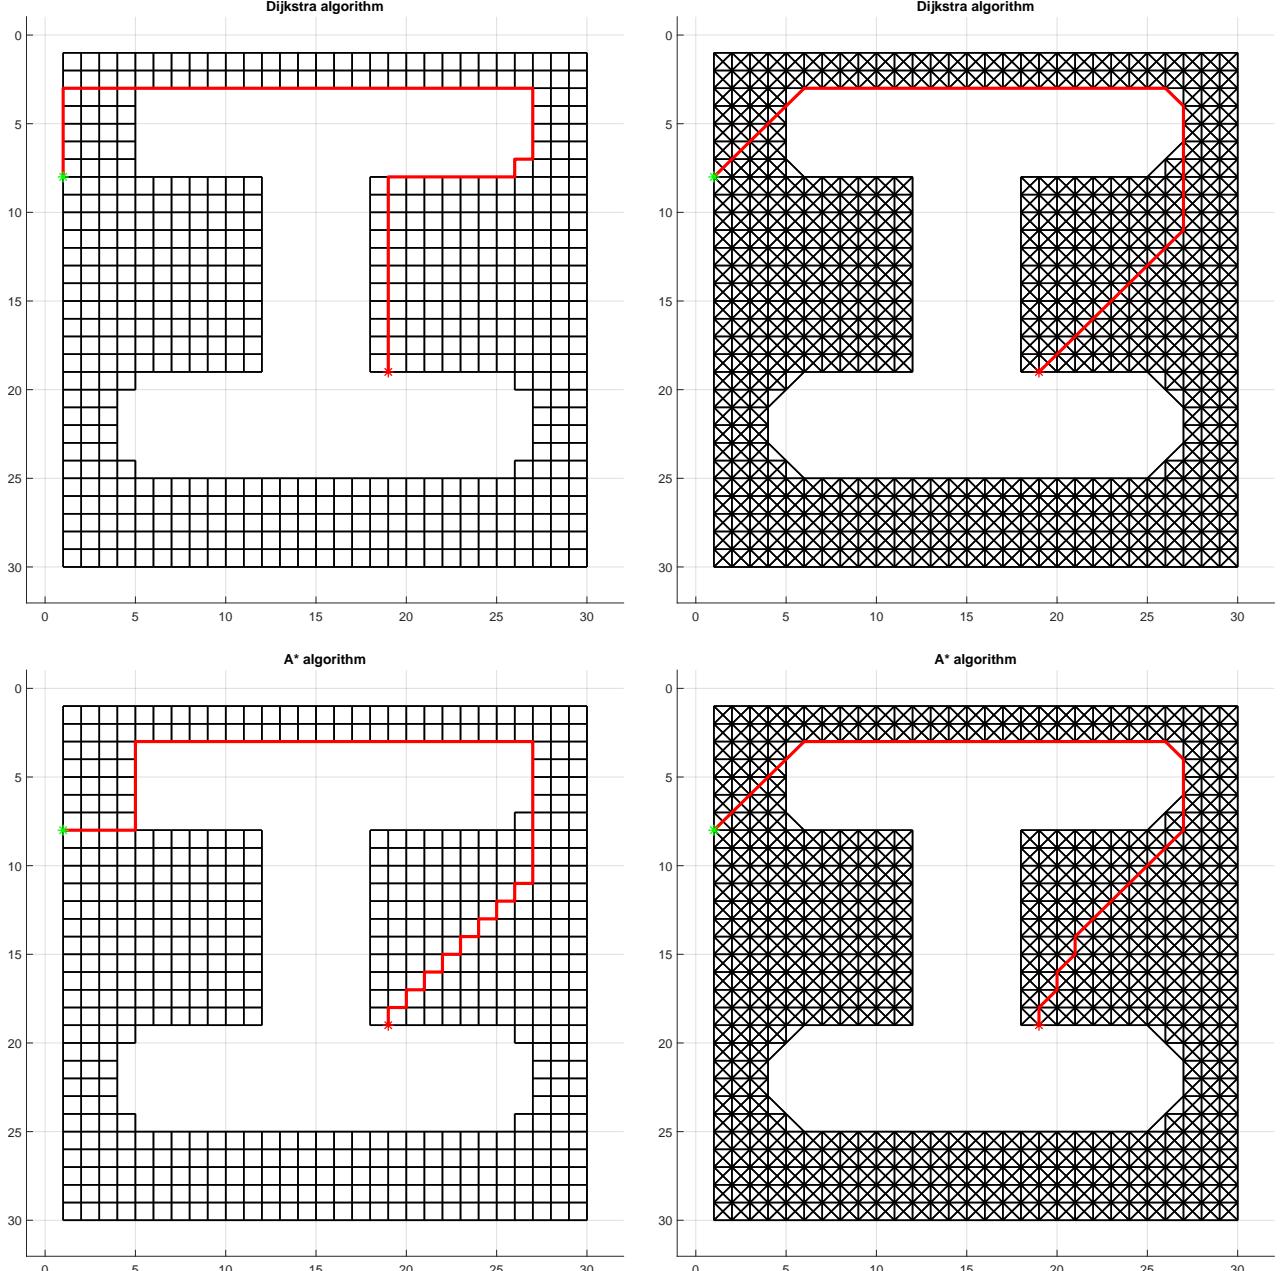


Figure 6: Map 4: Graph and path found by the algorithms.

| Algorithm | Diagonal | Elapsed time (ms) | Nodes analyzed | Path length (m) |
|------------------|-----------------|--------------------------|-----------------------|------------------------|
| Dijkstra | No | 725 | 650 | 55 |
| | Yes | 2799 | 650 | 47 |
| A* | No | 640 | 562 | 55 |
| | Yes | 2060 | 450 | 47 |

Table 4: Map 4: Results of the tests.

Once again, even from the tests on Map 4, similar conclusions as before can be drawn from the results presented in Table 4 and in Figure 6.

4 Conclusions

In this short report, we have presented the results of the fourth assignment of the course on Autonomous Vehicles. We have briefly explained the grid-based algorithm for path planning, and deepened into the implementation of both the Dijkstra and A* algorithms.

These algorithms have been tested on four different scenarios, with different obstacles and connectivity configurations (i.e., allowing or not diagonal movements). The results have shown that the A* algorithm is generally faster than the Dijkstra algorithm as it uses heuristics to guide the search towards the goal allowing for a more efficient exploration of the search space. This is also reflected in the number of nodes explored, which is significantly lower for the A* algorithm compared to the Dijkstra algorithm.

As also proved by the theoretical analysis performed for this grid-based algorithm, both algorithms are complete and optimal, meaning that they will always find the shortest path if one exists. Obtained results have reflected this theoretical analysis, as both algorithms have been able to find the optimal path in all tested scenarios.

One of the main limitations of the grid-based algorithm is that they can be computationally expensive, especially in large or hyperdimensional spaces. A quick solution to this problem is to use a coarser grid, which can reduce the number of nodes and edges in the graph, but at the cost of accuracy and possibly missing the optimal path, or even failing to find a path at all.

Because of the above limitation, the grid-based algorithm is not suitable for all applications, especially in dynamic environments where the obstacles can change over time. In these cases, more advanced algorithms based on combinatorial techniques or sampling-based methods, such as Rapidly-exploring Random Trees (RRT) or Probabilistic Roadmaps (PRM), may be more appropriate.

A MATLAB Graph implementation

During the implementation of the Dijkstra and A* algorithms, a need for a solid and efficient graph representation arose. In order to achieve this, a set of MATLAB classes was created to represent the graph, its nodes and its edges.

In the following, the complete code used for this purpose is shown. One can refer back to the graph search algorithm implementation to understand how these classes are used.

```
1 classdef Graph < handle
2 % GRAPH Class for representing a graph structure with nodes and weighted edges.
3 % Provides methods for adding and retrieving nodes, edges, and adjacency data.
4
5 properties
6     nodes Node
7     edges Edge
8     adjacents cell
9 end
10
11 methods
12     function obj = Graph()
13         obj.nodes = Node.empty();
14         obj.edges = Edge.empty();
15         obj.adjacents = {};
16     end
17
18     function ID = addNode(obj, state)
19         ID = length(obj.nodes) + 1;
20         obj.nodes(ID, 1) = Node(ID, state);
21         obj.adjacents{ID, 1} = [];
22     end
23
24     function ID = addEdge(obj, ID_A, ID_B, weight)
25         ID = length(obj.edges) + 1;
26         obj.edges(ID, 1) = Edge(ID, ID_A, ID_B, weight);
27         obj.adjacents{ID_A, 1}(end+1, 1) = ID_B;
28     end
29
30     function node = getNodeByID(obj, ID)
31         node = obj.nodes(ID);
32     end
33
34     function edge = getEdgeByID(obj, ID)
35         edge = obj.edges(ID);
36     end
37
38
39     function nodes = getNodesByState(obj, state)
40         nodes = Node.empty();
41         for node = obj.nodes'
42             if Node.compareStates(node.state, state)
43                 nodes(end+1) = node;
44             end
45         end
46     end
47
48
49     function edge = getEdgeByConnection(obj, ID_A, ID_B)
50         for edge = obj.edges([obj.edges.ID_A] == ID_A)'
51             if (edge.ID_B == ID_B)
52                 return
53             end
54         end
55     end
```

```

54     end
55 end
56
57
58 function adjacents = getAdjacents(obj, node_ID)
59     adjacents = obj.adjacents{node_ID};
60 end
61 end
62
63 end

```

Listing 3: Graph class used to represent a weighted and directed/undirected graph.

```

1 classdef Edge
2 % EDGE Represents a weighted connection between two nodes in a graph.
3 % Stores IDs of connected nodes and the edge weight.
4
5 properties
6     ID double
7     ID_A double
8     ID_B double
9     weight double
10 end
11
12 methods
13     function obj = Edge(ID, ID_A, ID_B, weight)
14         obj.ID = ID;
15         obj.ID_A = ID_A;
16         obj.ID_B = ID_B;
17         obj.weight = weight;
18     end
19
20     function isConnection = isConnectionTo(obj, ID_A, ID_B)
21         isConnection = obj.ID_A == ID_A & obj.ID_B == ID_B;
22     end
23 end
24
25 end

```

Listing 4: Edge class used to represent a weighted and directed edge in the graph.

```

1 classdef Node
2 % NODE Represents a graph node with a unique ID and a state vector.
3 % Includes utility methods for comparison and distance calculation.
4
5 properties
6     ID double
7     state double
8 end
9
10 methods
11     function obj = Node(ID, state)
12         obj.ID = ID;
13         obj.state = state;
14     end
15
16     function isEqual = eq(node_A, node_B)
17         isEqual = Node.compareStates(node_A.state, node_B.state);
18     end
19 end
20

```

```

21     methods (Static)
22         function dist = euclideanDistance(state_A, state_B)
23             dist = norm(state_A - state_B);
24         end
25
26         function equal = compareStates(state_A, state_B)
27             equal = isequal(state_A, state_B);
28         end
29     end
30
31 end

```

Listing 5: Node class used to represent a node in the graph.

As a final note, a code profiler was used to test the performance of the graph implementation. The main bottleneck was found in the `getEdgeByConnection` method of the `Graph` class, as shown in Figure 7.

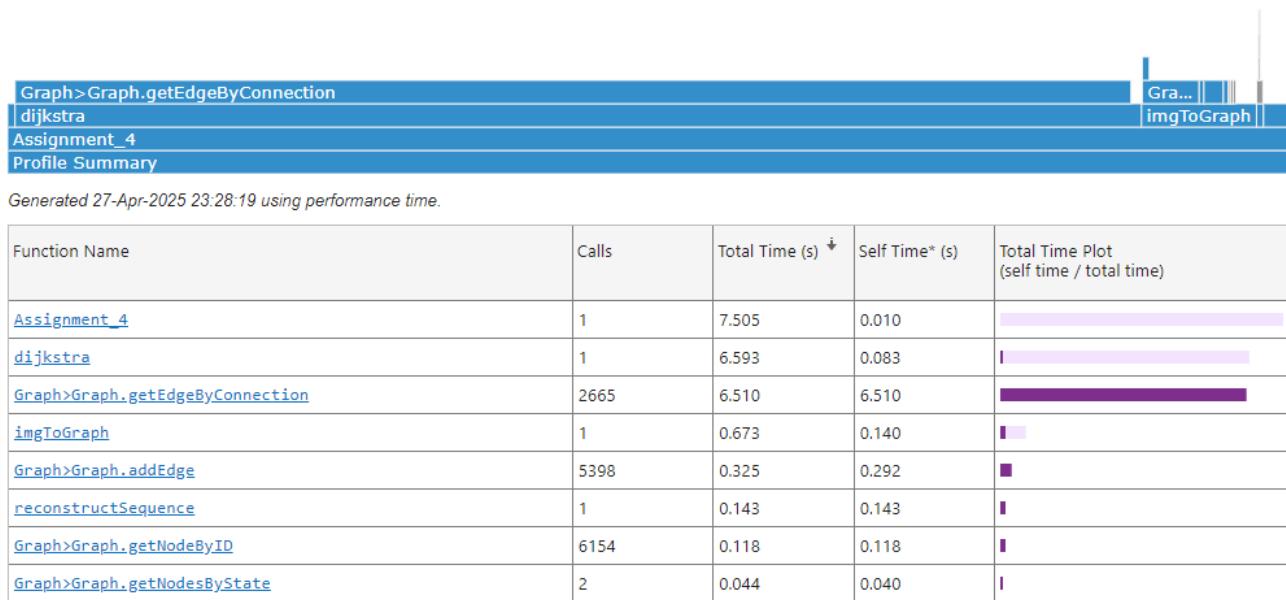


Figure 7: Profiler output for the graph implementation.

The author of this report is aware that the graph implementation can be improved in terms of performance by using perhaps a caching mechanism or a more efficient data structure. However, this was not the main focus of the assignment and further investigation is left as future work.

Autonomous Vehicles
Assignment V: Finite State Machines

Tommaso Bocchietti 10740309

A.Y. 2024/25



POLITECNICO
MILANO 1863

Contents

| | | |
|----------|------------------------------|-----------|
| 1 | Introduction | 3 |
| 2 | Finite State Machines | 3 |
| 2.1 | Turnstile FSM | 3 |
| 3 | Parking Gate System | 4 |
| 3.1 | Requests | 4 |
| 3.2 | System Overview | 4 |
| 3.2.1 | Vehicle FSM | 5 |
| 3.2.2 | Parking Gate FSM | 6 |
| 3.3 | Simulation Results | 9 |
| 4 | Conclusions | 10 |
| A | Appendix | 11 |

List of Figures

| | | |
|----|--|----|
| 1 | Turnstile FSM State Transition Diagram | 3 |
| 2 | Parking Gate System Overview | 4 |
| 3 | Parking Gate System FSM | 5 |
| 4 | Vehicle FSM | 6 |
| 5 | Control Block | 6 |
| 6 | Gate Block | 7 |
| 7 | PhotoCell Block | 8 |
| 8 | Semaphore Block | 9 |
| 9 | Simulation Results | 9 |
| 10 | Simulation Results - Detailed View | 10 |
| 11 | Original Parking Gate System Overview | 11 |

1 Introduction

The aim of this work is to gain insight into the modelling of systems via Finite State Machines (FSMs) and to understand the principles behind their implementation in the context of autonomous vehicles.

At first, a brief overview of the concepts of FSMs comprehensive of simple examples is provided. Then, the focus shifts to the modelling of a simple application, namely a parking gate system, which is designed using FSMs.

Tools As for the tools used, **Simulink** and in particular **Simulink/Stateflows**, is employed as the main tool for implementing the FSM and simulating the system. On the other hand, **MATLAB** is simply used to plot the results of the simulation.

2 Finite State Machines

Finite state machines (FSMs) are a powerful method of modelling systems whose output depends on the entire history of their inputs (as opposed to memoryless systems, where the output depends on the current input only). FSMs are widely used in various fields, including computer science, control systems, and digital circuit design. They can be classified into two main classes: deterministic finite state machines (DFSMs) and non-deterministic finite state machines (NDFSMs).

Mathematically, a finite state machine is defined as a 6-tuple $(S, I, O, \delta, \lambda, s_0)$, where:

- S is a finite set of states.
- I is a finite set of input symbols (input alphabet).
- O is a finite set of output symbols (output alphabet).
- $\delta : S \times I \rightarrow S$ is the state transition function, which maps a state and an input symbol to the next state.
- $\lambda : S \times I \rightarrow O$ is the output function, which maps a state and an input symbol to an output symbol.
- $s_0 \in S$ is the initial state.

The FSM processes a sequence of input symbols, transitioning between states according to the state transition function δ and producing output symbols according to the output function λ . The behavior of an FSM can be represented using state transition diagrams or state transition tables.

2.1 Turnstile FSM

As an example, we consider a turnstile that allows entry when a coin is inserted and locks when a person pushes through. The FSM for this turnstile can be represented by the state transition diagram in Figure 1 [1] and the state transition table in Table 1 [1].

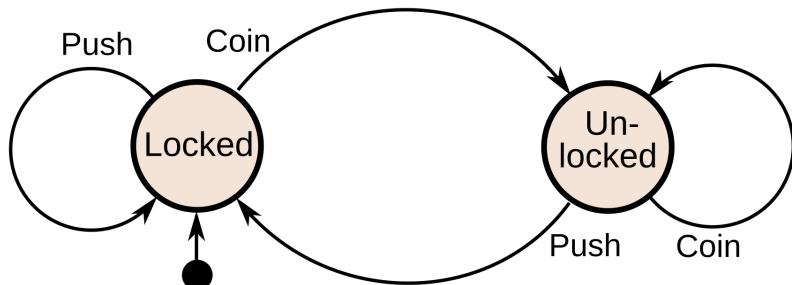


Figure 1: Turnstile FSM State Transition Diagram

| Current State | Input | Next State | Output |
|---------------|--------------|--------------------|---|
| Locked | Coin Push | Unlocked Locked | Unlocks the turnstile so that the customer can push through None |
| Unlocked | Coin Push | Unlocked Locked | None When the customer has pushed through, locks the turnstile |

Table 1: Turnstile FSM Transition Table

The FSM starts in the *Locked* state. When a coin is inserted, it transitions to the *Unlocked* state and allows the customer to push through. If the customer pushes through while in the *Unlocked* state, the FSM transitions back to the *Locked* state. If a coin is inserted while in the *Unlocked* state, the FSM remains in the same state.

3 Parking Gate System

As already mentioned in the previous sections, the focus of this assignment is to model a parking gate system using finite state machines (FSMs).

3.1 Requests

The requests that the parking gate system has to fulfil are the following:

- Analyze the provided *Vehicle FSM* and understand how it works;
- Define the *raise* and *lower* functions of the gate respecting $w_{max} = |4|[\text{°}/\text{s}]$ and $\dot{w}_{max} = |1|[\text{°}/\text{s}^2]$;
- Define the FSM to model the parking gate control system;
- Integrate the *Vehicle FSM* and the parking gate control system FSM;
- Provide results of the simulation of the integrated system;

Notice that despite the suggestion from the professor to start from the already provided *Vechicle FSM*, I decided to start from scratch in order to develop a deepened understanding of the FSMs, modelling also events, slightly more complex logics and composition / decomposition patterns.

3.2 System Overview

The parking gate system is composed of a single Finite State Machine (FSM) that is responsible for modelling at the same time the vehicle and the parking gate system.

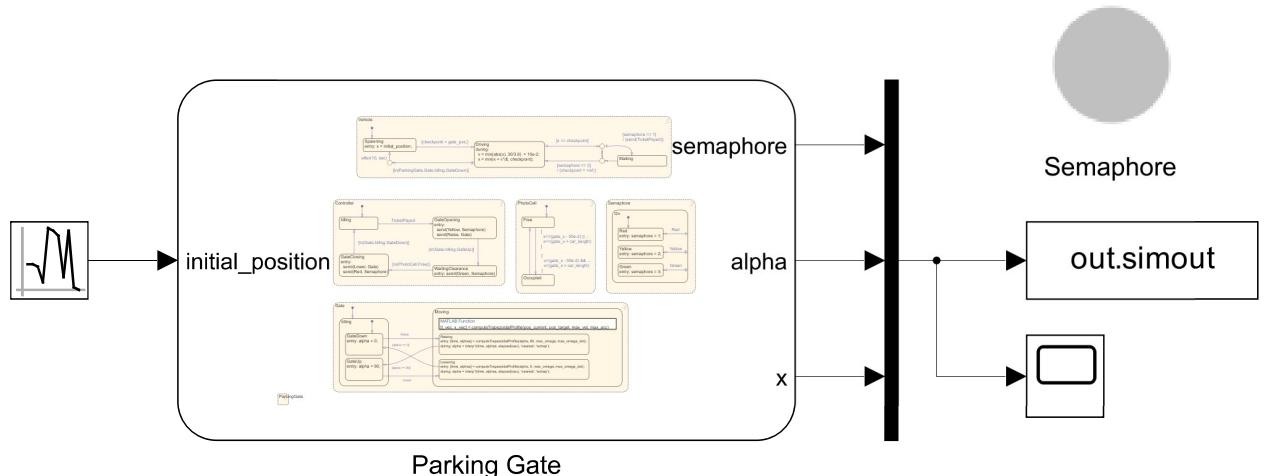


Figure 2: Parking Gate System Overview

Under the hood of the unique FSM, the system is composed of other two parallel FSMs, with precise and well-defined roles and responsibilities. The communication between the two FSMs is done through events and state transitions, which are used to inform each other about the current state of the system and to trigger the appropriate actions.

- The *Vehicle FSM* share the *vehicle_position* (x) with the *Parking Gate FSM* to inform it about the vehicle's position in the parking lot;
- The *Parking Gate FSM* share the *semaphore_state* with the *Vehicle FSM* to inform it about the state of the gate (closed, opening, opened);
- The *Vehicle FSM* also sends the event about the *TicketPayed* to the *Parking Gate FSM* to inform it that the vehicle has paid for the ticket and is ready to leave the parking lot;

Figure 2 shows the top level of the parking gate system, while Figure 3 shows the parking gate system FSM.

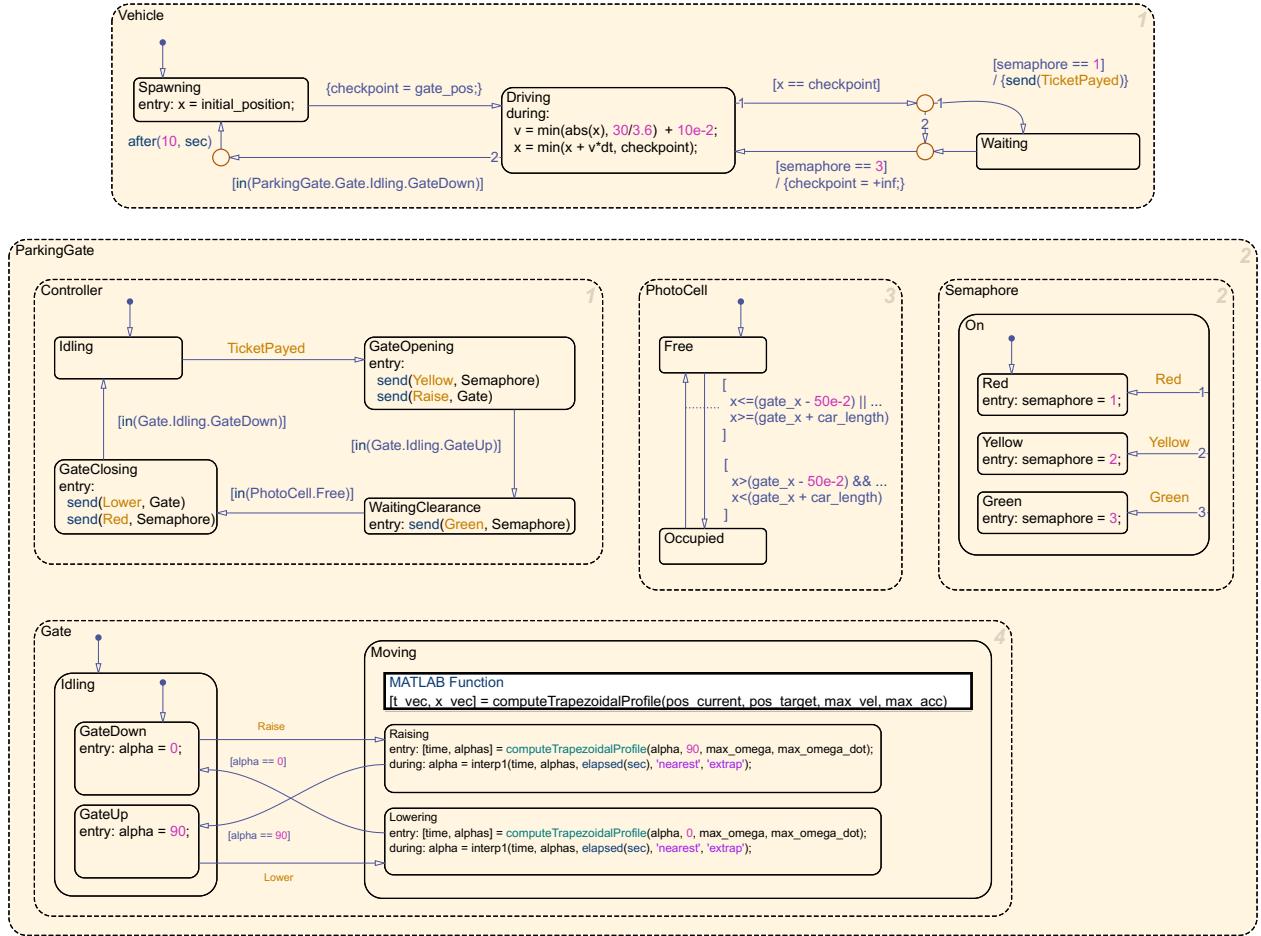


Figure 3: Parking Gate System FSM

3.2.1 Vehicle FSM

The vehicle FSM is responsible for simulating the vehicle's behavior in the parking lot. It has two main states: *Driving* and *Waiting*. At first, the vehicle is initialized and given a random position ahead of the gate. A checkpoint is also set so that it will stop at the gate before proceeding to the next state. Once at the gate, the vehicle looks for the semaphore state: if the semaphore is green, it directly returns to the *Driving* state and proceeds ahead to the next checkpoint which is now set to $+\infty$; instead, if the semaphore is red, it simulates the payment of the ticket by sending the event *TicketPayed* to the parking gate FSM and goes into the *Waiting* state, waiting for the semaphore to turn green. The vehicle proceed until the gate is closed again, when the external environment resets the vehicle to a new random position ahead of the gate to start the process again. Figure 4 shows the vehicle FSM.

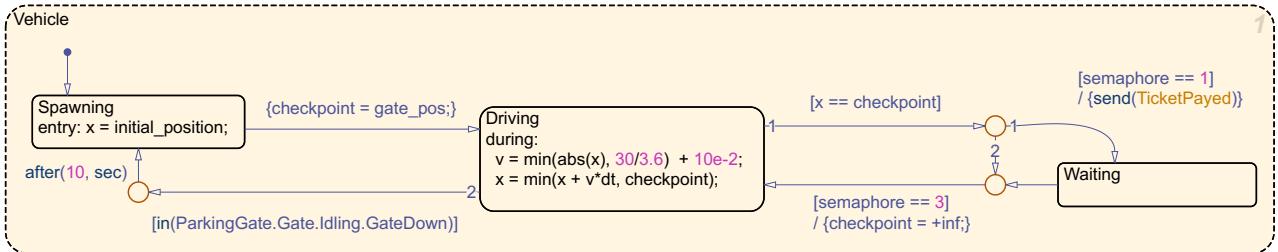


Figure 4: Vehicle FSM

Notice that the vehicle FSM has been modelled adopting the default *OR* logic, meaning that the vehicle can assume only one state at a time (i.e. *Driving* or *Waiting*, or *Spawning* when the vehicle is initialized).

3.2.2 Parking Gate FSM

The parking gate FSM is responsible for controlling at first the gate's opening and closing, and then the semaphore state. It has four main blocks: *Controller*, *Gate*, *Semaphore* and *PhotoCell*.

Controller The entry point of the parking gate FSM is the *Controller* block, which is responsible for managing the overall state of the system and coordinating the other blocks. It has been designed to be event-driven, meaning that it communicates with the subcomponents of the system by sending and receiving events. It has four main states: *Idling*, *GateOpening*, *WaitingClearance* and *GateClosing*. At first, the *Controller* block is in the *Idling* state, waiting for the vehicle to arrive and receive the *TicketPayed* event. Once the ticket is paid, the *Controller* block moves to the *GateOpening* state, where it sends the *Raise* event to the *Gate* block and the *Yellow* event to the *Semaphore* block. Once the gate is fully opened, the semaphore turns green and the *Controller* block moves to the *WaitingClearance* state, where it waits for the vehicle to pass through the gate. Notice that the system determines the vehicle's position by using the *PhotoCell* block, which is responsible for detecting the vehicle's presence in front of the gate. Once the vehicle has passed through the gate, the *Controller* block moves to the *GateClosing* state, where it sends the *Lower* event to the *Gate* block and the *Red* event to the *Semaphore* block. The *Controller* block then moves back to the *Idling* state, waiting for the next vehicle to arrive.

Figure 5 shows the *Controller* block of the parking gate FSM.

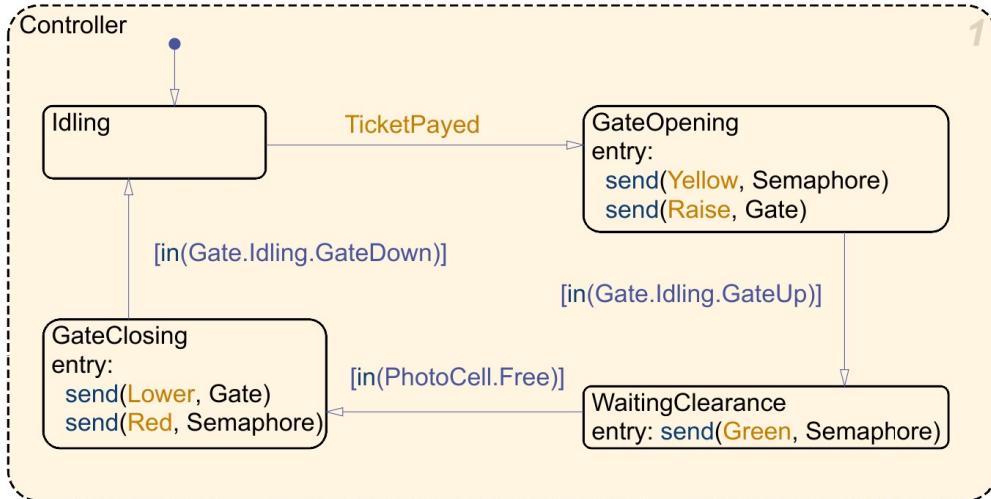


Figure 5: Control Block

Gate As already mentioned, the *Gate* block is responsible for controlling the gate's position and ensuring that it opens and closes correctly. It's designed based on two hierarchical states: *Idling* and *Moving*, which then are further divided into *GateUp* and *GateDown* states, and *Raising* and *Lowering* states, respectively. At first, the *Gate* block is in the *GateDown* state, where the gate is closed and the vehicle is not allowed to pass through. Once the *Raise* event is received from the *Controller* block, the *Gate* block moves to the *Raising* state, where it starts to open the gate. Once the gate is fully opened, the *Gate* block moves to the *GateUp* state waiting for

the *Lower* event from the *Controller* block that will then perform the same steps in reverse order, moving to the *Lowering* state and then to the *GateDown* state once the gate is fully closed.

Figure 6 shows the *Gate* block of the parking gate FSM.

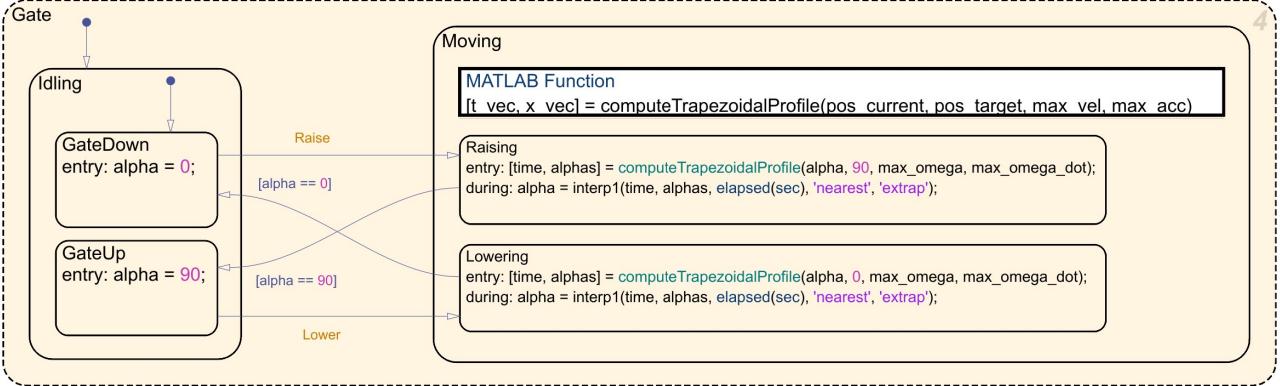


Figure 6: Gate Block

Notice that, in order to respect the constraints $w_{max} = |4|[\text{°}/\text{s}]$ and $\dot{w}_{max} = |1|[\text{°}/\text{s}^2]$, the position of the gate is updated based on the trapezoidal motion profile, generated when entering the *Raising* or *Lowering* states. The trapezoidal motion profile is defined by the following equations:

$$w(t) = \begin{cases} \frac{w_{max}}{t_1} t & 0 \leq t < t_1 \\ w_{max} & t_1 \leq t < t_2 \\ w_{max} - \frac{w_{max}}{t_3}(t - T) & t_2 \leq t < T \end{cases} \quad (1)$$

Where t_1 is the time when the gate reaches the maximum speed and t_2 is the time when the gate starts to decelerate. The trapezoidal motion profile is designed to ensure that the gate opens and closes smoothly, without any sudden jerks or stops. Listing 1 shows the implementation of the trapezoidal motion profile as MATLAB function.

```

1 function [t_vec, x_vec] = computeTrapezoidalProfile(pos_current, pos_target, max_vel,  
    max_acc)  
2 % COMPUTE_TRAPEZOIDAL_PROFILE Generates time and angle points for a trapezoidal profile  
3 % At first checks if a trapezoidal profile is feasible or is necessary to  
4 % fallback to a triangular one.  
5  
6 delta_pos = pos_target - pos_current;  
7 time_acc_phase = max_vel / max_acc;  
8 delta_pos_acc_phase = 0.5 * max_acc * time_acc_phase^2;  
9  
10 if 2 * delta_pos_acc_phase < abs(delta_pos)  
11     % Trapezoidal profile  
12     t1 = time_acc_phase;  
13     t2 = (abs(delta_pos) - 2 * delta_pos_acc_phase) / max_vel;  
14     t3 = time_acc_phase;  
15 else  
16     % Triangular profile  
17     max_vel = sqrt(abs(delta_pos) * max_acc);  
18     t1 = max_vel / max_acc;  
19     t2 = 0;  
20     t3 = max_vel / max_acc;  
21 end  
22  
23  
24 % Trapezoidal profile template (HP: s0 = 0, v0 = 0)  
25 acc = max_acc * sign(delta_pos);  
26 vel = max_vel * sign(delta_pos);  
27 profile_template = @(t) ...  
28     (t <= t1) .* (1/2 * acc * t.^2) + ...

```

```

29   (t > t1 & t <= (t1 + t2)) .*          (1/2 * acc * t1^2 + vel * (t - t1)) + ...
30   (t > (t1 + t2) & t <= (t1 + t2 + t3)) .* (1/2 * acc * t1^2 + vel * t2 + vel * (t - 
31   t2 + t1)) - 1/2 * acc * (t - (t1 + t2)).^2) + ...
32   (t > (t1 + t2 + t3)) .*          (1/2 * acc * t1^2 + vel * t2 + vel * t3 - 
33 % Profile generation
34 t_vec = linspace(0, t1 + t2 + t3, 100);
35 x_vec = profile_template(t_vec) + pos_current;
36
37 end

```

Listing 1: Trapezoidal Motion Profile

PhotoCell Despite being used only during the *WaitingClearance* state of the *Controller* block, the *PhotoCell* block is responsible for detecting the vehicle's presence in front of the gate. This, in fact, is the only way to understand if the vehicle has passed through the gate and the gate can be closed again safely. In order to continuously sample the vehicle's position, the *PhotoCell* block is designed to not be event-driven, but to move continuously between the two states *Free* and *Occupied* based on the vehicle's position. By doing so, an external agent can continuously monitor the vehicle's position by checking the state of the *PhotoCell* block. In order to determine the presence or not of the vehicle, the *PhotoCell* block uses the knowledge of the vehicle's length and position, the latter of which is shared by the *Vehicle FSM*.

Figure 7 shows the *PhotoCell* block of the parking gate FSM.

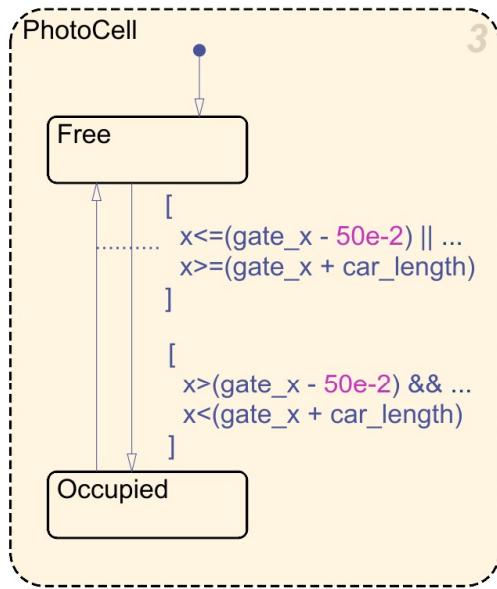


Figure 7: PhotoCell Block

Semaphore Last but not least, the *Semaphore* block is responsible for switching the semaphore state between red, yellow and green based on the events received from the *Controller* block. It has one single state *On* composed of three sub-states representative of the color of the semaphore. Based on the event received, the correct sub-state is chosen and the semaphore state is then shared with the *Vehicle FSM* to inform it about the state of the gate (open, closed, etc.). This is done in order to inform the vehicle about the state of the gate and allow it to proceed or stop based on the semaphore state, which is the only information that the vehicle has about the parking gate system (i.e. the vehicle doesn't know if the gate is open or closed, but only if the semaphore is red or green).

Figure 8 shows the *Semaphore* block of the parking gate FSM.

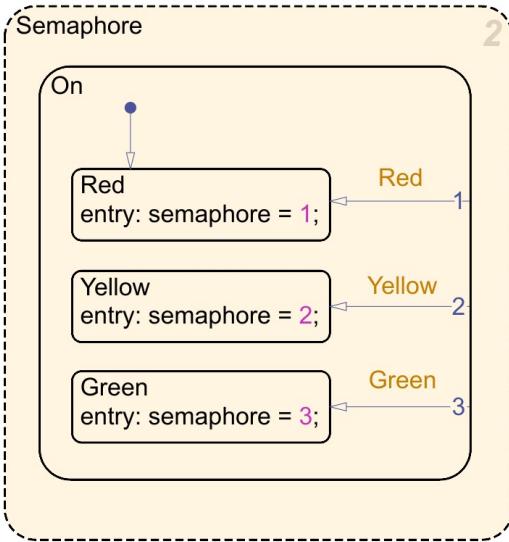


Figure 8: Semaphore Block

3.3 Simulation Results

The simulation has been run for 100s and the vehicle has been reset to a new random position right after the closure of the gate. The following figures show the results of the simulation.

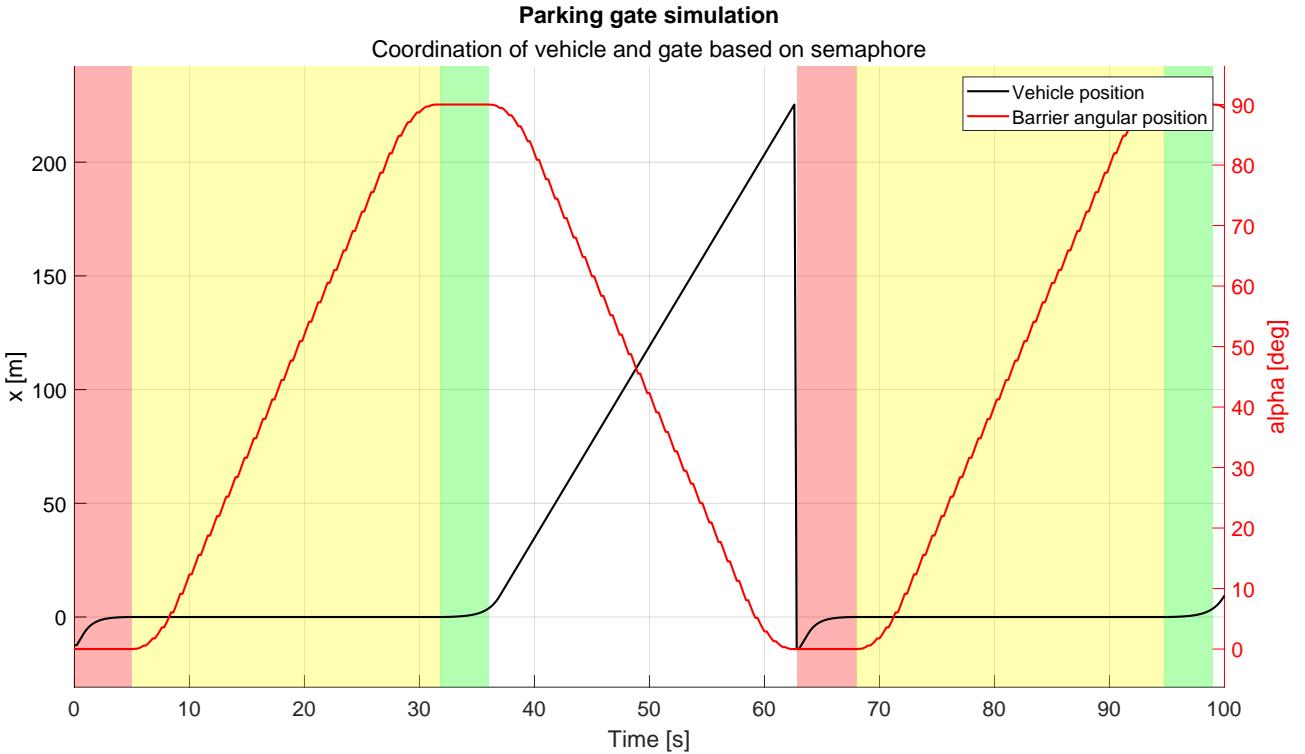


Figure 9: Simulation Results

Figure 9 shows the vehicle position (in black) and the gate angular position (in red) over time. The plot has also been divided into different regions based on the state of the *Semaphore* block (red, yellow and green). Notice that for clarity, the semaphore state has been added to the plot only when actually visible by the vehicle (i.e. when the vehicle is in front of the gate and not yet passed through it).

Instead, a more detailed view of the underlying coordination between the two FSMs is shown in Figure 10, where a visualization of the state transitions and events is shown. The plot report on the vertical axis the time of the simulation and on each vertical dashed line the state of each superstate of the FSMs is shown. One can appreciate how the two FSMs are coordinated and how the events are sent and received between them, along with the transitions between the states.

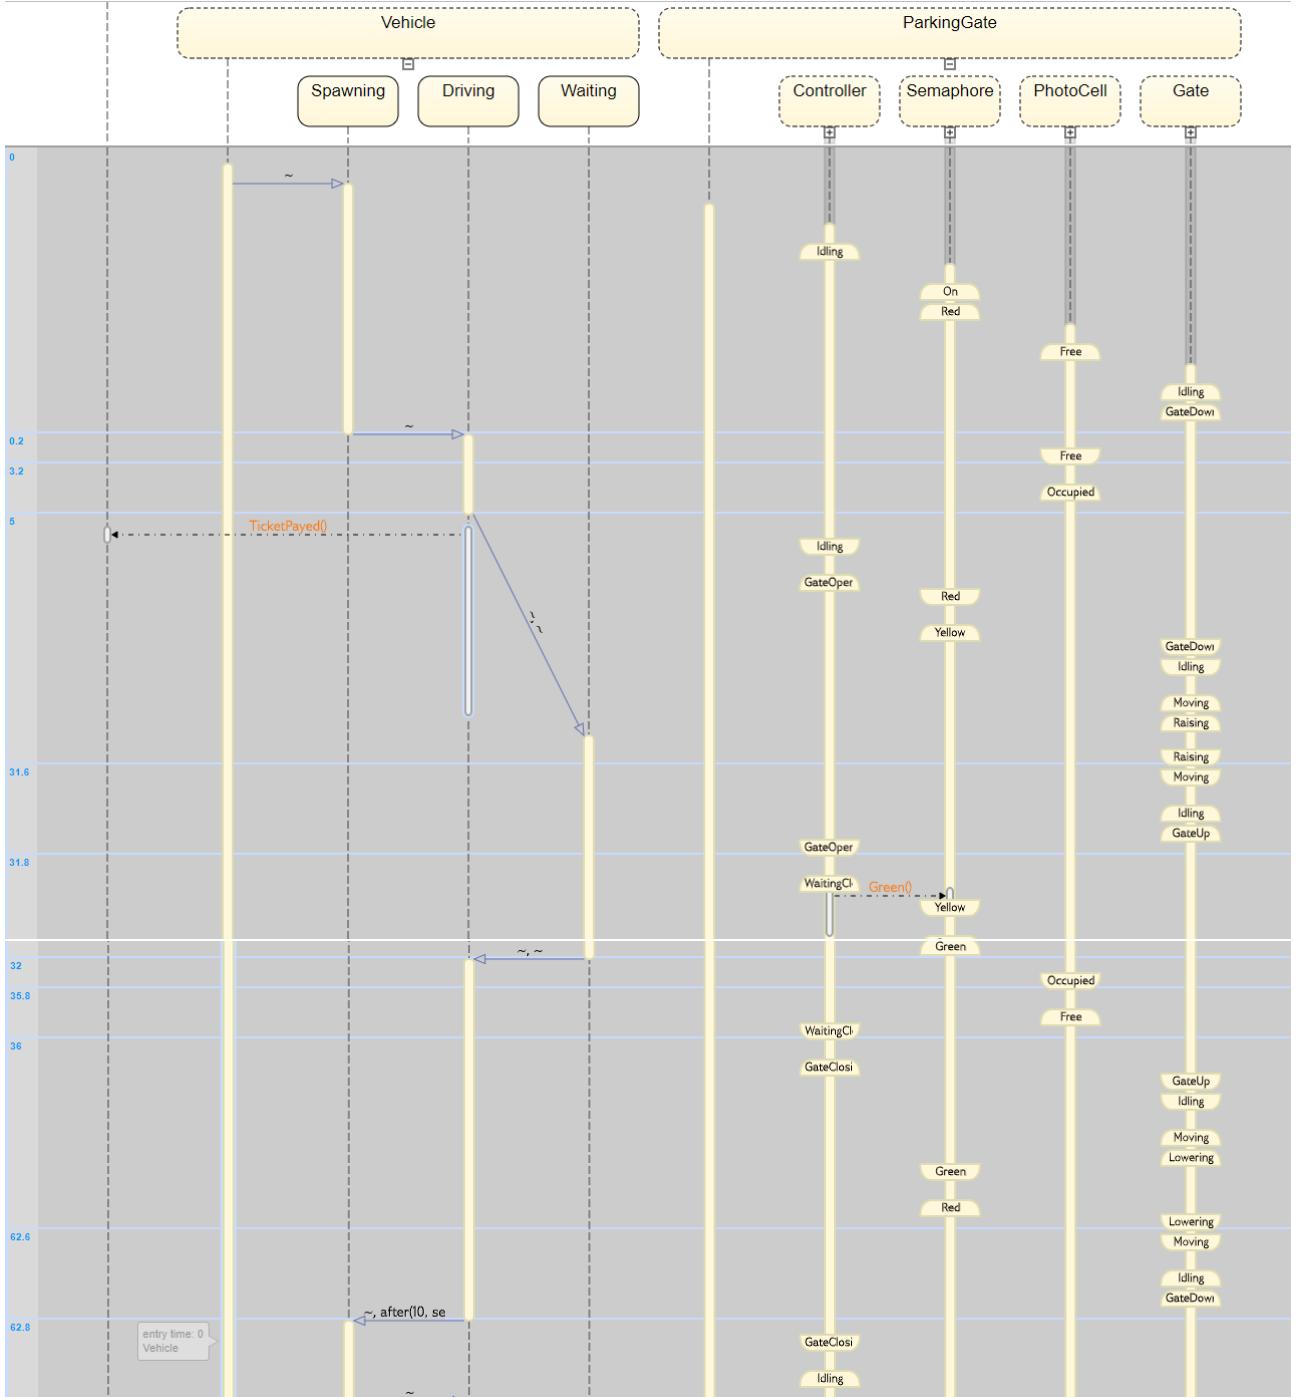


Figure 10: Simulation Results - Detailed View

4 Conclusions

In this short report, we have presented the results of the fifth assignment of the course on Autonomous Vehicles. We have briefly explained what FMSs are and how they can be used to model the behavior of generic systems. We have also presented a possible FMSs implementation about a car parking system, which is able to coordinate the behavior of the car and the parking gate. The results presented have shown the effectiveness of the proposed solution.

The current solution could be improved under many aspects. For example, one could think of implementing features like error handling, or ticket authentication and validation, or even a more complex parking system composed of multiple parking gates and cars that must coordinate over the network.

We believe that FMSs are a power tool that can often replace traditional programming paradigms, allowing to model complex systems in a more intuitive way. Further studies on this topic could be useful to understand the full potential of FMSs and their applications in the field of autonomous vehicles and robotics in general.

References

- [1] Wikipedia contributors. Finite-state machine — Wikipedia, the free encyclopedia. https://en.wikipedia.org/w/index.php?title=Finite-state_machine&oldid=1285525354, 2025. [Online; accessed 26-April-2025].

A Appendix

The system presented in Section 3.2, leverages one single FSM including under the hood both the vehicle and the parking gate FSMs. However, this solution is not the most elegant one given that the two FSMs are actually two separate entities that communicate with each other via external signals. For this reason, the first solution and implementation of the parking gate system was the one shown in Figure 11.

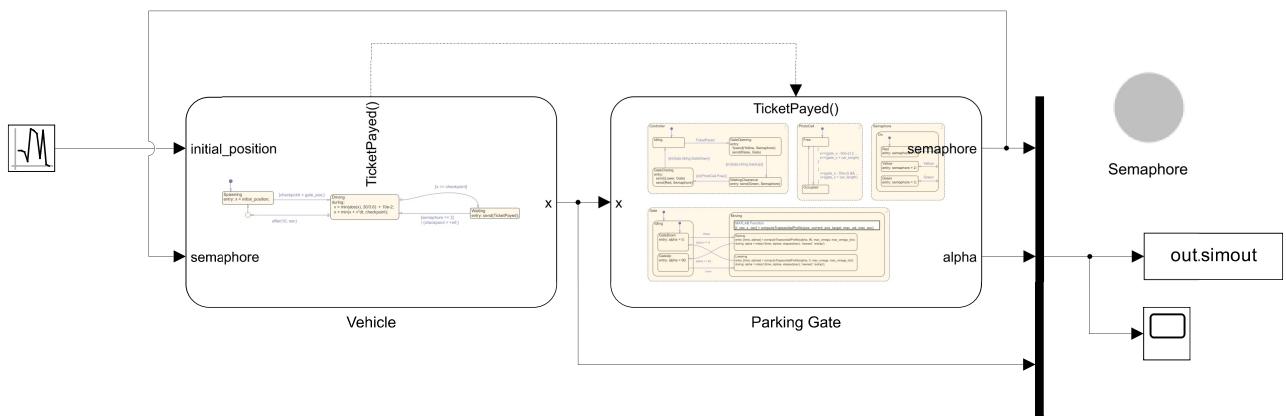


Figure 11: Original Parking Gate System Overview

Notice that the underlying components and states are exactly the same as the ones presented in the previous section. Nonetheless, the structure of Figure 11 resolved in unexpected behavior of the system and never generated the expected results.

Even with the help of debugging tools, it was impossible to understand the reason behind the unexpected behavior of the system and the author had to simplify the structure and fallback to the single FSM solution discussed during the document (compromise between elegance and functionality).

Autonomous Vehicles
Implementation of sampled-based motion planning algorithms

Tommaso Bocchietti 10740309

A.Y. 2024/25



POLITECNICO
MILANO 1863

Contents

| | | |
|----------|--------------------------------------|-----------|
| 1 | Introduction | 3 |
| 2 | Sampling-based planning | 3 |
| 2.1 | Single-query algorithms | 4 |
| 3 | Algorithm implementation | 5 |
| 3.1 | RRT algorithm | 5 |
| 3.2 | RRT* algorithm | 7 |
| 3.3 | RRT-kinematic algorithm | 9 |
| 4 | Algorithm Testing | 12 |
| 4.1 | Map scenarios | 12 |
| 4.2 | Results analysis | 13 |
| 4.2.1 | Map 1 | 14 |
| 4.2.2 | Map 2 | 15 |
| 4.2.3 | Map 3 | 17 |
| 4.2.4 | Map 4 | 18 |
| 5 | Grid vs Sample Based Planning | 19 |
| 5.1 | Path Asymptotic Optimality | 19 |
| 5.2 | Multidimensional Workspace | 21 |
| 6 | Conclusions | 23 |

List of Figures

| | | |
|----|--|----|
| 1 | Generic sampling-based algorithm flowchart | 4 |
| 2 | Example of the rewiring process in RRT* algorithm. Credit to [1]. | 9 |
| 3 | Bicycle kinematic model. Image credit to Thomas Fermi. | 11 |
| 4 | Test maps used for the algorithm testing. | 13 |
| 5 | Tree generated and path found by the RRT, RRT* and RRT Kinematic algorithms on map 1 | 14 |
| 6 | Tree generated and path found by the RRT, RRT* and RRT Kinematic algorithms on map 2 | 16 |
| 7 | Tree generated and path found by the RRT, RRT* and RRT Kinematic algorithms on map 3 | 17 |
| 8 | Tree generated and path found by the RRT, RRT* and RRT Kinematic algorithms on map 4 | 18 |
| 9 | Path generated by RRT* (left) and A* (right) on the same scenario. | 19 |
| 10 | Path generated by RRT* (left) and A* (right) on the same scenario. | 20 |
| 11 | Asymptotic optimality of RRT* algorithm. Both the images refer to 10000 iterations of the algorithm. | 20 |
| 12 | 3D grid of obstacles used as a scenario for the tests. | 21 |
| 13 | Path generated by A* (top), RRT (middle) and RRT* (bottom) algorithms on the same 3D scenario. | 22 |

1 Introduction

Motion planning is a fundamental aspect of autonomous vehicle navigation, enabling them to navigate complex environments while avoiding obstacles and reaching their destinations efficiently.

While in previous work we have focused on the implementation of two grid-based motion planning algorithms, namely Dijkstra and A*, this report aims to implement and analyze sampled-based motion planning algorithms. Specifically, we focus on Rapidly-exploring Random Trees (RRT) and two of its variants, RRT* and a kinematically-learnt version of RRT.

After a brief introduction to the sampled-based motion planning approach, we present the implementation of the algorithms, followed by a discussion of the results obtained by running them on a set of predefined map scenarios. An extensive comparison of the performance between the previously implemented A* algorithm and the RRT* algorithm is also reported, highlighting the strengths and weaknesses of each approach.

Report Structure The report is structured as follows:

- In Section 2, we provide a brief introduction to the sampled-based motion planning algorithms, giving an overview of the single-query and multi-query approaches, and discussing the RRT algorithm and its variants.
- In Section 3, we present all the three algorithms, RRT, RRT* and RRT-kinematics, explaining the main features of each one and the implementation details.
- In Section 4, we present the results obtained by running the algorithms on a set of predefined map scenarios, and we discuss their performance in terms of computation time, path length and number of nodes sampled.
- In Section 5, we compare the performance of the RRT* algorithm with the A* algorithm, highlighting the strengths and weaknesses of each approach.
- In Section 6, we summarize the main findings of the report and discuss the future work that can be done to improve the algorithms and their performance.

Tools As for the tools used, MATLAB is employed as the main platform for implementing the algorithms and performing the analysis of the results.

2 Sampling-based planning

In the field of robotics, a variety of algorithms are available for path planning, each with its own strengths and weaknesses. One of the most ancient and well known class of algorithms is the so called *Grid-based*, which is based on the discretization of the environment into a grid of cells and nodes. As we have already discovered during the previous assignment, this approach is particularly useful for environments with a known structure, such as indoor environments or outdoor environments with well-defined obstacles but is limited in terms of scalability and flexibility.

In contrast, *Sampling-based* algorithms are designed to work in high-dimensional spaces and can handle much more complex environments. For this reason, most of the modern path planning algorithms are based on sampling-based methods, which are able to efficiently explore the configuration space of the robot and find a feasible path from the start to the goal configuration.

Among the many robotics applications that benefit from sampling-based algorithms, we can mention:

- Terrestrial robots: path planning for autonomous vehicles in outdoor environments, such as urban areas or off-road terrains. The ability to handle large environments is crucial for the success of these applications.
- Robotic manipulation: path planning for robotic arms in cluttered environments, where the robot needs to avoid obstacles and reach a specific target position. The ability to handle high-dimensional configuration spaces is essential for these applications.

Sampling-based algorithms are further divided into two main classes: *single-query* and *multi-query* algorithms. Single-query algorithms are designed to find a path from a specific start configuration to a specific goal configuration, while multi-query algorithms are designed to find paths between multiple pairs of start and goal configurations.

In the following sections, we focus exclusively on the single-query algorithms, leaving the multi-query algorithms for future discussions.

2.1 Single-query algorithms

As it was for the grid-based ones, also the single-query sampling-based algorithms follow a structure that is in common to all of them. The main idea here is to start from a given configuration and then randomly sample points in the configuration space, connecting them to form a tree if the condition of collision-free is satisfied. The tree is then incrementally expanded until a path from the start configuration to the goal configuration is found.

The fundamental difference between grid-based and sampling-based algorithms is that the former rely on a discretized representation of the environment, while the latter rely on a continuous representation of the configuration space. This allows sampling-based algorithms to handle complex environments with arbitrary shapes and obstacles, as they do not rely on a fixed grid structure.

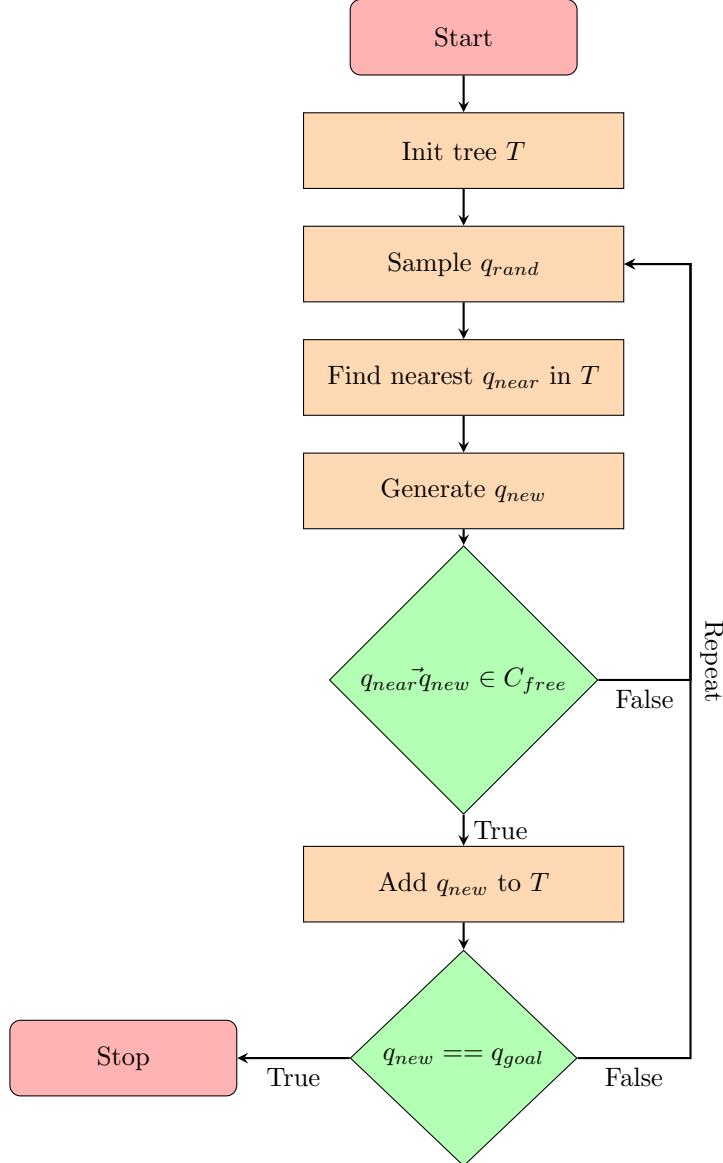


Figure 1: Generic sampling-based algorithm flowchart

The flowchart in Figure 1 summarizes the main steps of a generic sampling-based algorithm. One can notice that tree T is built incrementally at each iteration, starting from the initial configuration and adding new sampled points to the tree if they respect the collision-free condition. As already said, this is in fact the main difference between the sampling-based algorithms and the grid-based algorithms, which rely on a complete discretization of the environment made in advance.

However, the sampling-based algorithms suffers from a major drawback, which is the lack of completeness and optimality. In particular, the algorithms are not guaranteed to find a solution, even if one exists, and the solution found may not be optimal. As we discuss in the following sections, improvements have been made to the core structure to address these issues, but they still remain a challenge in the field of robotics and motion planning.

There exist many variations built on top of the core algorithm presented in Figure 1. The most common ones are:

- **Rapidly-exploring Random Tree (RRT)**: a single-query algorithm that just implements the core algorithm presented in Figure 1. It's the simplest, but also the most suboptimal algorithm among the ones presented in this project.
- **RRT***: an extension of the RRT algorithm that improves the optimality of the solution by rewiring the tree to find shorter paths. It is more computationally expensive than the RRT algorithm, but it is able to refine the solutions in terms of path length.
- **RRT-connect**: a variant of the RRT algorithm that builds two trees, one from the start configuration and one from the goal configuration, and connects them to find a path. It's usually faster than the RRT algorithm to find a solution given that the search space is explored from both ends and also benefits from some form of optimality.
- **RRT-kinematic**: a variant of the RRT algorithm that takes into account the kinematic constraints of the robot/vehicle during the planning process. This is particularly useful for vehicles with non-holonomic constraints, such as terrestrial vehicles or mobile robots. The algorithm generates a tree of feasible trajectories that respect the kinematic constraints of the robot.

The following sections focus on the in-depth analysis and implementation of the RRT, RRT* and RRT-kinematic algorithms. Notice that despite being a popular and widely used algorithm, the RRT-connect algorithm is not going to be implemented in this project due to time constraints.

3 Algorithm implementation

In this section, we present the implementation of the three algorithms we are going to analyze in this report. The algorithms are implemented in MATLAB and are designed to handle multiple dimensions problems, meaning that they can be used for both N-D environments. In the following only the core of each algorithm is presented, leaving the code of all the auxiliaries classes such as `Graph`, `Edge`, or `Node` among the files in the repository.

3.1 RRT algorithm

The most basic version of the single-query sampling-based algorithm is the Rapidly-exploring Random Tree (RRT) algorithm. The RRT algorithm is nothing more than the direct implementation of the core algorithm presented in Figure 1.

Listing 1 shows a possible implementation of the RRT algorithm.

```

1 function [exist, distances, parents] = rrt(G, map, node_initial, node_final)
2 % RRT Builds a tree using the Rapidly-exploring Random Tree (RRT) algorithm.
3 %
4 % [EXIST, DISTANCES, PARENTS] = RRT(G, MAP, NODE_INITIAL, NODE_FINAL)
5 % attempts to grow a tree from NODE_INITIAL to NODE_FINAL using random
6 % sampling and a step-limited steering function.
7 %
8 % Inputs:
9 %     G           - Graph object representing the RRT.
10 %     MAP         - Binary occupancy grid (1 = free, 0 = obstacle).
11 %     NODE_INITIAL - Node object with initial state.
12 %     NODE_FINAL   - Node object with final/goal state.
13 %
14 % Outputs:
15 %     EXIST      - Boolean flag indicating whether a path to the goal was found.
16 %     DISTANCES  - Vector of cumulative distances from the start to each node.
17 %     PARENTS    - Vector of parent node IDs for reconstructing the tree.
18 %
19 % The tree is expanded by selecting random points and connecting the closest
20 % node in the tree toward them, checking each connection for validity.
21
22 q_new = node_initial.state;
23 distances = zeros(1, 1);

```

```

24 parents = NaN(1, 1);
25 iter = 0;
26
27 while (Node.euclideanDistance(q_new, node_final.state) > G.step_length && iter < G.
28     max_iter)
29
30     % Sampling stage
31     q_rand = rand(1, ndims(map)) .* size(map);
32
33     % Finding the nearest node
34     IDs_near = G.getNearbyIDs(q_rand);
35     ID_near = IDs_near(1);
36     q_near = G.getNodeByID(ID_near).state;
37
38     % Steering stage
39     q_new = G.steerState(q_near, q_rand, G.step_length);
40
41     % Collision checking and tree update
42     if (G.isValidConnection(q_near, q_new, map))
43
44         ID_new = G.addNode(q_new);
45         dist = Node.euclideanDistance(q_near, q_new);
46         G.addEdge(ID_near, ID_new, dist);
47
48         parents(ID_new, 1) = ID_near;
49         distances(ID_new, 1) = distances(ID_near) + dist;
50
51     end
52
53     iter = iter + 1;
54
55 end
56
57 exist = ~(iter == G.max_iter);
58
59 % Connection with final node
60 if (exist)
61     ID_final = G.addNode(node_final.state);
62     IDs_near = G.getNearbyIDs(node_final.state);
63     ID_near = IDs_near(1);
64     q_near = G.getNodeByID(ID_near).state;
65     dist = Node.euclideanDistance(q_near, node_final.state);
66     G.addEdge(ID_near, ID_final, dist);
67     parents(ID_final, 1) = ID_near;
68     distances(ID_final, 1) = distances(ID_near) + dist;
69 end
70

```

Listing 1: Node class used to represent a node in the graph.

The algorithm takes as input the graph G , the map M , the start node s , and the goal node g . It starts by initializing the tree T with the start node and enters a loop where at first a random point/configuration is sampled from the workspace. Then, it finds the nearest node already in the tree and generates a new node by moving in the direction of the sampled point. If the new node is in collision with the obstacles, it's discarded and the algorithm skips to the next iteration. If the new node is valid, it's added to the tree. A final step is to check if the new node is close enough to the goal node, in which case the algorithm terminates and returns the path from the start node to the goal node, otherwise the algorithm continues to sample new points until a maximum number of iterations is reached, or the goal node is reached.

The RRT algorithm is simple and efficient, but lacks optimality in the generated path which is normally jagged and not smooth. The algorithm is also not complete, meaning that it may not find a solution even if one exists in the configuration space due to the maximum number of iterations limit. However, it has been proved that the

RRT algorithm is probabilistically complete, meaning that as the number of iterations increases, the probability of finding a solution approaches 1 (at the cost of longer computation time, of course).

3.2 RRT* algorithm

RRT* is an extension of the RRT algorithm that improves the optimality of the solution by rewiring the tree to find shorter paths. The main idea behind the RRT* algorithm is to keep track of the cost of reaching each node in the tree and to rewire the tree whenever a new node is added. This is done by checking if the new node can be used as a cheaper parent connection for any of the existing nodes in the tree.

Listing 2 shows a possible implementation of the RRT* algorithm.

```

1 function [exist, distances, parents] = rrtStar(G, map, node_initial, node_final)
2 % RRTSTAR Builds a tree using the RRT* (optimal Rapidly-exploring Random Tree) algorithm
3 %
4 % [EXIST, DISTANCES, PARENTS] = RRTSTAR(G, MAP, NODE_INITIAL, NODE_FINAL)
5 % attempts to find an optimal path from NODE_INITIAL to NODE_FINAL
6 % in the graph G by incrementally building a tree of collision-free paths
7 % using random sampling, rewiring, and cost minimization.
8 %
9 % Inputs:
10 %     G          - Graph object representing the RRT.
11 %     MAP        - Binary occupancy grid (1 = free, 0 = obstacle).
12 %     NODE_INITIAL - Node object with initial state.
13 %     NODE_FINAL   - Node object with final/goal state.
14 %
15 % Outputs:
16 %     EXIST      - Boolean flag indicating whether a path to the goal was found.
17 %     DISTANCES  - Vector of cumulative distances from the start to each node.
18 %     PARENTS    - Vector of parent node IDs for reconstructing the tree.
19 %
20 % The algorithm improves upon RRT by locally rewiring nodes to ensure that
21 % each node is connected through the lowest-cost path, producing asymptotically
22 % optimal solutions.
23
24 q_new = node_initial.state;
25 distances = zeros(1, 1);
26 parents = NaN(1, 1);
27 iter = 0;
28
29 while (Node.euclideanDistance(q_new, node_final.state) > G.step_length && iter < G.
30 max_iter)
31
32     % Sampling stage
33     q_rand = rand(1, ndims(map)) .* size(map);
34
35     % Finding the nearest node
36     IDs_near = G.getNearbyIDs(q_rand);
37     ID_near = IDs_near(1);
38     q_near = G.getNodeByID(ID_near).state;
39
40     % Steering stage
41     q_new = G.steerState(q_near, q_rand, G.step_length);
42
43     % Collision checking and tree update
44     if (G.isValidConnection(q_near, q_new, map))
45
46         % Node creation
47         ID_new = G.addNode(q_new);
48
49         % Neighborhood selection
50
51         % Rewire stage
52         for i = 1:length(distances)
53             if (distances(i) >= G.step_length)
54                 % Check if new node is closer
55                 if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
56                     % Update parent and distance
57                     parents(i) = ID_new;
58                     distances(i) = sqrt((q_new - G.nodes(i).state).^2);
59                 end
60             end
61         end
62
63         % Prune stage
64         for i = 1:length(distances)
65             if (distances(i) >= G.step_length)
66                 % Check if new node is closer
67                 if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
68                     % Update parent and distance
69                     parents(i) = ID_new;
70                     distances(i) = sqrt((q_new - G.nodes(i).state).^2);
71                 end
72             end
73         end
74
75         % Local optimization stage
76         for i = 1:length(distances)
77             if (distances(i) >= G.step_length)
78                 % Check if new node is closer
79                 if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
80                     % Update parent and distance
81                     parents(i) = ID_new;
82                     distances(i) = sqrt((q_new - G.nodes(i).state).^2);
83                 end
84             end
85         end
86
87         % Prune stage
88         for i = 1:length(distances)
89             if (distances(i) >= G.step_length)
90                 % Check if new node is closer
91                 if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
92                     % Update parent and distance
93                     parents(i) = ID_new;
94                     distances(i) = sqrt((q_new - G.nodes(i).state).^2);
95                 end
96             end
97         end
98
99         % Local optimization stage
100        for i = 1:length(distances)
101            if (distances(i) >= G.step_length)
102                % Check if new node is closer
103                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
104                    % Update parent and distance
105                    parents(i) = ID_new;
106                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
107                end
108            end
109        end
110
111        % Prune stage
112        for i = 1:length(distances)
113            if (distances(i) >= G.step_length)
114                % Check if new node is closer
115                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
116                    % Update parent and distance
117                    parents(i) = ID_new;
118                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
119                end
120            end
121        end
122
123        % Local optimization stage
124        for i = 1:length(distances)
125            if (distances(i) >= G.step_length)
126                % Check if new node is closer
127                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
128                    % Update parent and distance
129                    parents(i) = ID_new;
130                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
131                end
132            end
133        end
134
135        % Prune stage
136        for i = 1:length(distances)
137            if (distances(i) >= G.step_length)
138                % Check if new node is closer
139                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
140                    % Update parent and distance
141                    parents(i) = ID_new;
142                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
143                end
144            end
145        end
146
147        % Local optimization stage
148        for i = 1:length(distances)
149            if (distances(i) >= G.step_length)
150                % Check if new node is closer
151                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
152                    % Update parent and distance
153                    parents(i) = ID_new;
154                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
155                end
156            end
157        end
158
159        % Prune stage
160        for i = 1:length(distances)
161            if (distances(i) >= G.step_length)
162                % Check if new node is closer
163                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
164                    % Update parent and distance
165                    parents(i) = ID_new;
166                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
167                end
168            end
169        end
170
171        % Local optimization stage
172        for i = 1:length(distances)
173            if (distances(i) >= G.step_length)
174                % Check if new node is closer
175                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
176                    % Update parent and distance
177                    parents(i) = ID_new;
178                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
179                end
180            end
181        end
182
183        % Prune stage
184        for i = 1:length(distances)
185            if (distances(i) >= G.step_length)
186                % Check if new node is closer
187                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
188                    % Update parent and distance
189                    parents(i) = ID_new;
190                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
191                end
192            end
193        end
194
195        % Local optimization stage
196        for i = 1:length(distances)
197            if (distances(i) >= G.step_length)
198                % Check if new node is closer
199                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
200                    % Update parent and distance
201                    parents(i) = ID_new;
202                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
203                end
204            end
205        end
206
207        % Prune stage
208        for i = 1:length(distances)
209            if (distances(i) >= G.step_length)
210                % Check if new node is closer
211                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
212                    % Update parent and distance
213                    parents(i) = ID_new;
214                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
215                end
216            end
217        end
218
219        % Local optimization stage
220        for i = 1:length(distances)
221            if (distances(i) >= G.step_length)
222                % Check if new node is closer
223                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
224                    % Update parent and distance
225                    parents(i) = ID_new;
226                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
227                end
228            end
229        end
230
231        % Prune stage
232        for i = 1:length(distances)
233            if (distances(i) >= G.step_length)
234                % Check if new node is closer
235                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
236                    % Update parent and distance
237                    parents(i) = ID_new;
238                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
239                end
240            end
241        end
242
243        % Local optimization stage
244        for i = 1:length(distances)
245            if (distances(i) >= G.step_length)
246                % Check if new node is closer
247                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
248                    % Update parent and distance
249                    parents(i) = ID_new;
250                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
251                end
252            end
253        end
254
255        % Prune stage
256        for i = 1:length(distances)
257            if (distances(i) >= G.step_length)
258                % Check if new node is closer
259                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
260                    % Update parent and distance
261                    parents(i) = ID_new;
262                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
263                end
264            end
265        end
266
267        % Local optimization stage
268        for i = 1:length(distances)
269            if (distances(i) >= G.step_length)
270                % Check if new node is closer
271                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
272                    % Update parent and distance
273                    parents(i) = ID_new;
274                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
275                end
276            end
277        end
278
279        % Prune stage
280        for i = 1:length(distances)
281            if (distances(i) >= G.step_length)
282                % Check if new node is closer
283                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
284                    % Update parent and distance
285                    parents(i) = ID_new;
286                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
287                end
288            end
289        end
290
291        % Local optimization stage
292        for i = 1:length(distances)
293            if (distances(i) >= G.step_length)
294                % Check if new node is closer
295                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
296                    % Update parent and distance
297                    parents(i) = ID_new;
298                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
299                end
300            end
301        end
302
303        % Prune stage
304        for i = 1:length(distances)
305            if (distances(i) >= G.step_length)
306                % Check if new node is closer
307                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
308                    % Update parent and distance
309                    parents(i) = ID_new;
310                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
311                end
312            end
313        end
314
315        % Local optimization stage
316        for i = 1:length(distances)
317            if (distances(i) >= G.step_length)
318                % Check if new node is closer
319                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
320                    % Update parent and distance
321                    parents(i) = ID_new;
322                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
323                end
324            end
325        end
326
327        % Prune stage
328        for i = 1:length(distances)
329            if (distances(i) >= G.step_length)
330                % Check if new node is closer
331                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
332                    % Update parent and distance
333                    parents(i) = ID_new;
334                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
335                end
336            end
337        end
338
339        % Local optimization stage
340        for i = 1:length(distances)
341            if (distances(i) >= G.step_length)
342                % Check if new node is closer
343                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
344                    % Update parent and distance
345                    parents(i) = ID_new;
346                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
347                end
348            end
349        end
350
351        % Prune stage
352        for i = 1:length(distances)
353            if (distances(i) >= G.step_length)
354                % Check if new node is closer
355                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
356                    % Update parent and distance
357                    parents(i) = ID_new;
358                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
359                end
360            end
361        end
362
363        % Local optimization stage
364        for i = 1:length(distances)
365            if (distances(i) >= G.step_length)
366                % Check if new node is closer
367                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
368                    % Update parent and distance
369                    parents(i) = ID_new;
370                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
371                end
372            end
373        end
374
375        % Prune stage
376        for i = 1:length(distances)
377            if (distances(i) >= G.step_length)
378                % Check if new node is closer
379                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
380                    % Update parent and distance
381                    parents(i) = ID_new;
382                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
383                end
384            end
385        end
386
387        % Local optimization stage
388        for i = 1:length(distances)
389            if (distances(i) >= G.step_length)
390                % Check if new node is closer
391                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
392                    % Update parent and distance
393                    parents(i) = ID_new;
394                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
395                end
396            end
397        end
398
399        % Prune stage
400        for i = 1:length(distances)
401            if (distances(i) >= G.step_length)
402                % Check if new node is closer
403                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
404                    % Update parent and distance
405                    parents(i) = ID_new;
406                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
407                end
408            end
409        end
410
411        % Local optimization stage
412        for i = 1:length(distances)
413            if (distances(i) >= G.step_length)
414                % Check if new node is closer
415                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
416                    % Update parent and distance
417                    parents(i) = ID_new;
418                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
419                end
420            end
421        end
422
423        % Prune stage
424        for i = 1:length(distances)
425            if (distances(i) >= G.step_length)
426                % Check if new node is closer
427                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
428                    % Update parent and distance
429                    parents(i) = ID_new;
430                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
431                end
432            end
433        end
434
435        % Local optimization stage
436        for i = 1:length(distances)
437            if (distances(i) >= G.step_length)
438                % Check if new node is closer
439                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
440                    % Update parent and distance
441                    parents(i) = ID_new;
442                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
443                end
444            end
445        end
446
447        % Prune stage
448        for i = 1:length(distances)
449            if (distances(i) >= G.step_length)
450                % Check if new node is closer
451                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
452                    % Update parent and distance
453                    parents(i) = ID_new;
454                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
455                end
456            end
457        end
458
459        % Local optimization stage
460        for i = 1:length(distances)
461            if (distances(i) >= G.step_length)
462                % Check if new node is closer
463                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
464                    % Update parent and distance
465                    parents(i) = ID_new;
466                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
467                end
468            end
469        end
470
471        % Prune stage
472        for i = 1:length(distances)
473            if (distances(i) >= G.step_length)
474                % Check if new node is closer
475                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
476                    % Update parent and distance
477                    parents(i) = ID_new;
478                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
479                end
480            end
481        end
482
483        % Local optimization stage
484        for i = 1:length(distances)
485            if (distances(i) >= G.step_length)
486                % Check if new node is closer
487                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
488                    % Update parent and distance
489                    parents(i) = ID_new;
490                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
491                end
492            end
493        end
494
495        % Prune stage
496        for i = 1:length(distances)
497            if (distances(i) >= G.step_length)
498                % Check if new node is closer
499                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
500                    % Update parent and distance
501                    parents(i) = ID_new;
502                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
503                end
504            end
505        end
506
507        % Local optimization stage
508        for i = 1:length(distances)
509            if (distances(i) >= G.step_length)
510                % Check if new node is closer
511                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
512                    % Update parent and distance
513                    parents(i) = ID_new;
514                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
515                end
516            end
517        end
518
519        % Prune stage
520        for i = 1:length(distances)
521            if (distances(i) >= G.step_length)
522                % Check if new node is closer
523                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
524                    % Update parent and distance
525                    parents(i) = ID_new;
526                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
527                end
528            end
529        end
530
531        % Local optimization stage
532        for i = 1:length(distances)
533            if (distances(i) >= G.step_length)
534                % Check if new node is closer
535                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
536                    % Update parent and distance
537                    parents(i) = ID_new;
538                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
539                end
540            end
541        end
542
543        % Prune stage
544        for i = 1:length(distances)
545            if (distances(i) >= G.step_length)
546                % Check if new node is closer
547                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
548                    % Update parent and distance
549                    parents(i) = ID_new;
550                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
551                end
552            end
553        end
554
555        % Local optimization stage
556        for i = 1:length(distances)
557            if (distances(i) >= G.step_length)
558                % Check if new node is closer
559                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
560                    % Update parent and distance
561                    parents(i) = ID_new;
562                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
563                end
564            end
565        end
566
567        % Prune stage
568        for i = 1:length(distances)
569            if (distances(i) >= G.step_length)
570                % Check if new node is closer
571                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
572                    % Update parent and distance
573                    parents(i) = ID_new;
574                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
575                end
576            end
577        end
578
579        % Local optimization stage
580        for i = 1:length(distances)
581            if (distances(i) >= G.step_length)
582                % Check if new node is closer
583                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
584                    % Update parent and distance
585                    parents(i) = ID_new;
586                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
587                end
588            end
589        end
590
591        % Prune stage
592        for i = 1:length(distances)
593            if (distances(i) >= G.step_length)
594                % Check if new node is closer
595                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
596                    % Update parent and distance
597                    parents(i) = ID_new;
598                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
599                end
600            end
601        end
602
603        % Local optimization stage
604        for i = 1:length(distances)
605            if (distances(i) >= G.step_length)
606                % Check if new node is closer
607                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
608                    % Update parent and distance
609                    parents(i) = ID_new;
610                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
611                end
612            end
613        end
614
615        % Prune stage
616        for i = 1:length(distances)
617            if (distances(i) >= G.step_length)
618                % Check if new node is closer
619                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
620                    % Update parent and distance
621                    parents(i) = ID_new;
622                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
623                end
624            end
625        end
626
627        % Local optimization stage
628        for i = 1:length(distances)
629            if (distances(i) >= G.step_length)
630                % Check if new node is closer
631                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
632                    % Update parent and distance
633                    parents(i) = ID_new;
634                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
635                end
636            end
637        end
638
639        % Prune stage
640        for i = 1:length(distances)
641            if (distances(i) >= G.step_length)
642                % Check if new node is closer
643                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
644                    % Update parent and distance
645                    parents(i) = ID_new;
646                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
647                end
648            end
649        end
650
651        % Local optimization stage
652        for i = 1:length(distances)
653            if (distances(i) >= G.step_length)
654                % Check if new node is closer
655                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
656                    % Update parent and distance
657                    parents(i) = ID_new;
658                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
659                end
660            end
661        end
662
663        % Prune stage
664        for i = 1:length(distances)
665            if (distances(i) >= G.step_length)
666                % Check if new node is closer
667                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
668                    % Update parent and distance
669                    parents(i) = ID_new;
670                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
671                end
672            end
673        end
674
675        % Local optimization stage
676        for i = 1:length(distances)
677            if (distances(i) >= G.step_length)
678                % Check if new node is closer
679                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
680                    % Update parent and distance
681                    parents(i) = ID_new;
682                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
683                end
684            end
685        end
686
687        % Prune stage
688        for i = 1:length(distances)
689            if (distances(i) >= G.step_length)
690                % Check if new node is closer
691                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
692                    % Update parent and distance
693                    parents(i) = ID_new;
694                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
695                end
696            end
697        end
698
699        % Local optimization stage
700        for i = 1:length(distances)
701            if (distances(i) >= G.step_length)
702                % Check if new node is closer
703                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
704                    % Update parent and distance
705                    parents(i) = ID_new;
706                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
707                end
708            end
709        end
710
711        % Prune stage
712        for i = 1:length(distances)
713            if (distances(i) >= G.step_length)
714                % Check if new node is closer
715                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
716                    % Update parent and distance
717                    parents(i) = ID_new;
718                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
719                end
720            end
721        end
722
723        % Local optimization stage
724        for i = 1:length(distances)
725            if (distances(i) >= G.step_length)
726                % Check if new node is closer
727                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
728                    % Update parent and distance
729                    parents(i) = ID_new;
730                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
731                end
732            end
733        end
734
735        % Prune stage
736        for i = 1:length(distances)
737            if (distances(i) >= G.step_length)
738                % Check if new node is closer
739                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
740                    % Update parent and distance
741                    parents(i) = ID_new;
742                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
743                end
744            end
745        end
746
747        % Local optimization stage
748        for i = 1:length(distances)
749            if (distances(i) >= G.step_length)
750                % Check if new node is closer
751                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
752                    % Update parent and distance
753                    parents(i) = ID_new;
754                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
755                end
756            end
757        end
758
759        % Prune stage
760        for i = 1:length(distances)
761            if (distances(i) >= G.step_length)
762                % Check if new node is closer
763                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
764                    % Update parent and distance
765                    parents(i) = ID_new;
766                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
767                end
768            end
769        end
770
771        % Local optimization stage
772        for i = 1:length(distances)
773            if (distances(i) >= G.step_length)
774                % Check if new node is closer
775                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
776                    % Update parent and distance
777                    parents(i) = ID_new;
778                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
779                end
780            end
781        end
782
783        % Prune stage
784        for i = 1:length(distances)
785            if (distances(i) >= G.step_length)
786                % Check if new node is closer
787                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
788                    % Update parent and distance
789                    parents(i) = ID_new;
790                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
791                end
792            end
793        end
794
795        % Local optimization stage
796        for i = 1:length(distances)
797            if (distances(i) >= G.step_length)
798                % Check if new node is closer
799                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
800                    % Update parent and distance
801                    parents(i) = ID_new;
802                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
803                end
804            end
805        end
806
807        % Prune stage
808        for i = 1:length(distances)
809            if (distances(i) >= G.step_length)
810                % Check if new node is closer
811                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
812                    % Update parent and distance
813                    parents(i) = ID_new;
814                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
815                end
816            end
817        end
818
819        % Local optimization stage
820        for i = 1:length(distances)
821            if (distances(i) >= G.step_length)
822                % Check if new node is closer
823                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
824                    % Update parent and distance
825                    parents(i) = ID_new;
826                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
827                end
828            end
829        end
830
831        % Prune stage
832        for i = 1:length(distances)
833            if (distances(i) >= G.step_length)
834                % Check if new node is closer
835                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
836                    % Update parent and distance
837                    parents(i) = ID_new;
838                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
839                end
840            end
841        end
842
843        % Local optimization stage
844        for i = 1:length(distances)
845            if (distances(i) >= G.step_length)
846                % Check if new node is closer
847                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
848                    % Update parent and distance
849                    parents(i) = ID_new;
850                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
851                end
852            end
853        end
854
855        % Prune stage
856        for i = 1:length(distances)
857            if (distances(i) >= G.step_length)
858                % Check if new node is closer
859                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
860                    % Update parent and distance
861                    parents(i) = ID_new;
862                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
863                end
864            end
865        end
866
867        % Local optimization stage
868        for i = 1:length(distances)
869            if (distances(i) >= G.step_length)
870                % Check if new node is closer
871                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
872                    % Update parent and distance
873                    parents(i) = ID_new;
874                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
875                end
876            end
877        end
878
879        % Prune stage
880        for i = 1:length(distances)
881            if (distances(i) >= G.step_length)
882                % Check if new node is closer
883                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
884                    % Update parent and distance
885                    parents(i) = ID_new;
886                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
887                end
888            end
889        end
890
891        % Local optimization stage
892        for i = 1:length(distances)
893            if (distances(i) >= G.step_length)
894                % Check if new node is closer
895                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
896                    % Update parent and distance
897                    parents(i) = ID_new;
898                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
899                end
900            end
901        end
902
903        % Prune stage
904        for i = 1:length(distances)
905            if (distances(i) >= G.step_length)
906                % Check if new node is closer
907                if (q_new - G.nodes(i).state).^2 <= G.step_length.^2)
908                    % Update parent and distance
909                    parents(i) = ID_new;
910                    distances(i) = sqrt((q_new - G.nodes(i).state).^2);
911                end
912            end
913        end
914
```

```

49     IDs_near = G.getNearbyIDs(q_new, G.rewiring_radius)';
50     % IDs_near = G.getNearbyIDs(q_new, G.rewiring_radius*(log(numel(G.nodes)) /
51     numel(G.nodes))^(1/(ndims(map)+1)) )';
52
53     % Parent choice & edge creation
54     distances(ID_new, 1) = +inf;
55     for ID_near = IDs_near
56         q_near = G.getNodeByID(ID_near).state;
57         dist = distances(ID_near) + Node.euclideanDistance(q_near, q_new);
58         if (dist < distances(ID_new, 1) && G.isValidConnection(q_near, q_new, map))
59             distances(ID_new, 1) = dist;
60             parents(ID_new, 1) = ID_near;
61         end
62     end
63     node_parent = G.getNodeByID(parents(ID_new));
64     G.addEdge(node_parent.ID, ID_new, Node.euclideanDistance(node_parent.state,
65     q_new));
66
67     % Rewiring stage
68     for ID_near = IDs_near
69         q_near = G.getNodeByID(ID_near).state;
70         dist = distances(ID_new) + Node.euclideanDistance(q_near, q_new);
71         if (dist < distances(ID_near) && G.isValidConnection(q_new, q_near, map))
72             edge = G.getEdgeByConnection(parents(ID_near), ID_near);
73             G.edges(edge.ID).ID_A = ID_new;
74             G.edges(edge.ID).ID_B = ID_near;
75             G.edges(edge.ID).weight = Node.euclideanDistance(q_new, q_near);
76             distances(ID_near) = dist;
77             parents(ID_near) = ID_new;
78         end
79     end
80
81     iter = iter + 1;
82
83 end
84
85 exist = ~(iter == G.max_iter);
86
87 % Connection with final node
88 if (exist)
89     ID_final = G.addNode(node_final.state);
90     IDs_near = G.getNearbyIDs(node_final.state);
91     ID_near = IDs_near(1);
92     q_near = G.getNodeByID(ID_near).state;
93     dist = Node.euclideanDistance(q_near, node_final.state);
94     G.addEdge(ID_near, ID_final, dist);
95     parents(ID_final, 1) = ID_near;
96     distances(ID_final, 1) = distances(ID_near) + dist;
97 end
98
99 end

```

Listing 2: Code for the implementation of RRT* algorithm.

Again, the core structure of the algorithm is similar to the RRT algorithm, but with a few extra steps. The algorithm starts by initializing the tree T with the start node and enters a loop where it samples a random point in the configuration space. Then, it finds the nearest node in the tree and generates a new node by moving towards the sampled point. If the new node is in collision with the obstacles, it is discarded and the algorithm continues to sample a new point. If the new node is valid, the algorithm at first connect it to the nearest node in the tree, and then checks if any of its neighbors can also be connected to it with a lower cost-to-go. This

analysis, called rewiring, help to construct a more optimal tree with smoother connections and consequently shorter paths. Figure 2 shows an example of the rewiring process.

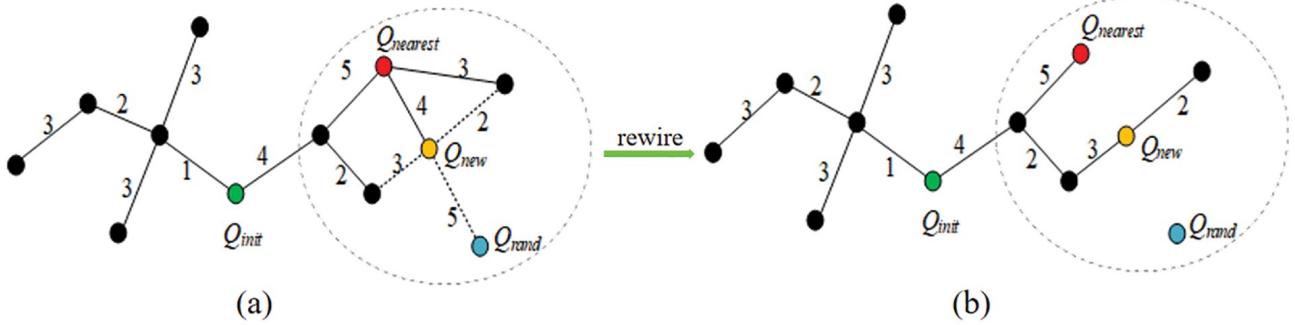


Figure 2: Example of the rewiring process in RRT* algorithm. Credit to [1].

As it is going to be shown in the next section, the RRT* algorithm is more efficient than the RRT algorithm in terms of path length, at the cost negligible increase in computation time.

Optimality convergence of RRT* As a final note, we recall that thanks to [2], it has been proved that the RRT* algorithm converges to the optimal solution as the number of sampled points goes to infinity and the rewiring radius is dynamically adjusted as:

$$r_n = \gamma \left(\frac{\log n}{n} \right)^{\frac{1}{d+1}} \quad (1)$$

Where n is the number of sampled points, d is the dimension of the configuration space, and γ is a constant that depends on d .

Line 50 of the code in Listing 2 shows a possible implementation of this dynamic rewiring radius.

3.3 RRT-kinematic algorithm

Finally, we move to the last algorithm we are going to implement, which is the RRT-kinematic algorithm. The RRT-kinematic algorithm is a variant of the RRT algorithm that takes into account the kinematic constraints of the robot during the planning process. This is particularly useful for vehicles with non-holonomic constraints, such as terrestrial vehicles or mobile robots.

Listing 3 shows a possible implementation of the RRT-kinematic algorithm.

```

1 function [exist, distances, parents] = rrtKinematic(G, map, node_initial, node_final,
2   vehicle_model)
3 %
4 % RRTKINEMATIC Builds a tree using RRT with vehicle kinematic constraints.
5 %
6 %
7 %
8 %
9 % Inputs:
10 %   G           - Graph object to store the RRT tree.
11 %   MAP         - Binary occupancy map (1 = free, 0 = obstacle).
12 %   NODE_INITIAL - Starting node with state [x, y, theta].
13 %   NODE_FINAL   - Goal node with target state [x, y, theta].
14 %   VEHICLE_MODEL - Function handle simulating vehicle motion over time.
15 %
16 %
17 %
18 %
19 %
20 % Outputs:
21 %   EXIST      - Boolean indicating whether a valid path was found.
22 %   DISTANCES  - Cumulative cost from the start node to each tree node.
23 %   PARENTS    - Parent indices used to reconstruct the tree structure.
24 %
25 % This version of RRT uses motion primitives derived from a forward vehicle model
26 % to ensure that the generated paths respect the system kinematic constraints.

```

```

22
23 v_options = linspace(-1, 1, 2);
24 v_options = 1;
25 delta_options = linspace(-pi/6, pi/6, 3);
26
27 q_new = node_initial.state;
28 distances = zeros(1, 1);
29 parents = NaN(1, 1);
30 iter = 0;
31
32 while (Node.euclideanDistance(q_new(1:2), node_final.state(1:2)) > G.step_length && iter
33 < G.max_iter)
34
35 % Sampling stage
36 q_rand = [max(rand(1, 2) .* size(map), 1), rand() * 2 * pi];
37
38 % Finding the nearest node
39 IDs_near = G.getNearbyIDs(q_rand);
40 ID_near = IDs_near(1);
41 q_near = G.getNodeByID(ID_near).state;
42
43 % Kinematics constraints
44 best_state = [];
45 min_dist = Inf;
46 for v = v_options
47     for delta = delta_options
48
49         % Forward integration of vehicle motion
50         q_new = vehicle_model(q_near, [v, delta], G.step_length / abs(v));
51
52         % Collision checking and choice of nearest state to sampled one
53         if (G.isValidConnection(q_near, q_new, map))
54             dist = Node.euclideanDistance(q_new(1:2), q_rand(1:2));
55             if (dist < min_dist)
56                 best_state = q_new;
57                 min_dist = dist;
58             end
59         end
60     end
61 end
62
63 if isempty(best_state)
64     continue;
65 end
66
67 % Tree update
68 q_new = best_state;
69 ID_new = G.addNode(q_new);
70 dist = Node.euclideanDistance(q_near(1:2), q_new(1:2));
71 G.addEdge(ID_near, ID_new, dist);
72
73 parents(ID_new, 1) = ID_near;
74 distances(ID_new, 1) = distances(ID_near) + dist;
75
76 iter = iter + 1;
77
78 end
79
80 exist = ~(iter == G.max_iter);
81

```

```

82 % Connection with final node
83 if (exist)
84     ID_final = G.addNode(node_final.state);
85     IDs_near = G.getNearbyIDs(node_final.state);
86     ID_near = IDs_near(1);
87     q_near = G.getNodeByID(ID_near).state;
88     dist = Node.euclideanDistance(q_near(1:2), node_final.state(1:2));
89     G.addEdge(ID_near, ID_final, dist);
90     parents(ID_final, 1) = ID_near;
91     distances(ID_final, 1) = distances(ID_near) + dist;
92 end
93
94 end

```

Listing 3: Code for the implementation of RRT-kinematic algorithm.

Again, we recognize the same core structure of the RRT algorithm, but with a few extra steps. In particular, the choice of the new node is not done by simply moving towards the sampled point, but rather by generating a trajectory from the nearest node in the tree to the sampled point leveraging the kinematic model of the robot. In the specific case of the implemented algorithm, we consider dealing with a *Bicycle* model, which is a common model used to represent the motion of a car-like vehicle. Starting from the nearest node in the tree, a random control input is sampled and a trajectory is generated by integrating the kinematic model of the robot over a fixed time horizon. Collision checking is then performed and among all the tested trajectories, the one that is valid and has the lowest cost is selected as the new node to be added to the tree. By doing so, we are sure that the generated path is feasible by the robot, respecting its kinematic constraints.

Bicycle kinematic model Even if the kinematic model of the vehicle is not the main focus of this report, we briefly present it here for completeness. The kinematic model of the vehicle is described by the following equations:

$$\begin{aligned}\dot{x} &= v \cdot \cos(\theta) \\ \dot{y} &= v \cdot \sin(\theta) \\ \dot{\theta} &= \frac{v}{L} \cdot \tan(\delta)\end{aligned}\quad (2)$$

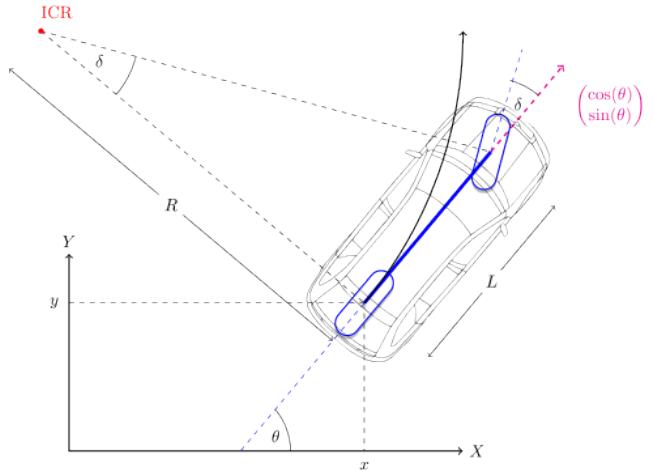


Figure 3: Bicycle kinematic model. Image credit to Thomas Fermi.

Where x and y are the position of the vehicle in the workspace, θ is the heading angle of the vehicle, v is the linear velocity, L is the wheelbase of the vehicle (distance between the front and rear axles), and δ is the steering angle.

One can clearly see that the kinematic model is composed by a set of differential equations. In order to simulate in time based on the control inputs and current state of the vehicle, we need to integrate these equations. Different integration methods can be used, such as Euler integration or Runge-Kutta integration. However, given the simplicity and expected slow dynamics of the model, we can use a simple Euler integration with a fixed time step Δt .

Listing 4 shows a possible implementation of the kinematic model of the vehicle.

```

1 function q_new = modelBicycle(q, u, T)
2 % MODELBICYCLE Simulates bicycle model kinematics over time T.
3 %
4 % Q_NEW = MODELBICYCLE(Q, U, T) updates the state of a vehicle using the

```

```

5 % kinematic bicycle model given the current state Q, control input U, and
6 % time duration T.
7 %
8 % Inputs:
9 %   Q - Current state vector [x, y, theta], where:
10 %     x, y - Position of the vehicle in 2D space
11 %     theta - Heading angle (radians)
12 %
13 %   U - Control input vector [speed, steering_angle], where:
14 %     speed - Forward velocity
15 %     steering_angle - Steering angle (radians)
16 %
17 %   T - Duration over which to apply the control input (seconds)
18 %
19 % Output:
20 %   Q_NEW - New state vector [x, y, theta] after time T.
21
22 % States
23 x = q(1);
24 y = q(2);
25 theta = q(3);
26
27 % Inputs
28 speed = u(1);
29 steering_angle = u(2);
30
31 % Parameters
32 wheelbase = 2.0;
33
34 time = linspace(0, T, 100);
35 dt = diff(time(1:2));
36 for t = time
37   x = x + speed * cos(theta) * dt;
38   y = y + speed * sin(theta) * dt;
39   theta = theta + (speed / wheelbase) * tan(steering_angle) * dt;
40 end
41
42 q_new = [x y theta];
43
44 end

```

Listing 4: Code for the kinematic model of the vehicle.

4 Algorithm Testing

We can now proceed to test the algorithms we have implemented in the previous section. In the following, the maps scenarios adopted are presented first, and then the results of the tests are shown.

4.1 Map scenarios

Figures 4 show the four image maps used to test the algorithm. Even if the images here below are of size 300x300, we preferred to scale down to 30x30 in order to have a direct comparison between the results obtained previously with the grid-based algorithms. Notice also that while the white areas are free space, the one marked in black or gray are to be considered as obstacles. Finally, the green dot is the starting point of the vehicle, while the red one is the goal.

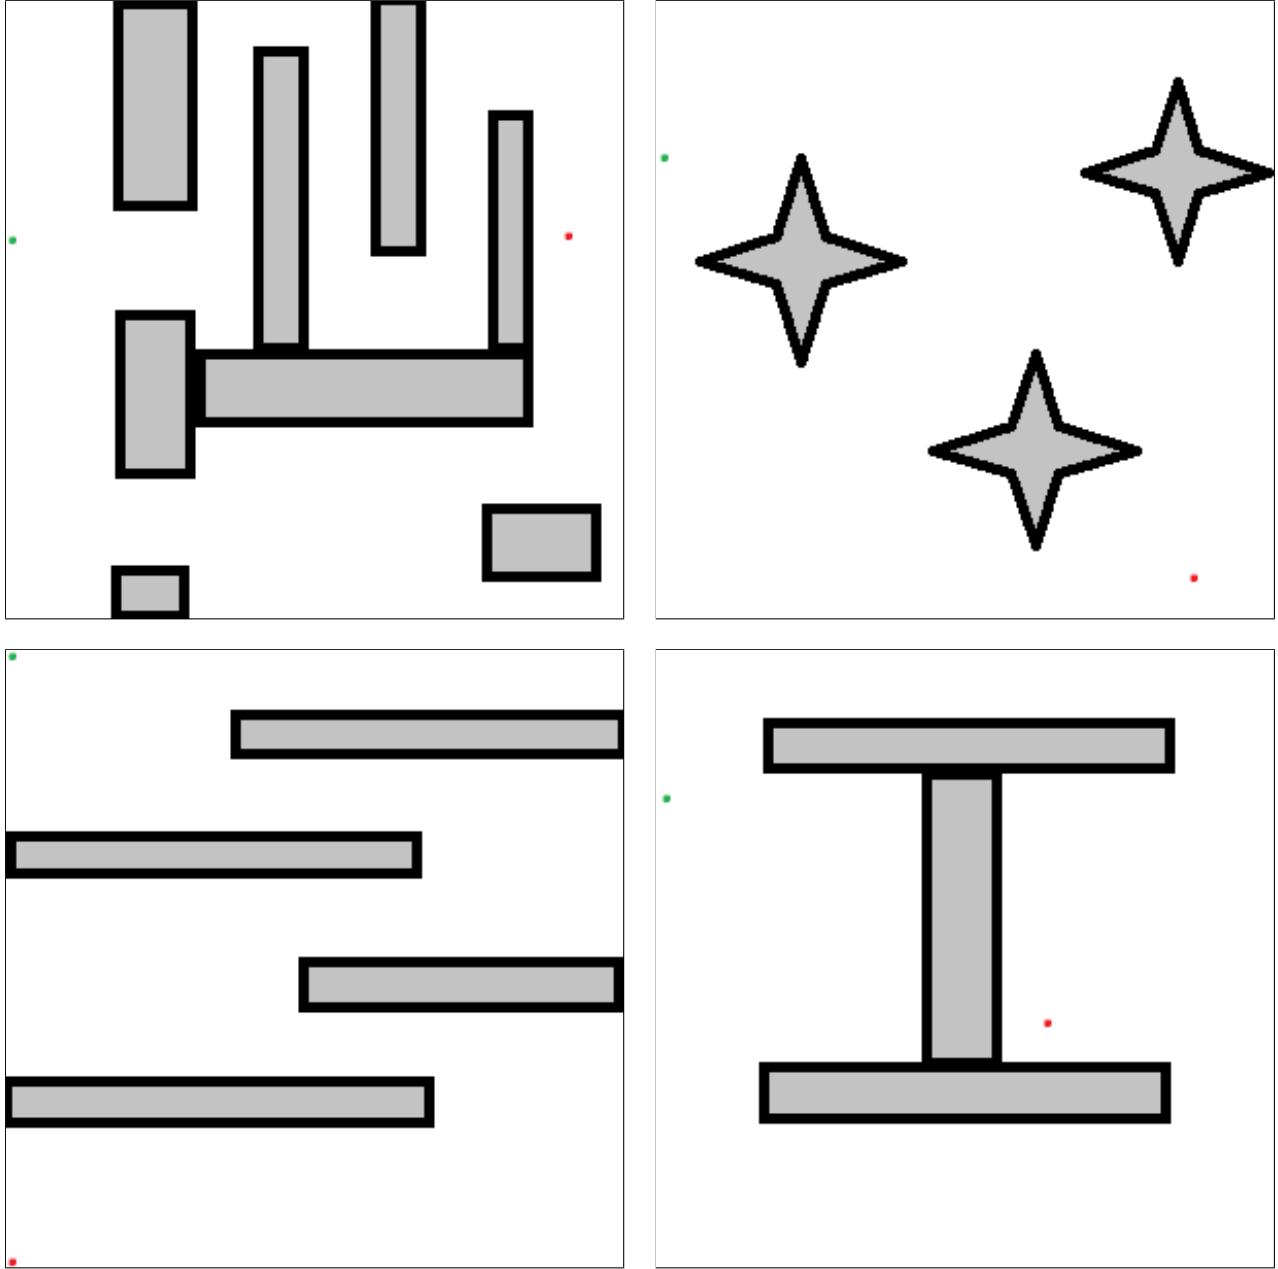


Figure 4: Test maps used for the algorithm testing.

4.2 Results analysis

In the following, results obtained running both RRT, RRT* and RRT kinematic algorithms on each of the scenario are presented. For comparison purposes, the results obtained running the A* algorithm on the same maps are also reported, even if an in-depth analysis of the differences and performances of the two families of algorithms is left to the next section (Section 5).

The plotting of both the (geometrical) connected graph and the best path found is also shown. Tables presenting the main metrics of the tests are also shown. The metrics are the following:

- **Elapsed time:** here only the time needed to run the searching algorithm is considered (i.e. we are including the tree generation but excluding the plotting of the results);
- **Tree dimension:** number of nodes in the tree generated by the algorithm. In case of the A* algorithm, this is the number of nodes in the graph generated during the discretization/grid generation phase;
- **Path length:** length of the path found by the algorithm. This also correspond to the path cost, given that the cost of each edge is equal to its geometrical length.

Notice that the due to the randomness nature of sampling-based algorithms, the reported results must be

interpreted as an average of 5 runs of each algorithm on each map. Moreover, the tuning parameters of the algorithms such as step size and rewiring radius have been set on all the algorithms variants to be:

| Parameter | Value |
|------------------|-------|
| Step size | 1.0 |
| Rewiring radius* | 1.5 |

Table 1: Tuning parameters of the algorithms. The rewiring radius is only used in the RRT* algorithm.

Units are considered to be in pixels dimensions, where we recall the size of the map to be 30x30.

4.2.1 Map 1

Results obtained running the three algorithms on the first map are shown in Figures 5 and Table 2. The path found by the algorithm is shown in red, while the tree/graph is shown in black.

Notice that the fourth image (bottom right) is actually the results obtained running the A* algorithm, and it is shown here just for comparison purposes.

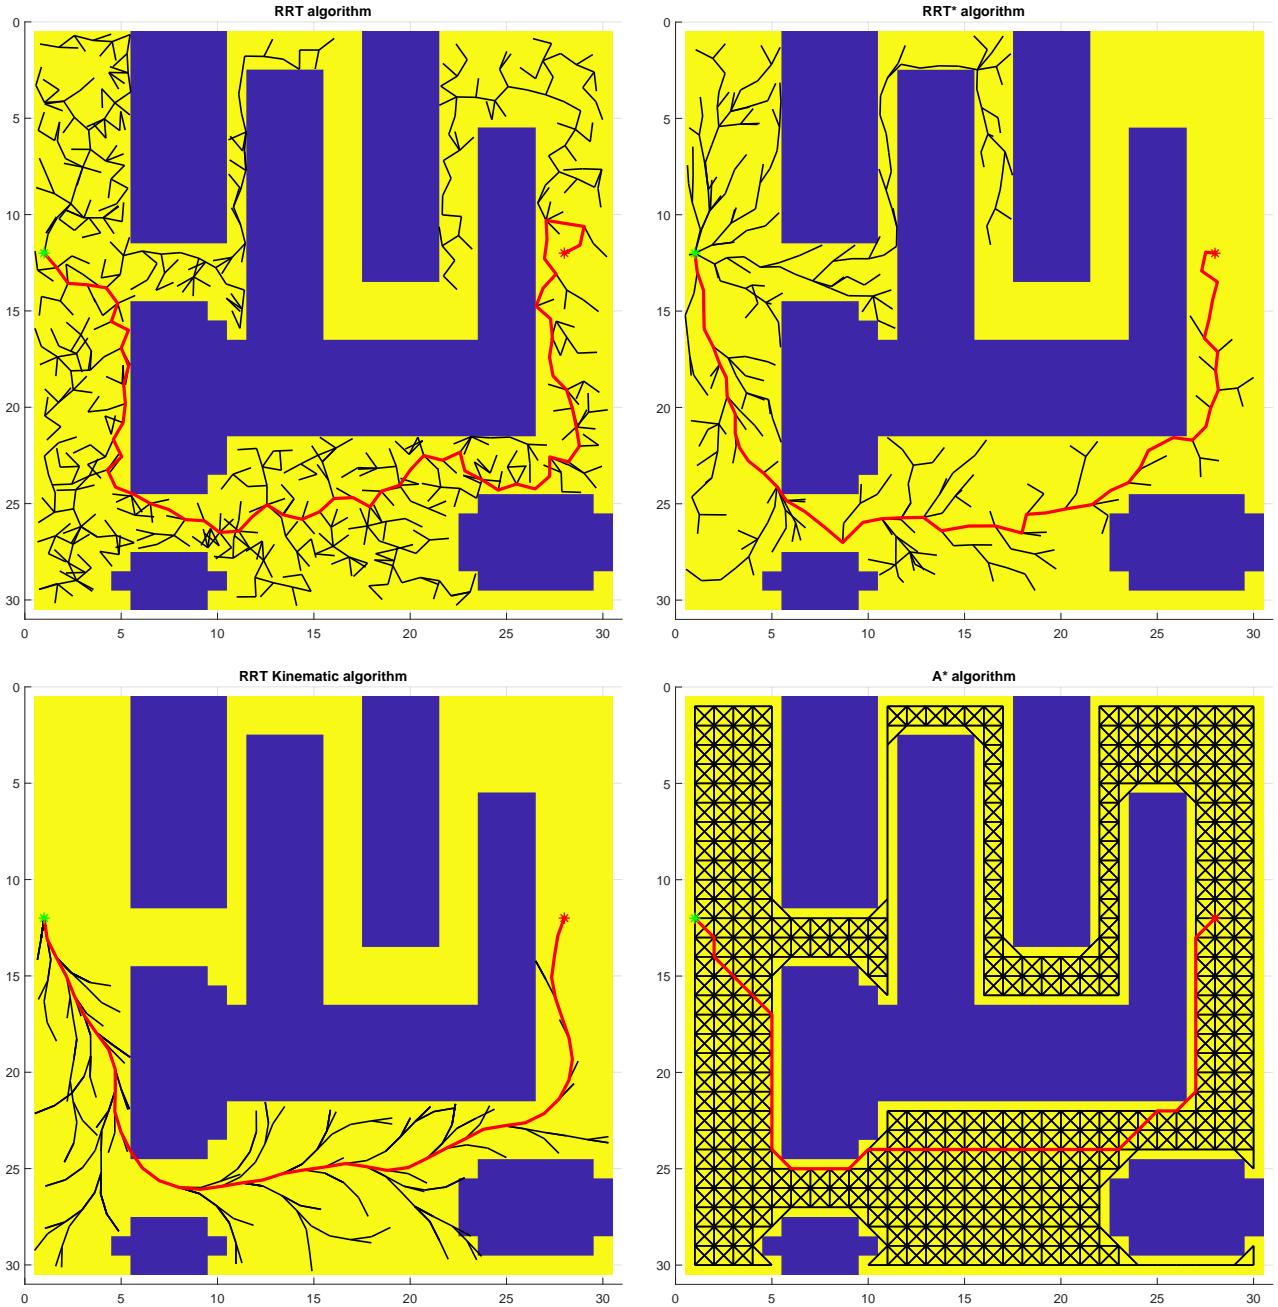


Figure 5: Tree generated and path found by the RRT, RRT* and RRT Kinematic algorithms on map 1.

| Algorithm | Elapsed time (ms) | Tree dimension | Path length (m) |
|---------------|-------------------|----------------|-----------------|
| RRT | 145 | 338 | 57 |
| RRT* | 202 | 271 | 49 |
| RRT Kinematic | 129 | 457 | 49 |
| A* | 1987 | 3332 | 47 |

Table 2: Results of the algorithms on map 1.

The results clearly show the core problematics of the sampling-based algorithms: the path found is not optimal. Comparing the results of the RRT and RRT* algorithms, we can see that the latter is able to find a shorter path, but it's still not able to reach the lower limit which is the one found by the A* algorithm. On the other hand, sampling-based algorithms demonstrate to be in general 10x faster than grid-based algorithms, as we can see from the elapsed time. Also, under the space complexity point of view, the search space is significantly smaller for the sampling-based algorithms. This suggests that the sampling-based algorithms are more efficient in terms of space complexity, as they do not need to explore the whole workspace to find a feasible path.

4.2.2 Map 2

Results obtained running the three algorithms on the second map are shown in Figures 6 and Table 3. The path found by the algorithm is shown in red, while the tree/graph is shown in black. Again, we have included the results of the A* algorithm for comparison purposes.

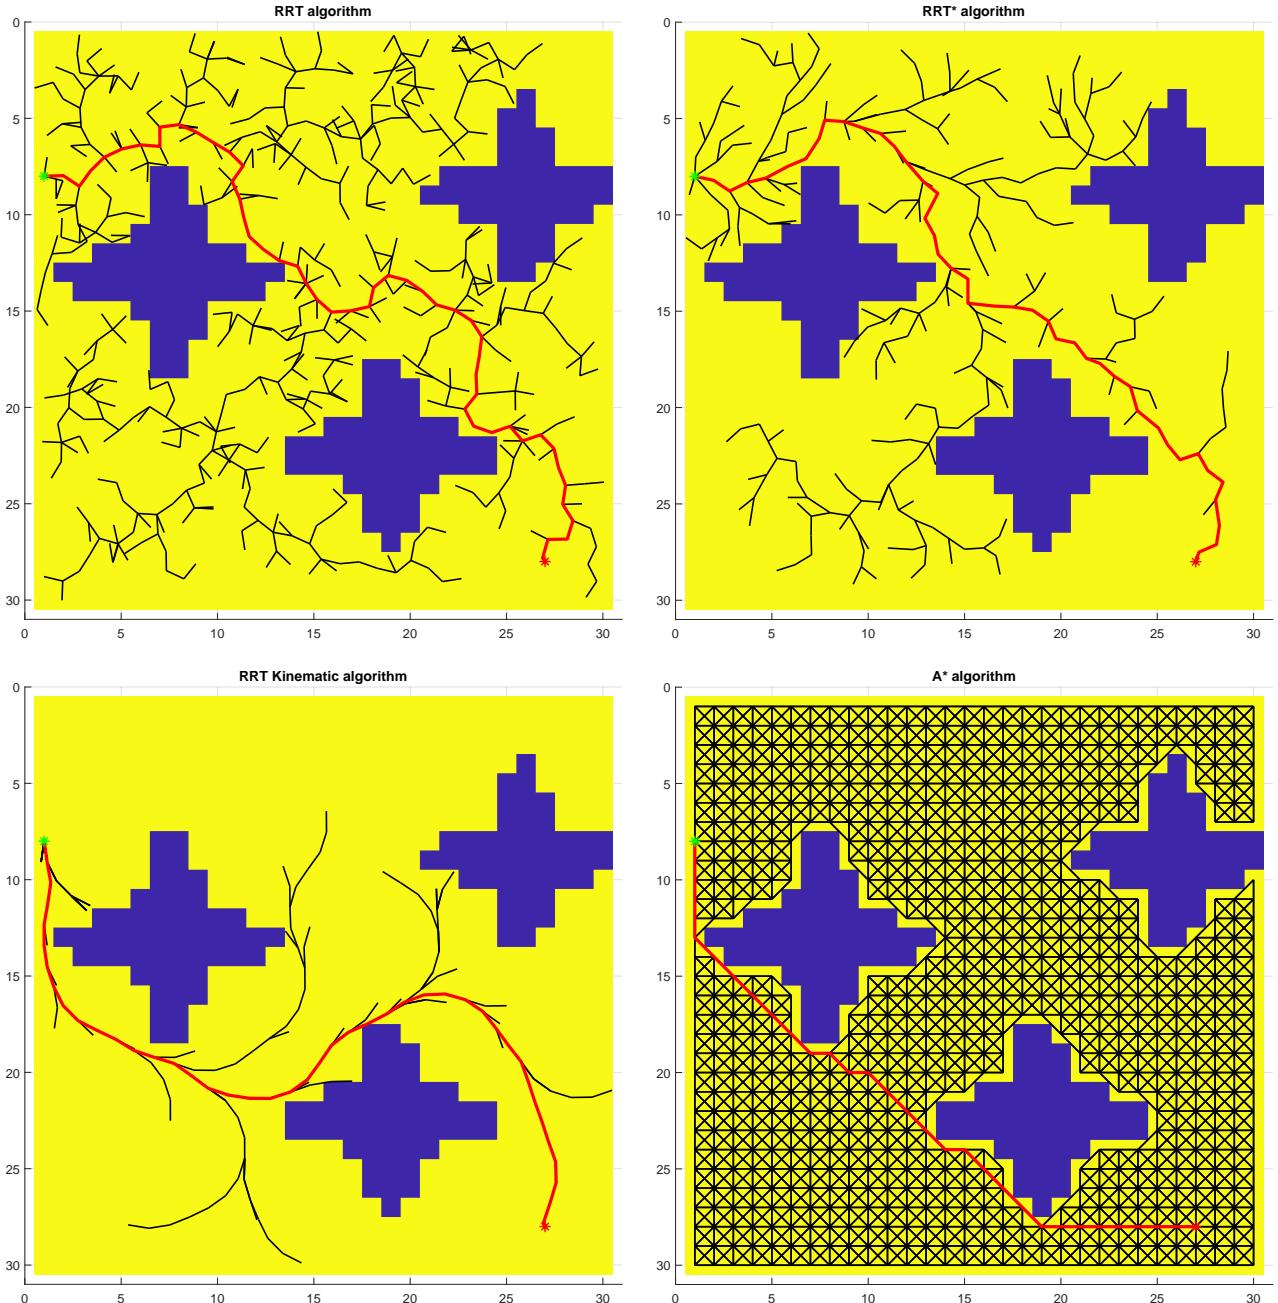


Figure 6: Tree generated and path found by the RRT, RRT* and RRT Kinematic algorithms on map 2.

| Algorithm | Elapsed time (ms) | Tree dimension | Path length (m) |
|---------------|-------------------|----------------|-----------------|
| RRT | 199 | 581 | 49 |
| RRT* | 360 | 383 | 40 |
| RRT Kinematic | 282 | 615 | 42 |
| A* | 2025 | 5398 | 37 |

Table 3: Results of the algorithms on map 2.

Once again, we observe that the path found by the sampling-based algorithms is suboptimal. Similar to before, the RRT* algorithm is the one able to get the closest to the optimal path found by the A* algorithm but still, it is not able to reach it. As for the elapsed time, we still recognize a factor of around 8 in speedup of the sampling-based algorithms with respect to the grid-based ones. We also recall that this particular map and start/goal configuration is the one where the A* algorithm was able to find the optimal path in the shortest time, thanks to the heuristic adopted and the lower number of possible dead-ends-street during the search. We believe that the speed-up factor might be even higher in more complex maps, where the A* algorithm would have to explore a larger number of nodes before finding the optimal path.

4.2.3 Map 3

Results obtained running the three algorithms on the third map are shown in Figures 7 and Table 4. The path found by the algorithm is shown in red, while the tree/graph is shown in black. Again, we have included the results of the A* algorithm for comparison purposes.

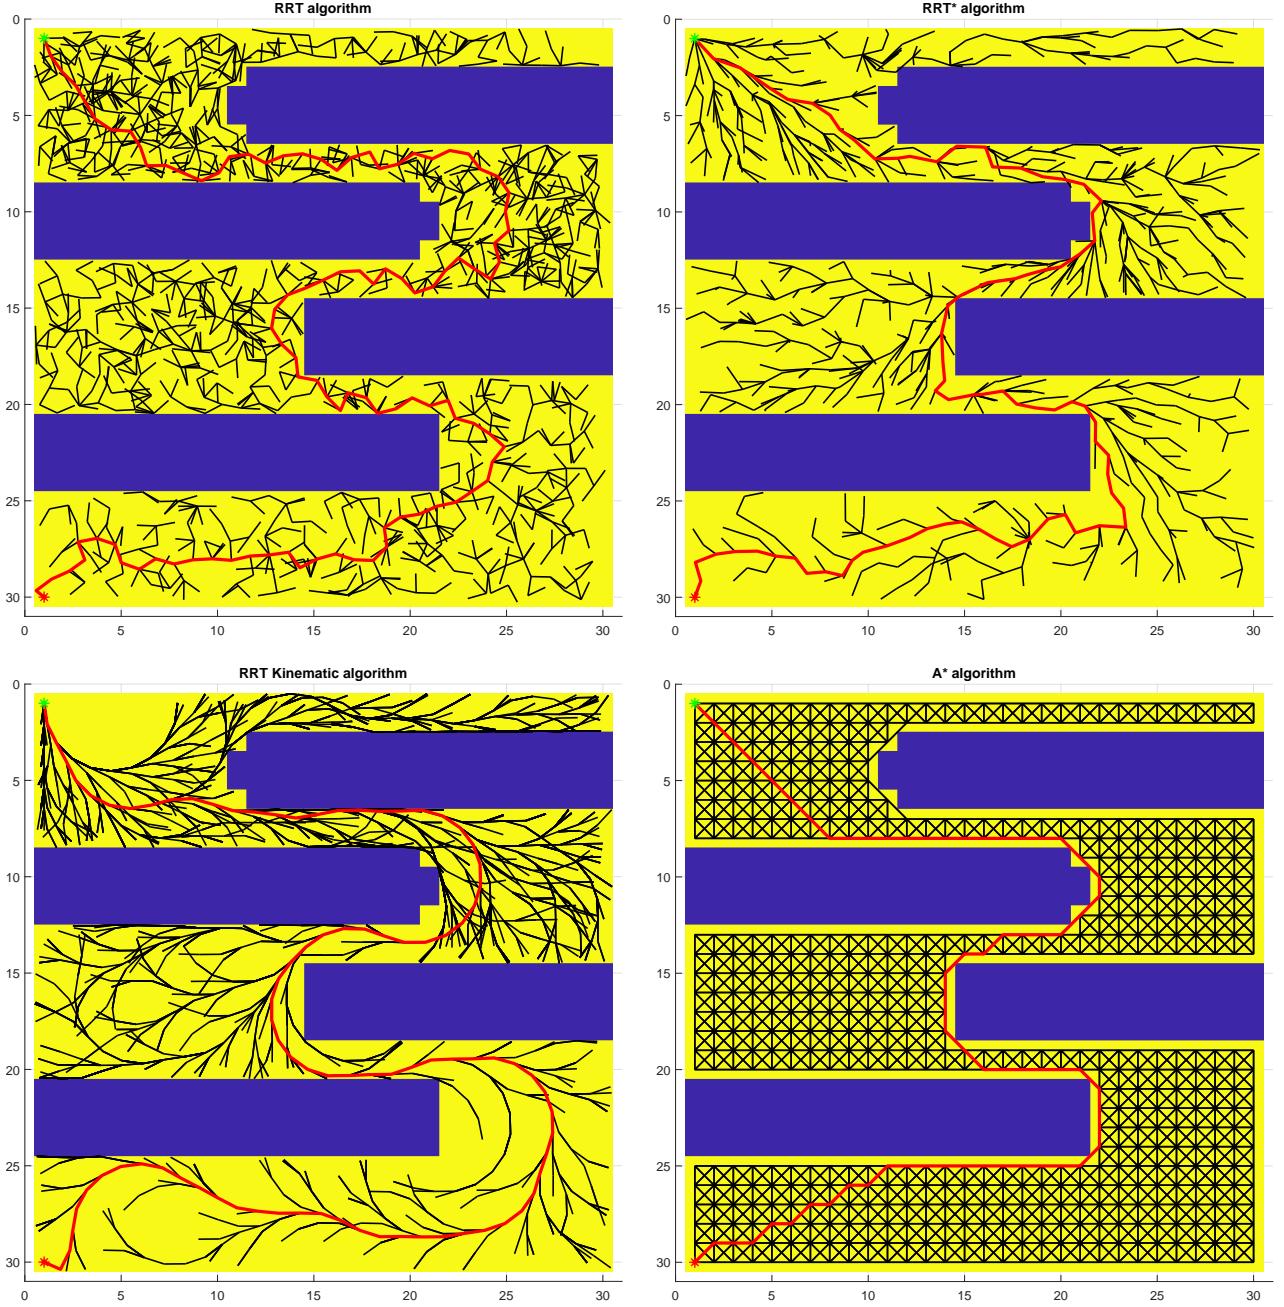


Figure 7: Tree generated and path found by the RRT, RRT* and RRT Kinematic algorithms on map 3.

| Algorithm | Elapsed time (ms) | Tree dimension | Path length (m) |
|---------------|-------------------|----------------|-----------------|
| RRT | 552 | 1085 | 97 |
| RRT* | 1208 | 932 | 74 |
| RRT Kinematic | 68601 | 10338 | 93 |
| A* | 2216 | 3936 | 74 |

Table 4: Results of the algorithms on map 3.

From the test performed on this map, we gain further information about the performance of the algorithms in cluttered environments. In this case, the difference in computation time between the sampling-based algorithms and the A* algorithm is not as significant as in the previous cases. The sampling approach is in fact heavily

slowed down by the four obstacles, that impose the need of numerous samples to progress in the successive sectors of the map up to the goal. This is also reflected in the distribution of the nodes in the tree, which is much denser in the upper sectors of the map (near the starting point) than in the lower ones (near the goal). RRT-kinematic clearly shows this behavior, highlighting the trapping of the algorithm in front of each obstacle. A* algorithm, on the other hand, still suffers from local dead-ends searching, but it's much quicker in finding a way to get around the obstacles.

For completeness, we also report that in the case of the RRT-kinematic algorithm, high variance in the results was observed. In particular, among the runs performed, we have one case where the algorithm was able to find a path in 8.6 seconds, while in the worse case it took 115.7 seconds.

4.2.4 Map 4

Results obtained running the three algorithms on the fourth map are shown in Figures 8 and Table 5. The path found by the algorithm is shown in red, while the tree/graph is shown in black. Again, we have included the results of the A* algorithm for comparison purposes.

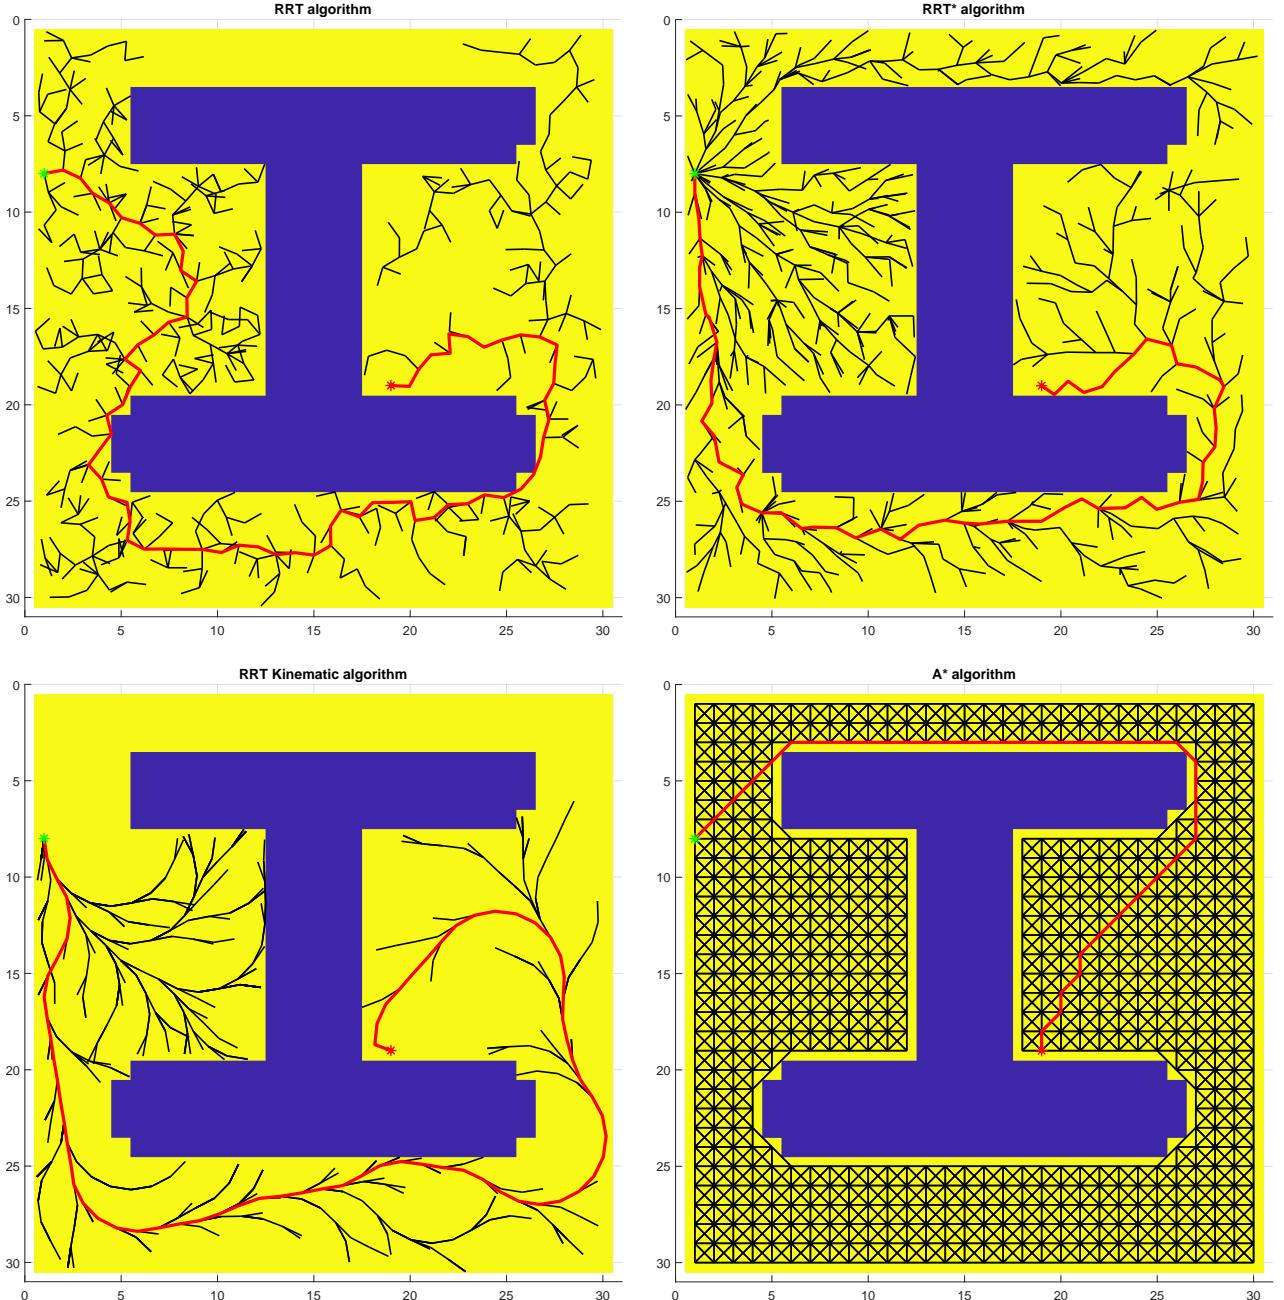


Figure 8: Tree generated and path found by the RRT, RRT*, and RRT Kinematic algorithms on map 4.

| Algorithm | Elapsed time (ms) | Tree dimension | Path length (m) |
|---------------|-------------------|----------------|-----------------|
| RRT | 164 | 604 | 67 |
| RRT* | 485 | 618 | 45 |
| RRT Kinematic | 782 | 1649 | 64 |
| A* | 2607 | 4520 | 47 |

Table 5: Results of the algorithms on map 4.

The results obtained on this map are similar to the ones obtained on the previous ones. It's interesting to notice how the RRT was not able to explore the upper region of the map, in which instead the A* finds its optimal path. Once again, the concept of obstacles blocking the path of the sampling-based algorithms is highlighted similarly to the previous case.

5 Grid vs Sample Based Planning

So far we have limited our analysis and comparisons between the grid-based and the sample-based algorithms to a fixed set of predefined scenarios. In this section, we further investigate the performance and capabilities of the two classes of algorithms by running specific tests that are designed to highlight their strengths and weaknesses. Thanks to the flexibility provided by the rewiring mechanism of RRT*, in this section we are going to compare the performance of the A* against the RRT* algorithm.

5.1 Path Asymptotic Optimality

As already largely discussed in the previous sections, the main difference between grid-based and sample-based algorithms is that the former are guaranteed to find the optimal path, while the latter are not. Formally, this statement has also been expressed in the literature by [2], which has proven that RRT can produce arbitrarily bad paths with non-negligible probability.

We can also observe this behavior in our implementation of the RRT* algorithm. Figure 9 shows the path generated by RRT* on the left and the path generated by A* on the right.

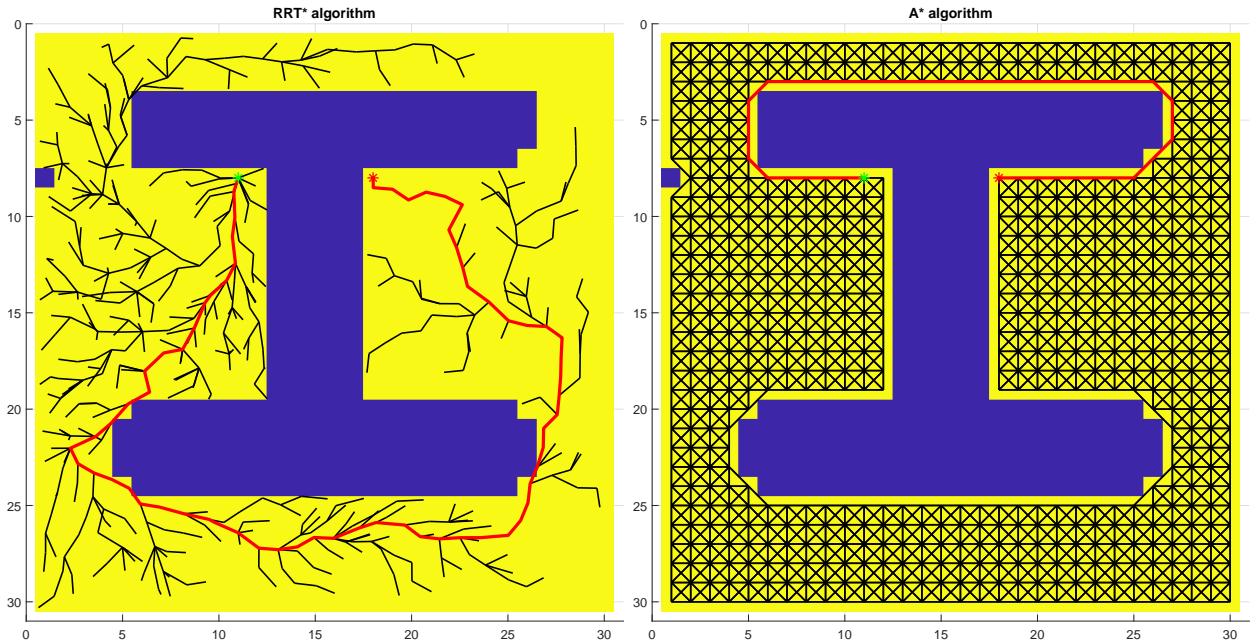


Figure 9: Path generated by RRT* (left) and A* (right) on the same scenario.

The path generated by RRT* is clearly suboptimal, choosing a much longer path only because fewer obstacles are present in that area and so higher chances of connectivity for the randomly generated nodes are present. A similar case is shown in Figure 10.

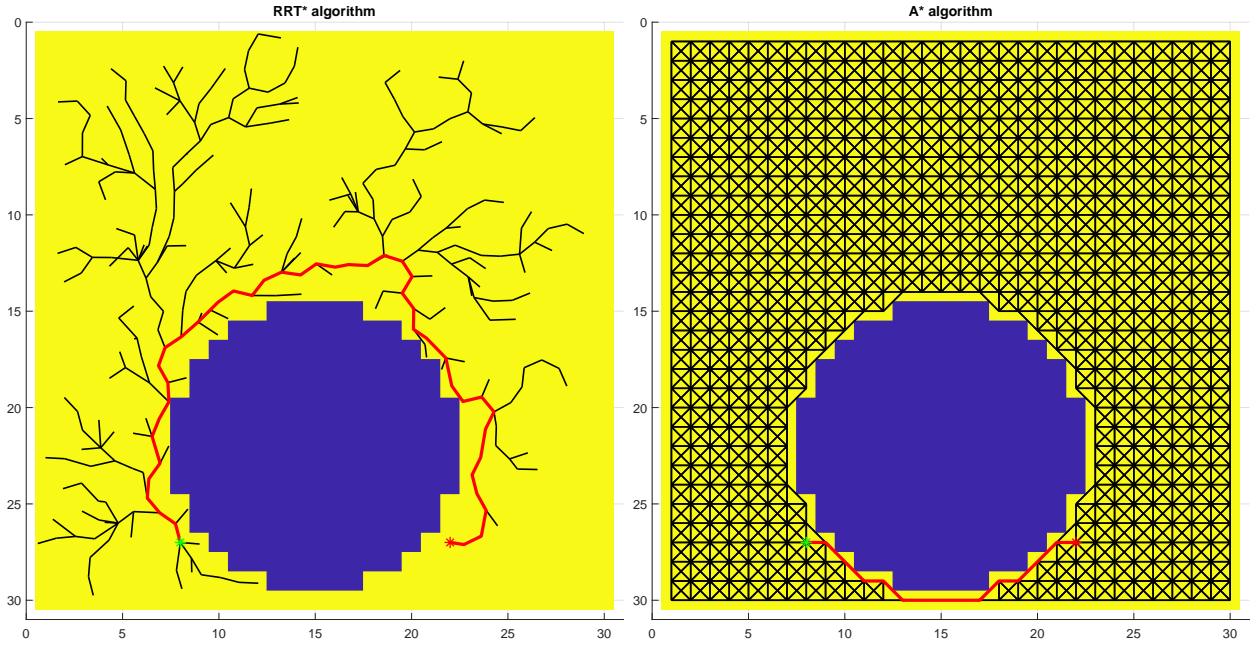


Figure 10: Path generated by RRT* (left) and A* (right) on the same scenario.

However, the RRT* algorithm is asymptotically optimal, meaning that the higher the number of nodes that the algorithm place, the more optimal the path becomes. This is a very important property, as it allows trading off between the time spent in finding a solution and the quality of the solution itself.

By removing the interruption of the algorithm after the first solution is found, we can observe how the path generated by RRT* becomes more and more optimal as the number of iterations increases. Figure 11 shows the path generated by RRT* after 1000 iterations on two different scenarios. One can appreciate how, despite the number of iterations, in the case of the circle the path is still suboptimal.

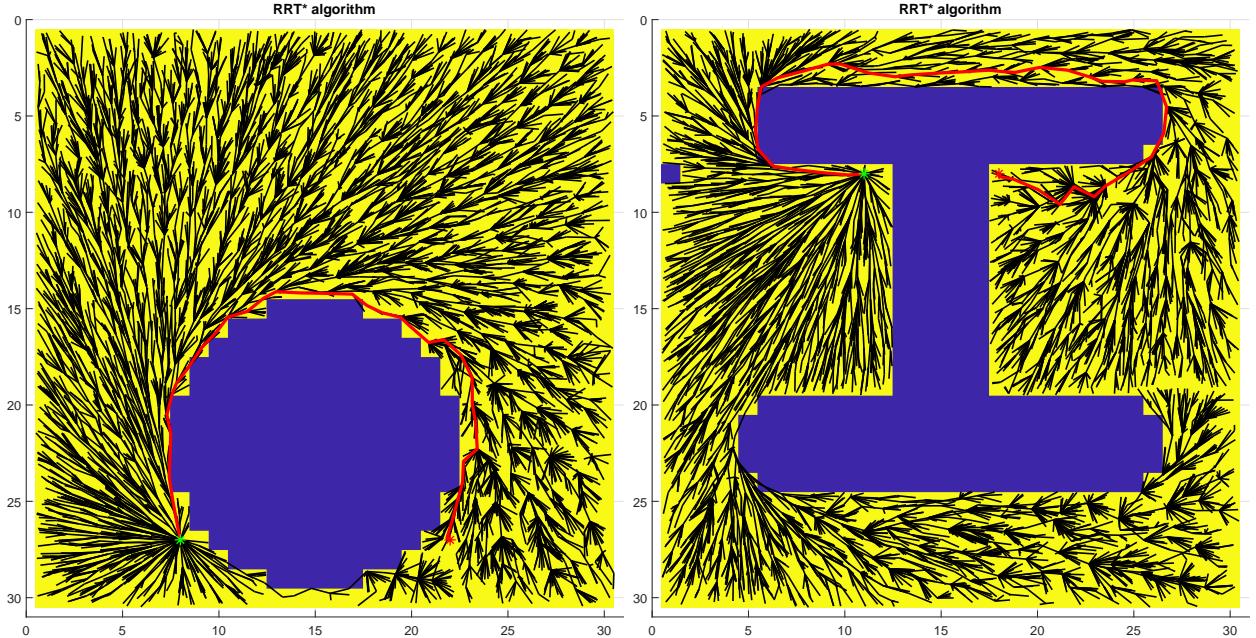


Figure 11: Asymptotic optimality of RRT* algorithm. Both the images refer to 10000 iterations of the algorithm.

In Table 6 are reported the path length and the time taken by RRT* algorithm as a function of the number of iterations performed. Results are relative to the right-hand side scenario of Figure 11.

| Iterations | Path length (m) | Elapsed time (ms) |
|------------|-----------------|-------------------|
| 1000 | 68 | 324 |
| 5000 | 60 | 12775 |
| 10000 | 44 | 69069 |

Table 6: Path length and time taken by RRT* algorithm after different number of iterations.

As expected, the path length decreases as the number of iterations increases, but the time taken by the algorithm increases significantly. Thinking about deploying the algorithm on a real robot which needs the planner component to run in real-time or nearly real-time, we clearly understand that the number of iterations must be limited, and asymptotic optimality is not a property that can be exploited in practice.

5.2 Multidimensional Workspace

In this section we analyze the performance of the two algorithms on a 3D scenario, adding a third dimension to the problem. The scenario is composed of a 3D grid of obstacles, which are randomly generated in the space. Figure 12 shows a 3D view of the scenario and a stack of slicing along the Z axis.

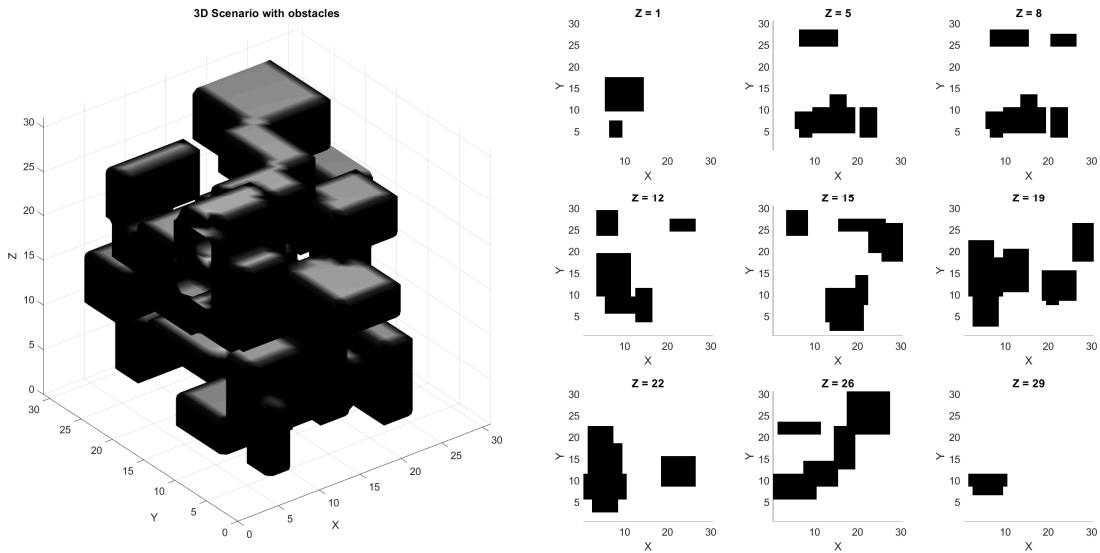


Figure 12: 3D grid of obstacles used as a scenario for the tests.

With zero-to-little modifications to the algorithms, we run both the grid-based and the sample-based algorithms on the same 3D scenario. Thanks to the increased dimensions of the problem, we clearly observe the huge difference in performance between the two class of algorithms. Table 7 summarizes the results of the tests, while Figure 13 shows the paths generated by A*, RRT and RRT* algorithms.

| Algorithm | Elapsed time (ms) | Tree dimension | Path length (m) |
|-----------|----------------------------|----------------|-----------------|
| A* | 2435229 (~ 40 [min]) | 20803 | 87 |
| RRT | 55 | 191 | 69 |
| RRT* | 273 | 397 | 69 |

Table 7: Path length and time taken by each algorithm on the same 3D scenario.

One can notice that the path generated by A* is longer than the one generated by RRT and RRT*. This can be explained by the fact that the graph used to represent the workspace has been generated without allowing diagonal connections. This means that, considering a 3D unitary volume of the grid, for the A* to go from one corner to the opposite one, it needs to go through all the 3 cells edges, while instead the RRT and RRT* algorithms are able to directly connect the two corners, without going through all the intermediate cells. Thanks to this observation, a simple math shows that on average, if the diagonal connection were allowed, the path length would have been reduced by a factor of $\sqrt{3}$. This means that the path length generated by A* would have been around $\frac{87}{\sqrt{3}} \approx 50$ m, which is lower than the path length generated by RRT and RRT*, indeed proving once again the concept of optimality of the grid-based algorithms.

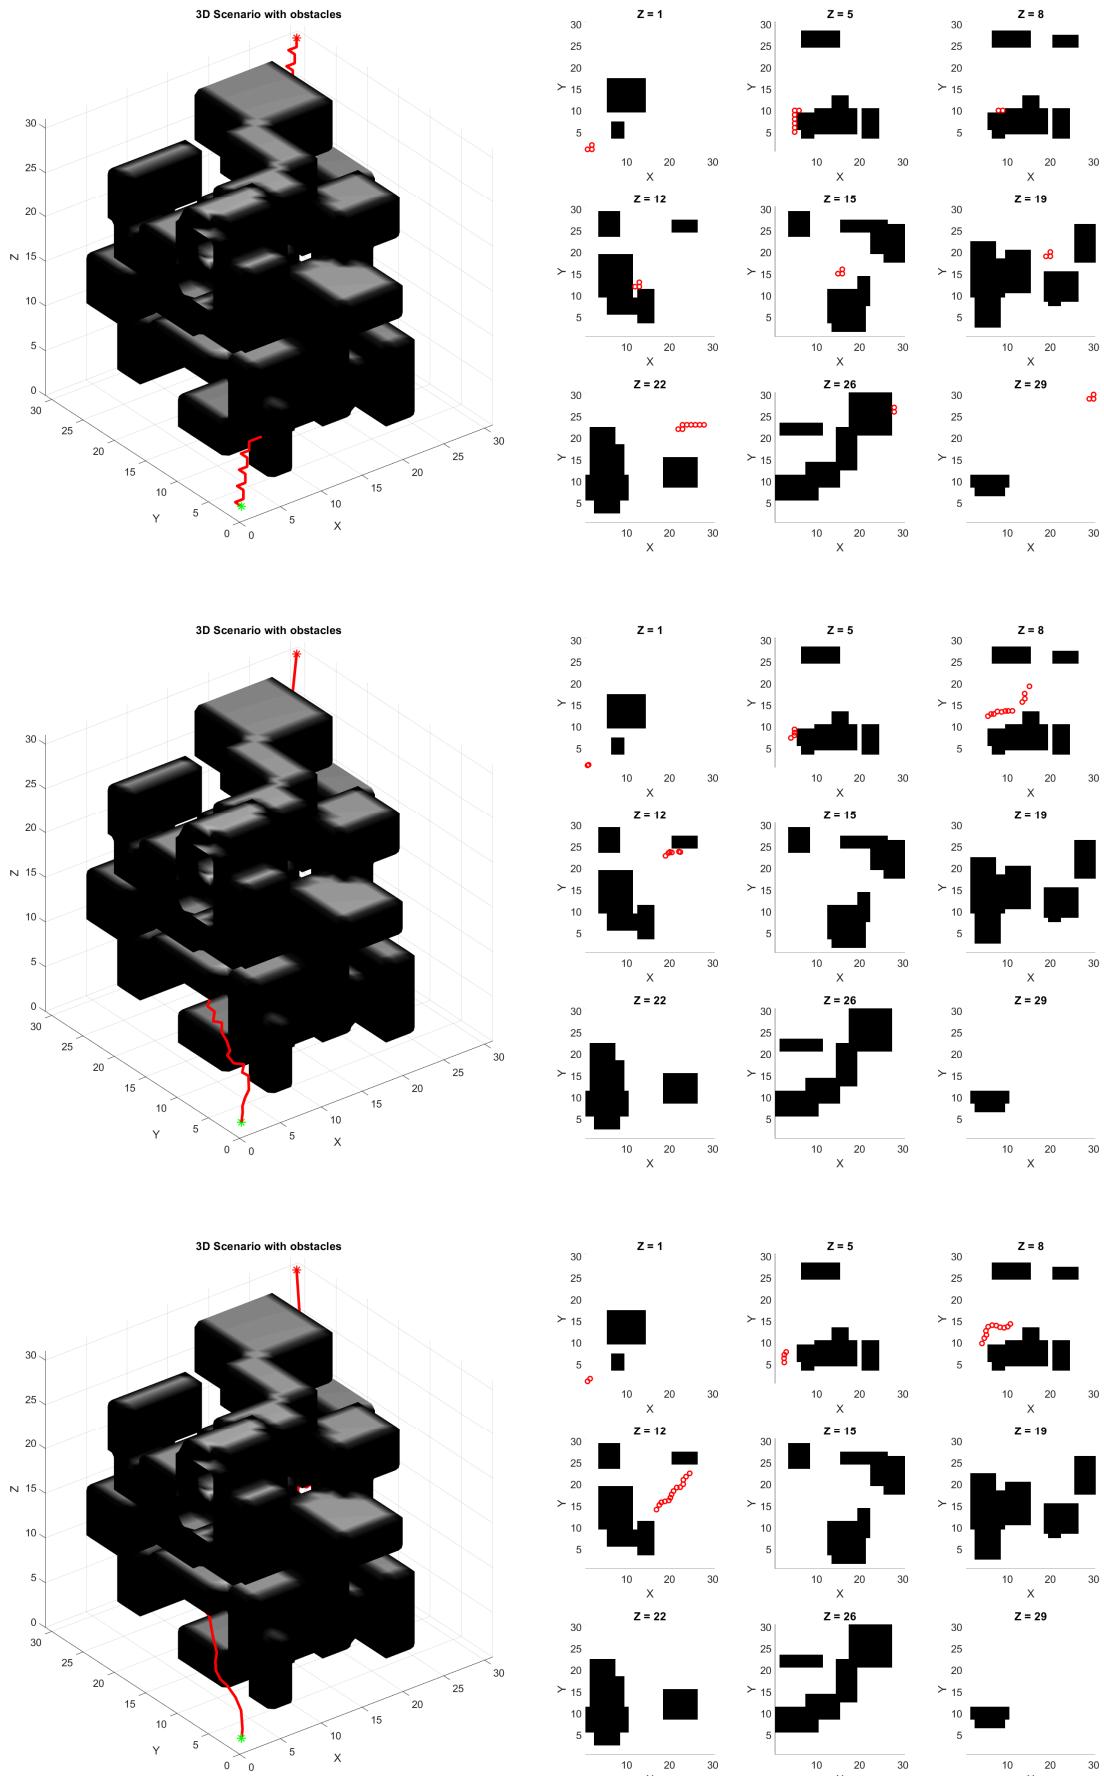


Figure 13: Path generated by A* (top), RRT (middle) and RRT* (bottom) algorithms on the same 3D scenario.

From Table 7 it's clear how the use of grid-based algorithm for single-query planning problem is unfeasible. These tests, again, highlight the strong adaptability of the sample-based algorithms to high-dimensional problems, while grid-based algorithms are not able to cope with the increased complexity of the problem.

Notice that a minor modification to both the RRT and RRT* algorithms have been applied in order to further reduce the computational time. In particular, we have found that the original algorithms excelled at fast exploration of the workspace, reaching in little amount of time every distinct volume closed by obstacles. However, once arrived at the volume of the goal configuration, they failed in creating the final connection and close the planning problem. In order to avoid useless sampling, we have added a simple logic inside the main algorithm loop that checks if the euclidean connection between the lastly added node and the goal is collision-free or not. By doing so, the time required for planning drops significantly.

Listing 5 shows the implementation for this strategy.

```

1 while (Node.euclideanDistance(q_new, node_final.state) > G.step_length && iter < G.
2     max_iter)
3     %
4     if (G.isValidConnection(q_new, node_final.state, map))
5         % We have free space along the connection with the goal
6         % Stop sampleing and directly add the goal node to the tree
7         break
8     end
9 end

```

Listing 5: Implementation of the fast goal connection for sampling-based algorithm.

6 Conclusions

In this project we have implemented tree sampling-based single-query motion planning algorithms, specifically RRT, RRT*, and RRT-kinematic. We have run them on a set of predefined scenarios, and we have compared their performance in terms of path length, computation time, and number of nodes. A comparison with grid-based planning algorithms has also been performed, showing that the tree sampling-based algorithms are more efficient in terms of computation time but lacks the optimality property of grid-based ones.

We are confident that an optimized implementation of both the classes of algorithms would lead to a more efficient and effective solution for motion planning problems. However, we believe that our implementation has already been capable of providing a good overview of the performance of the algorithms and their potential applications in real-world scenarios. In particular, a massive improvement in the performance of the tree sampling-based algorithms can be achieved by using a more efficient data structure for the tree, such as a **KD-tree**. This would allow for a faster nearest neighbor search, which is a key operation in the RRT algorithm. Moreover, we are confident to say that the combination of a quick exploration of the space and a local grid-based search is the key to achieving optimality in motion planning problems. As demonstrated in the multidimensional case, the capabilities of sampling-based algorithms to quickly explore vast volumes of space in a short time are notable. However, once the first sampled is placed in the vicinity of the goal, many more samples are needed before effectively connecting the goal node. This, could be solved by using a local grid-based search, which would work as a local optimizer, reducing the effort needed to simply connect the goal node to the tree.

In conclusion, sampling-based algorithms represent a robust and scalable approach for tackling motion planning problems, especially in high-dimensional or dynamic environments, where fast replanning is often required. Due to their weakness in terms of optimality, hybrid approaches that combine the strengths of both sampling-based and grid-based methods are likely to be the most effective in practice. The integration of these two paradigms could lead to planners that are not only computationally efficient but also capable of producing high-quality solutions across a wide range of applications.

We believe that the future of motion planning lies in the development of hybrid algorithms that can leverage the strengths of both sampling-based and grid-based methods.

References

- [1] Xiangkui Jiang, Zihao Wang, and Chao Dong. A path planning algorithm based on improved rrt sampling region. *Computers, Materials and Continua*, 80(3):4303–4323, 2024.
- [2] Sertac Karaman and Emilio Frazzoli. Sampling-based algorithms for optimal motion planning. *International Journal of Robotic Research - IJRR*, 30:846–894, 06 2011.